

The Big Picture

Thierry Sans


Goals of this lecture

- Define what an Operating System is
- Explain how an OS works in a nutshell
- Bridge the gap between hardware (CSCB58) and systems programming (CSCB09)
- Give an overview of the course content and projects

The big picture in 5 pieces


The need for bootstrapping	
The need for concurrency	project 1
The need for user programs	project 2
The need for virtual memory	project 3
The need for a filesystem	project 4

0x FF FF FF FF



Simple Computer Architecture

Memory + CPU



for a more accurate and detailed map of the x86 memory
look at [https://wiki.osdev.org/Memory_Map_\(x86\)](https://wiki.osdev.org/Memory_Map_(x86))

Each processor has its Instruction Set Architecture (ISA)

Processor executes instructions stored in memory

- ➡ Each instruction is a bit string that the processor understands as an operation
 - arithmetic
 - read/write bit strings
 - bit logic
 - jumps
- ✓ ~2000 instructions on modern x86-64 processors

0x FF FF FF FF

I/O

stack

stack pointer (esp)

heap

heap

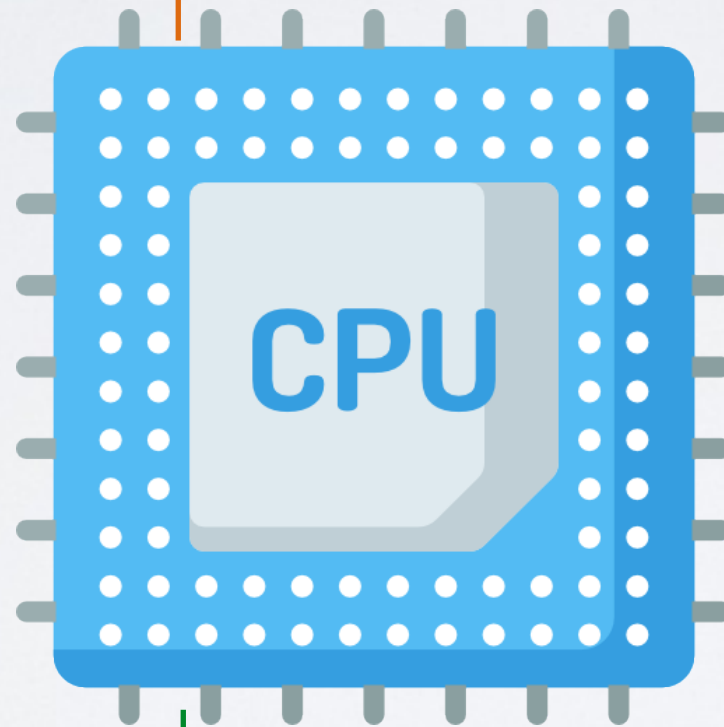
code (text)

instruction pointer (eip)

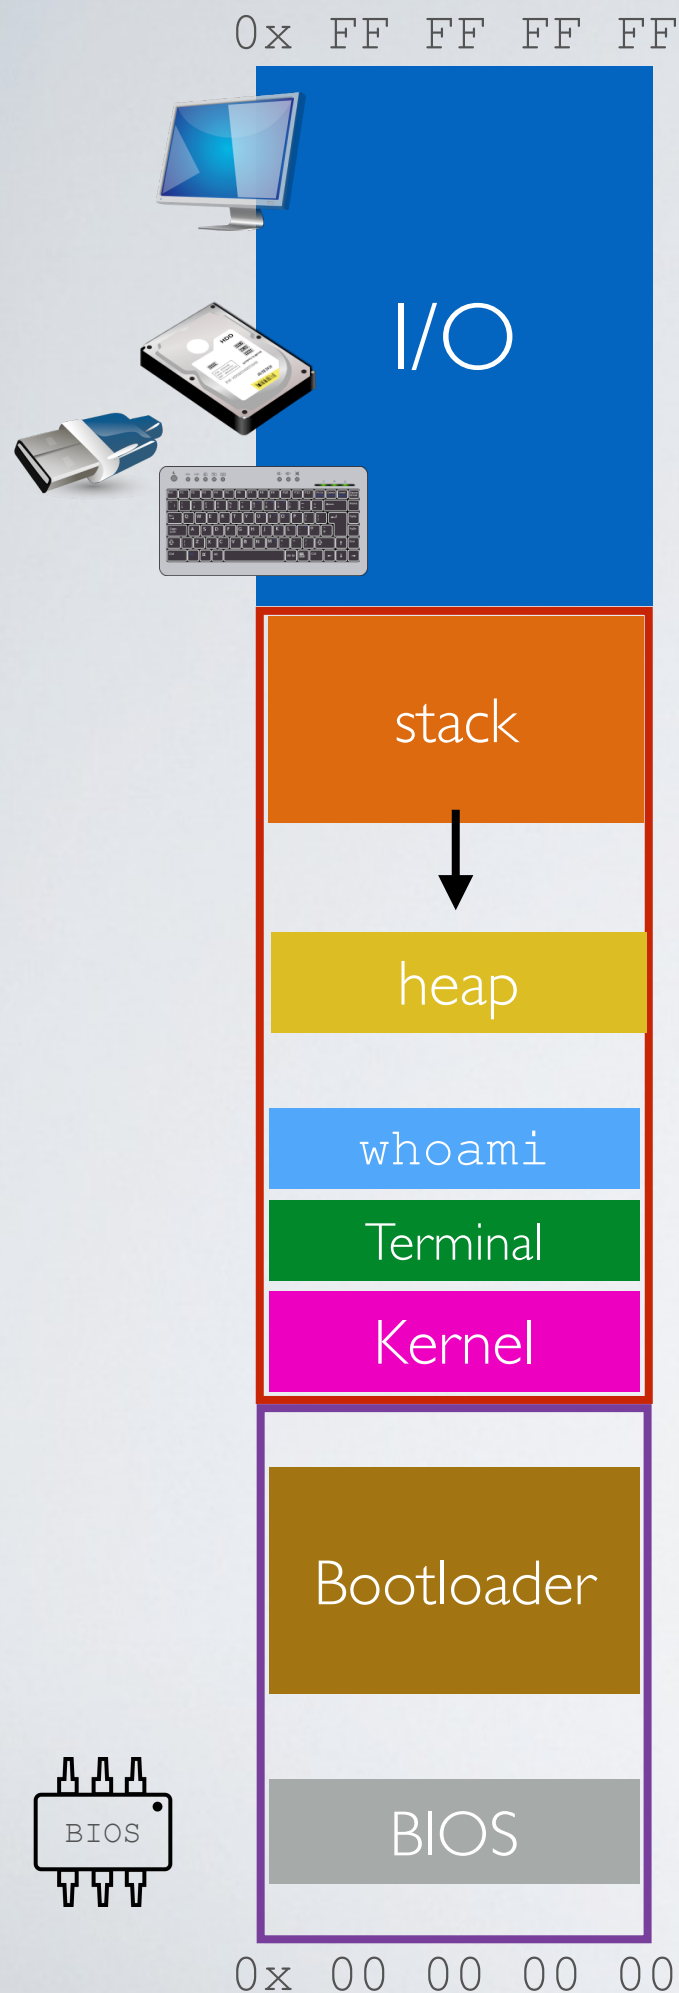
Boot

0x 00 00 00 00

Running one program



The need for **bootstrapping**



Bootstrapping

Step 5: using the terminal, users can execute programs (e.g Bash terminal) ... and repeat

Step 4: the kernel starts the user-interface program (e.g Bash terminal)

Step 3: the bootloader loads the OS kernel in RAM

Step 2: the BIOS loads the **bootloader** from a device (hard-drive, USB, network ...) based on the configuration


Step 1: Power -on! The CPU starts executing code contained in the **BIOS** (basic input/output system)

The need for **concurrency**

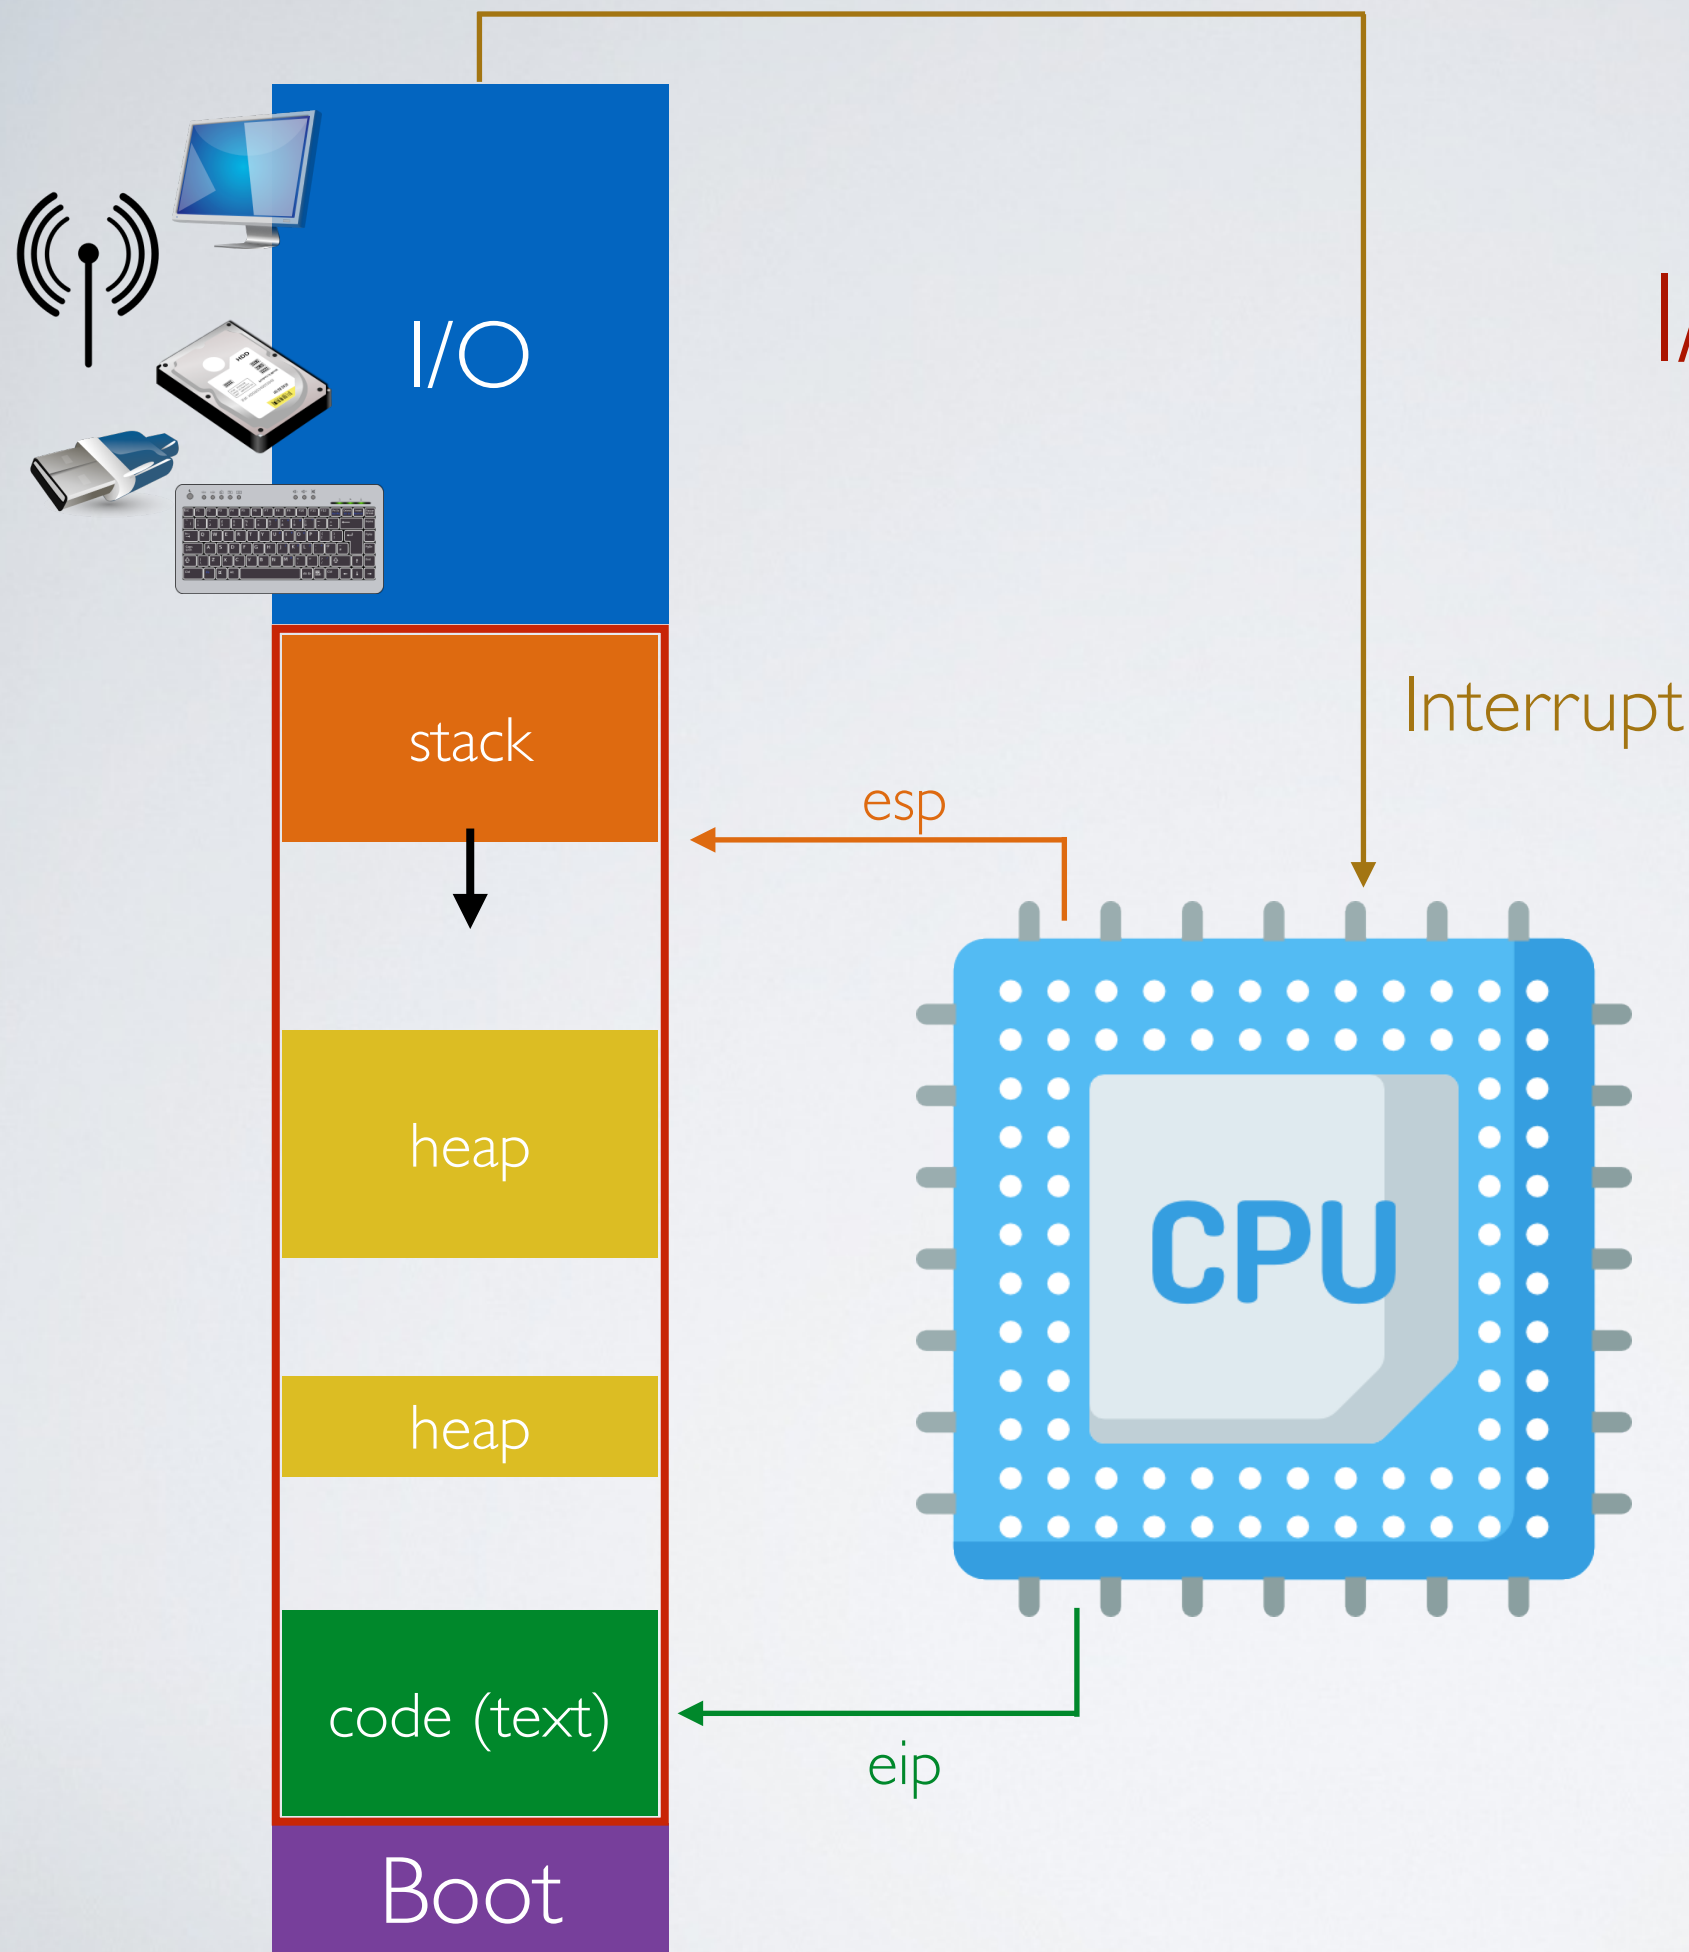
Running multiple programs one after the other



Problem: the CPU is waiting for I/O (polling)

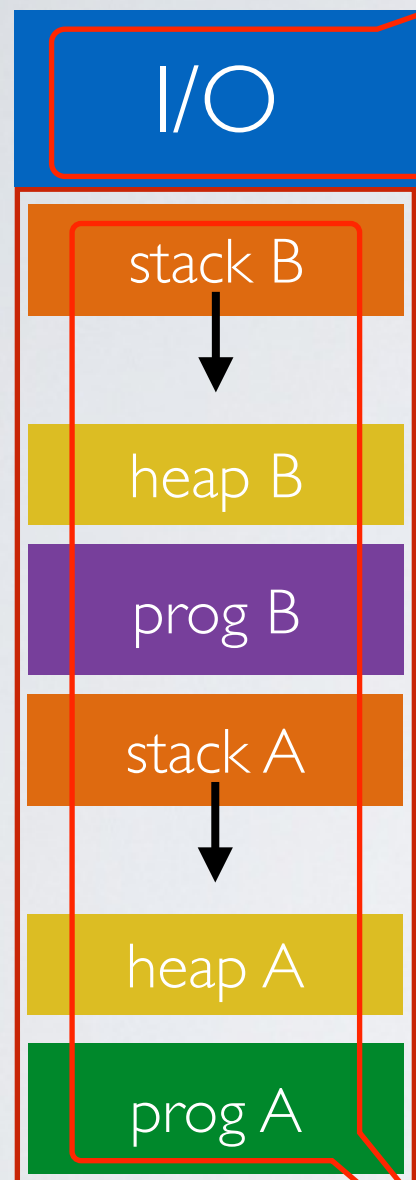


Problem: the programs must co-exists in memory
(coming next with virtual memory)



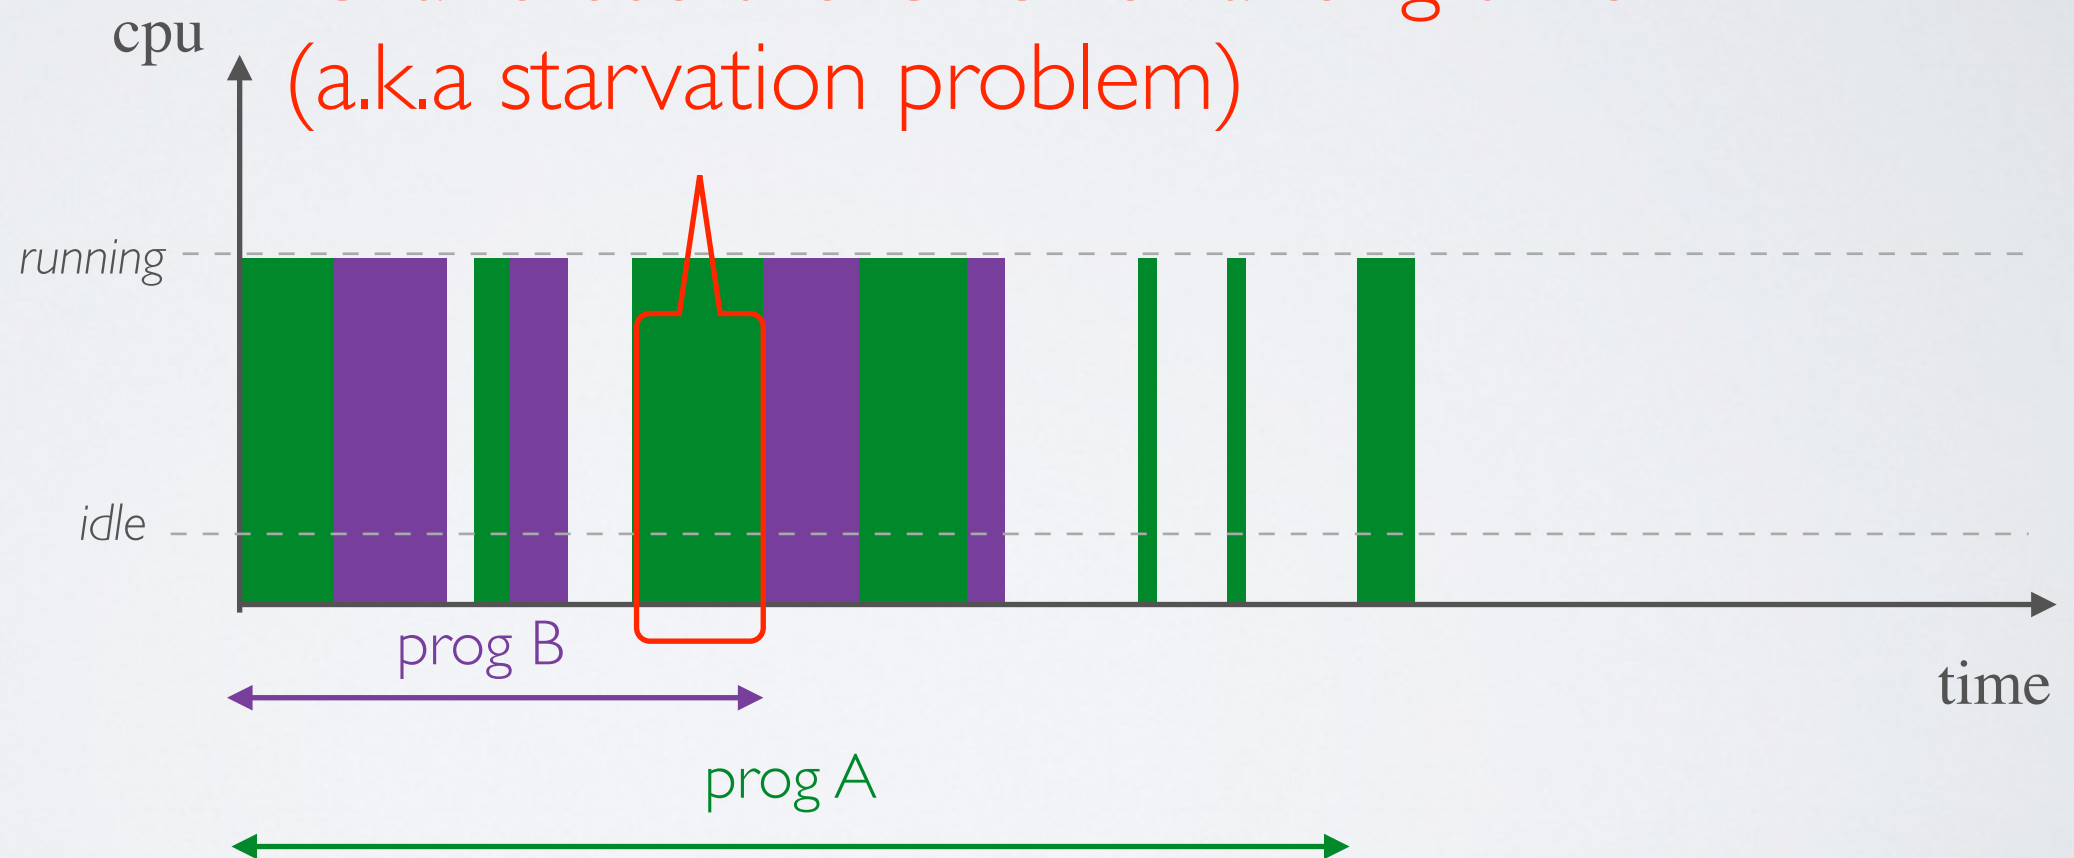
I/O with interrupts

Running multiple programs concurrently



Problem: concurrent access to I/O devices must be synchronized

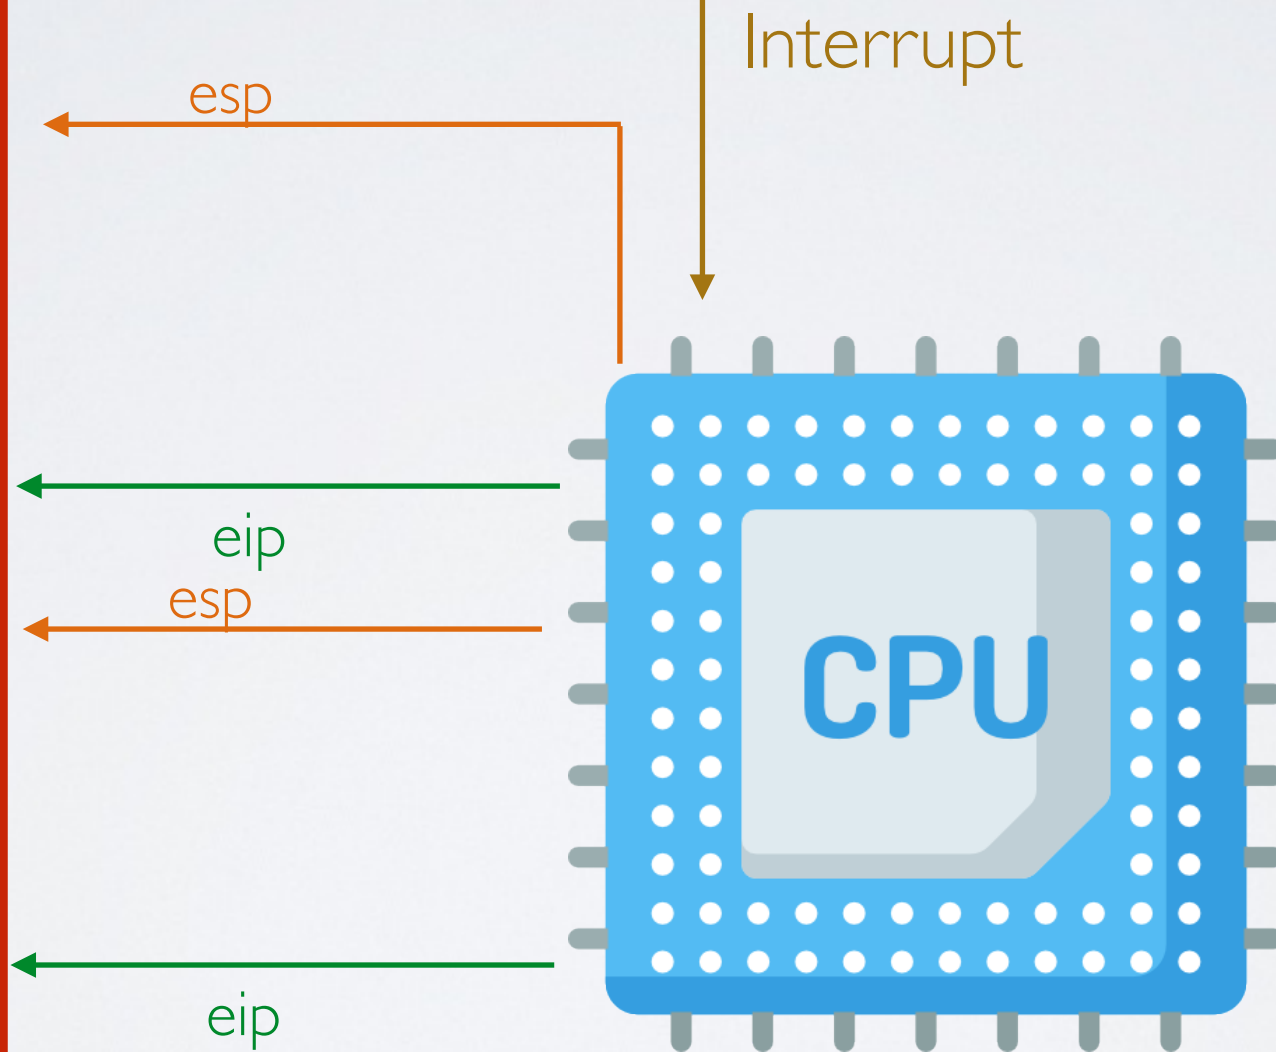
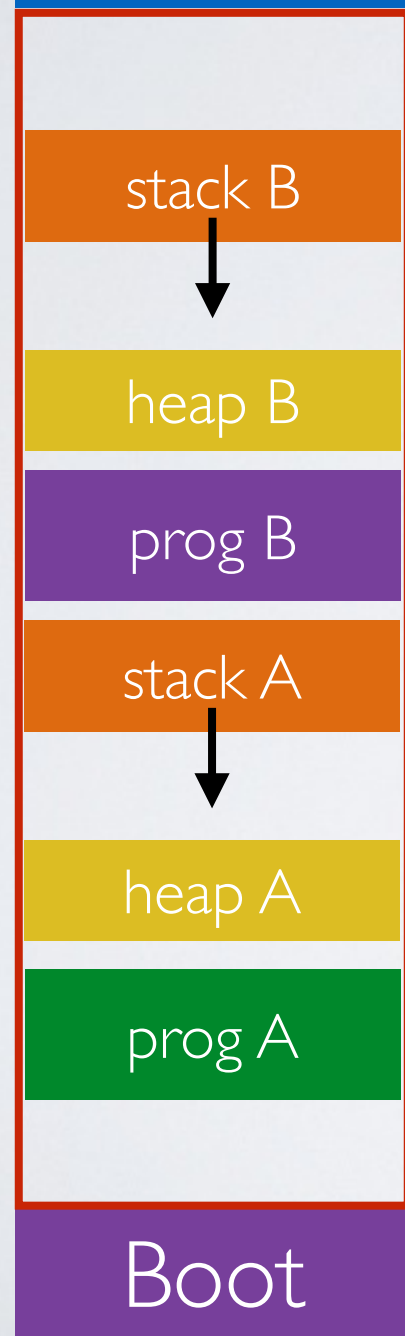
Problem: what if the program does not do any IO and use the CPU for a long time (a.k.a starvation problem)



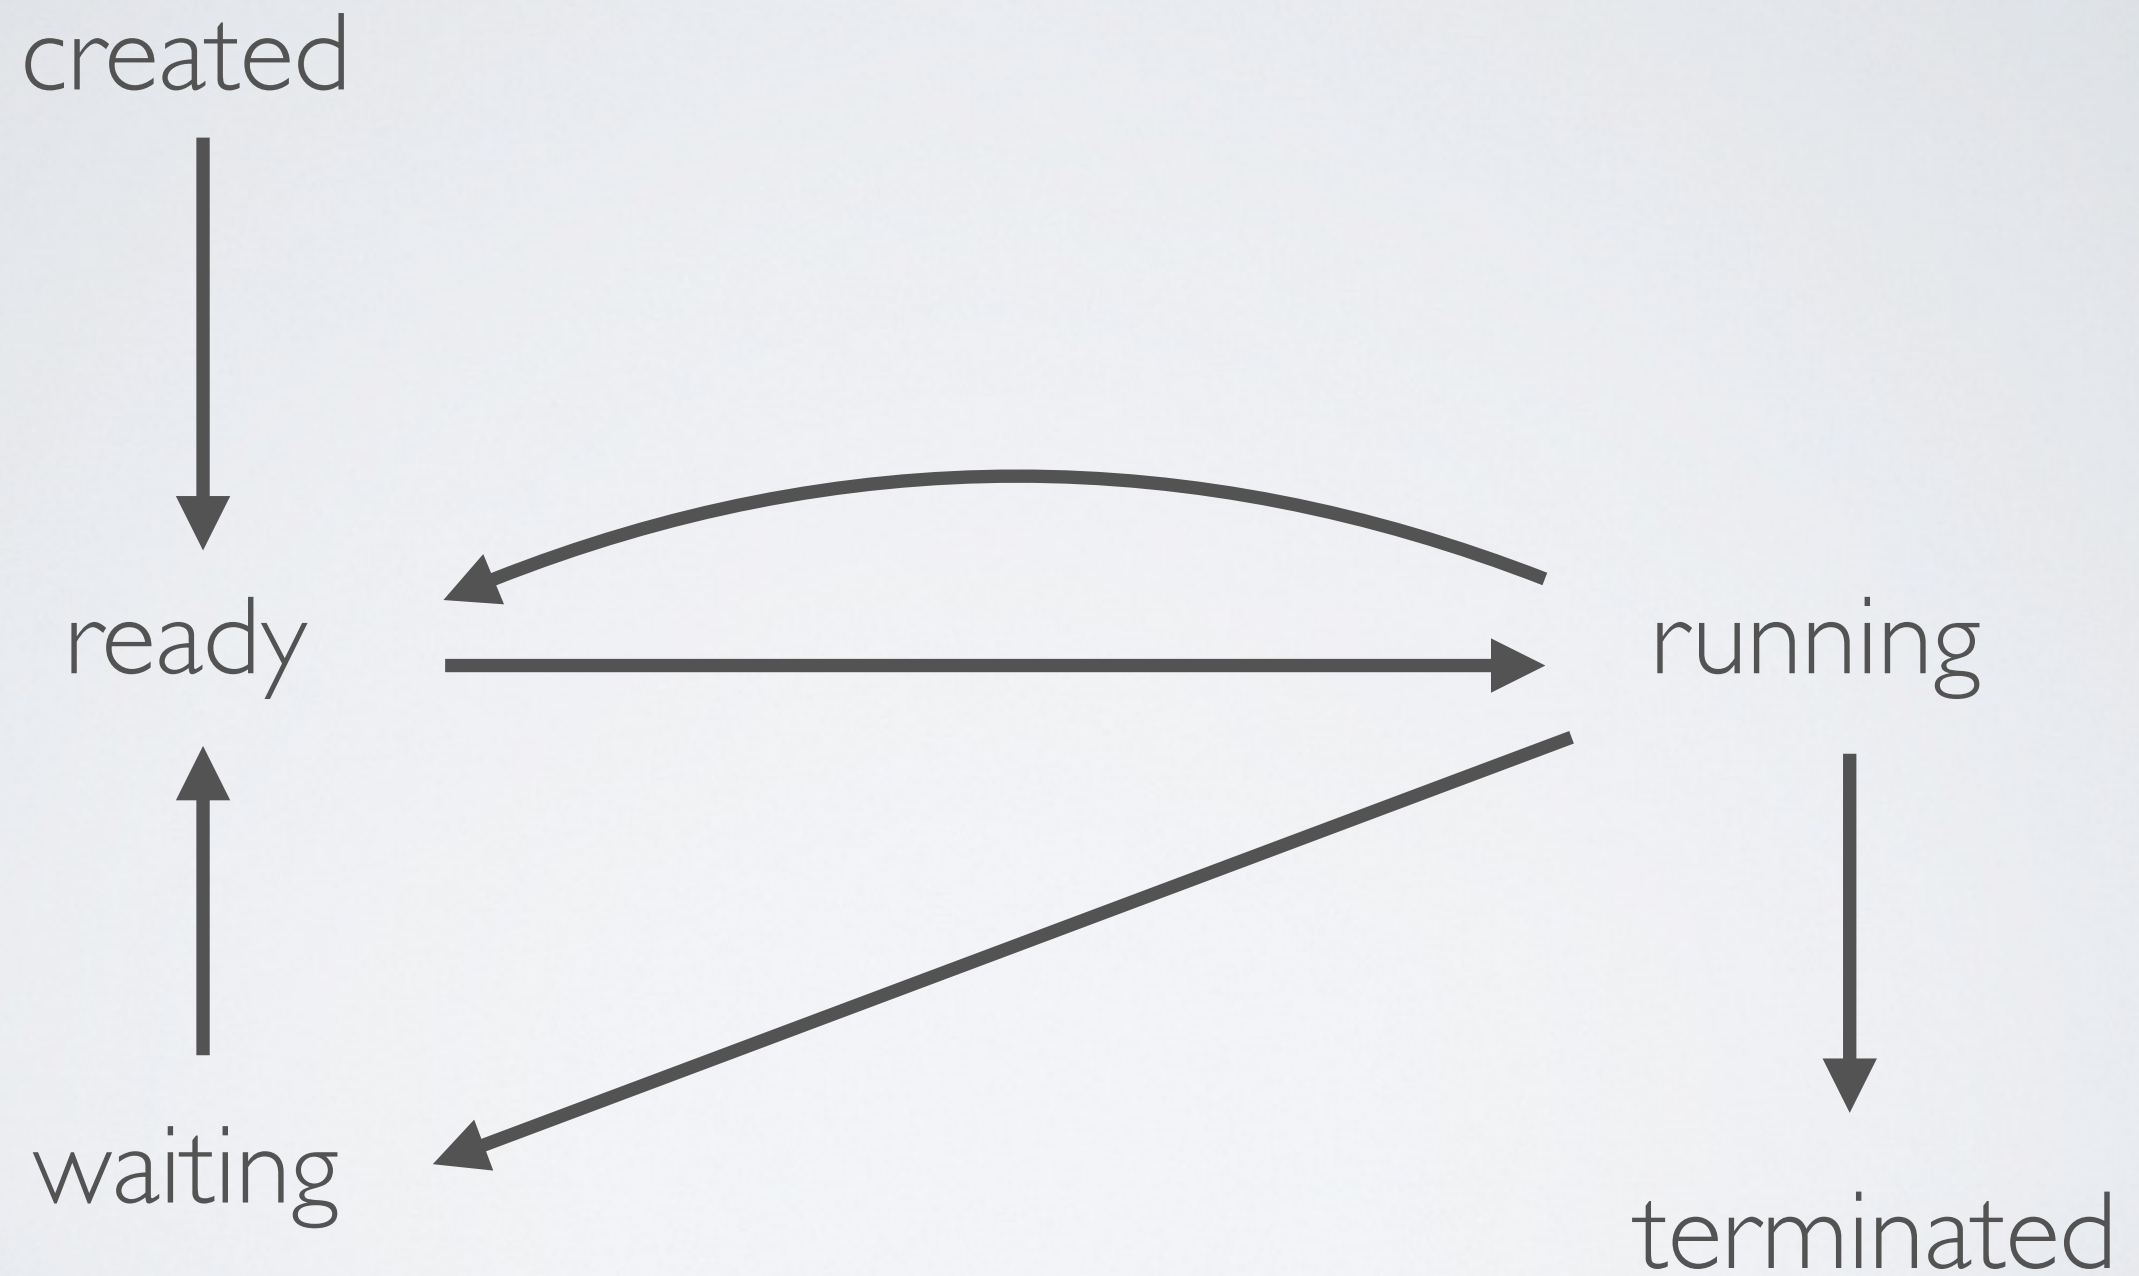
Problem: the programs and their stacks must co-exists in memory (coming next with virtual memory)



Using the clock
to trigger an interrupt



Program States



Other problems that we are going to address during the semester

- **Scheduling**

Decide which process to execute when several are ready to be run

- **Synchronization**

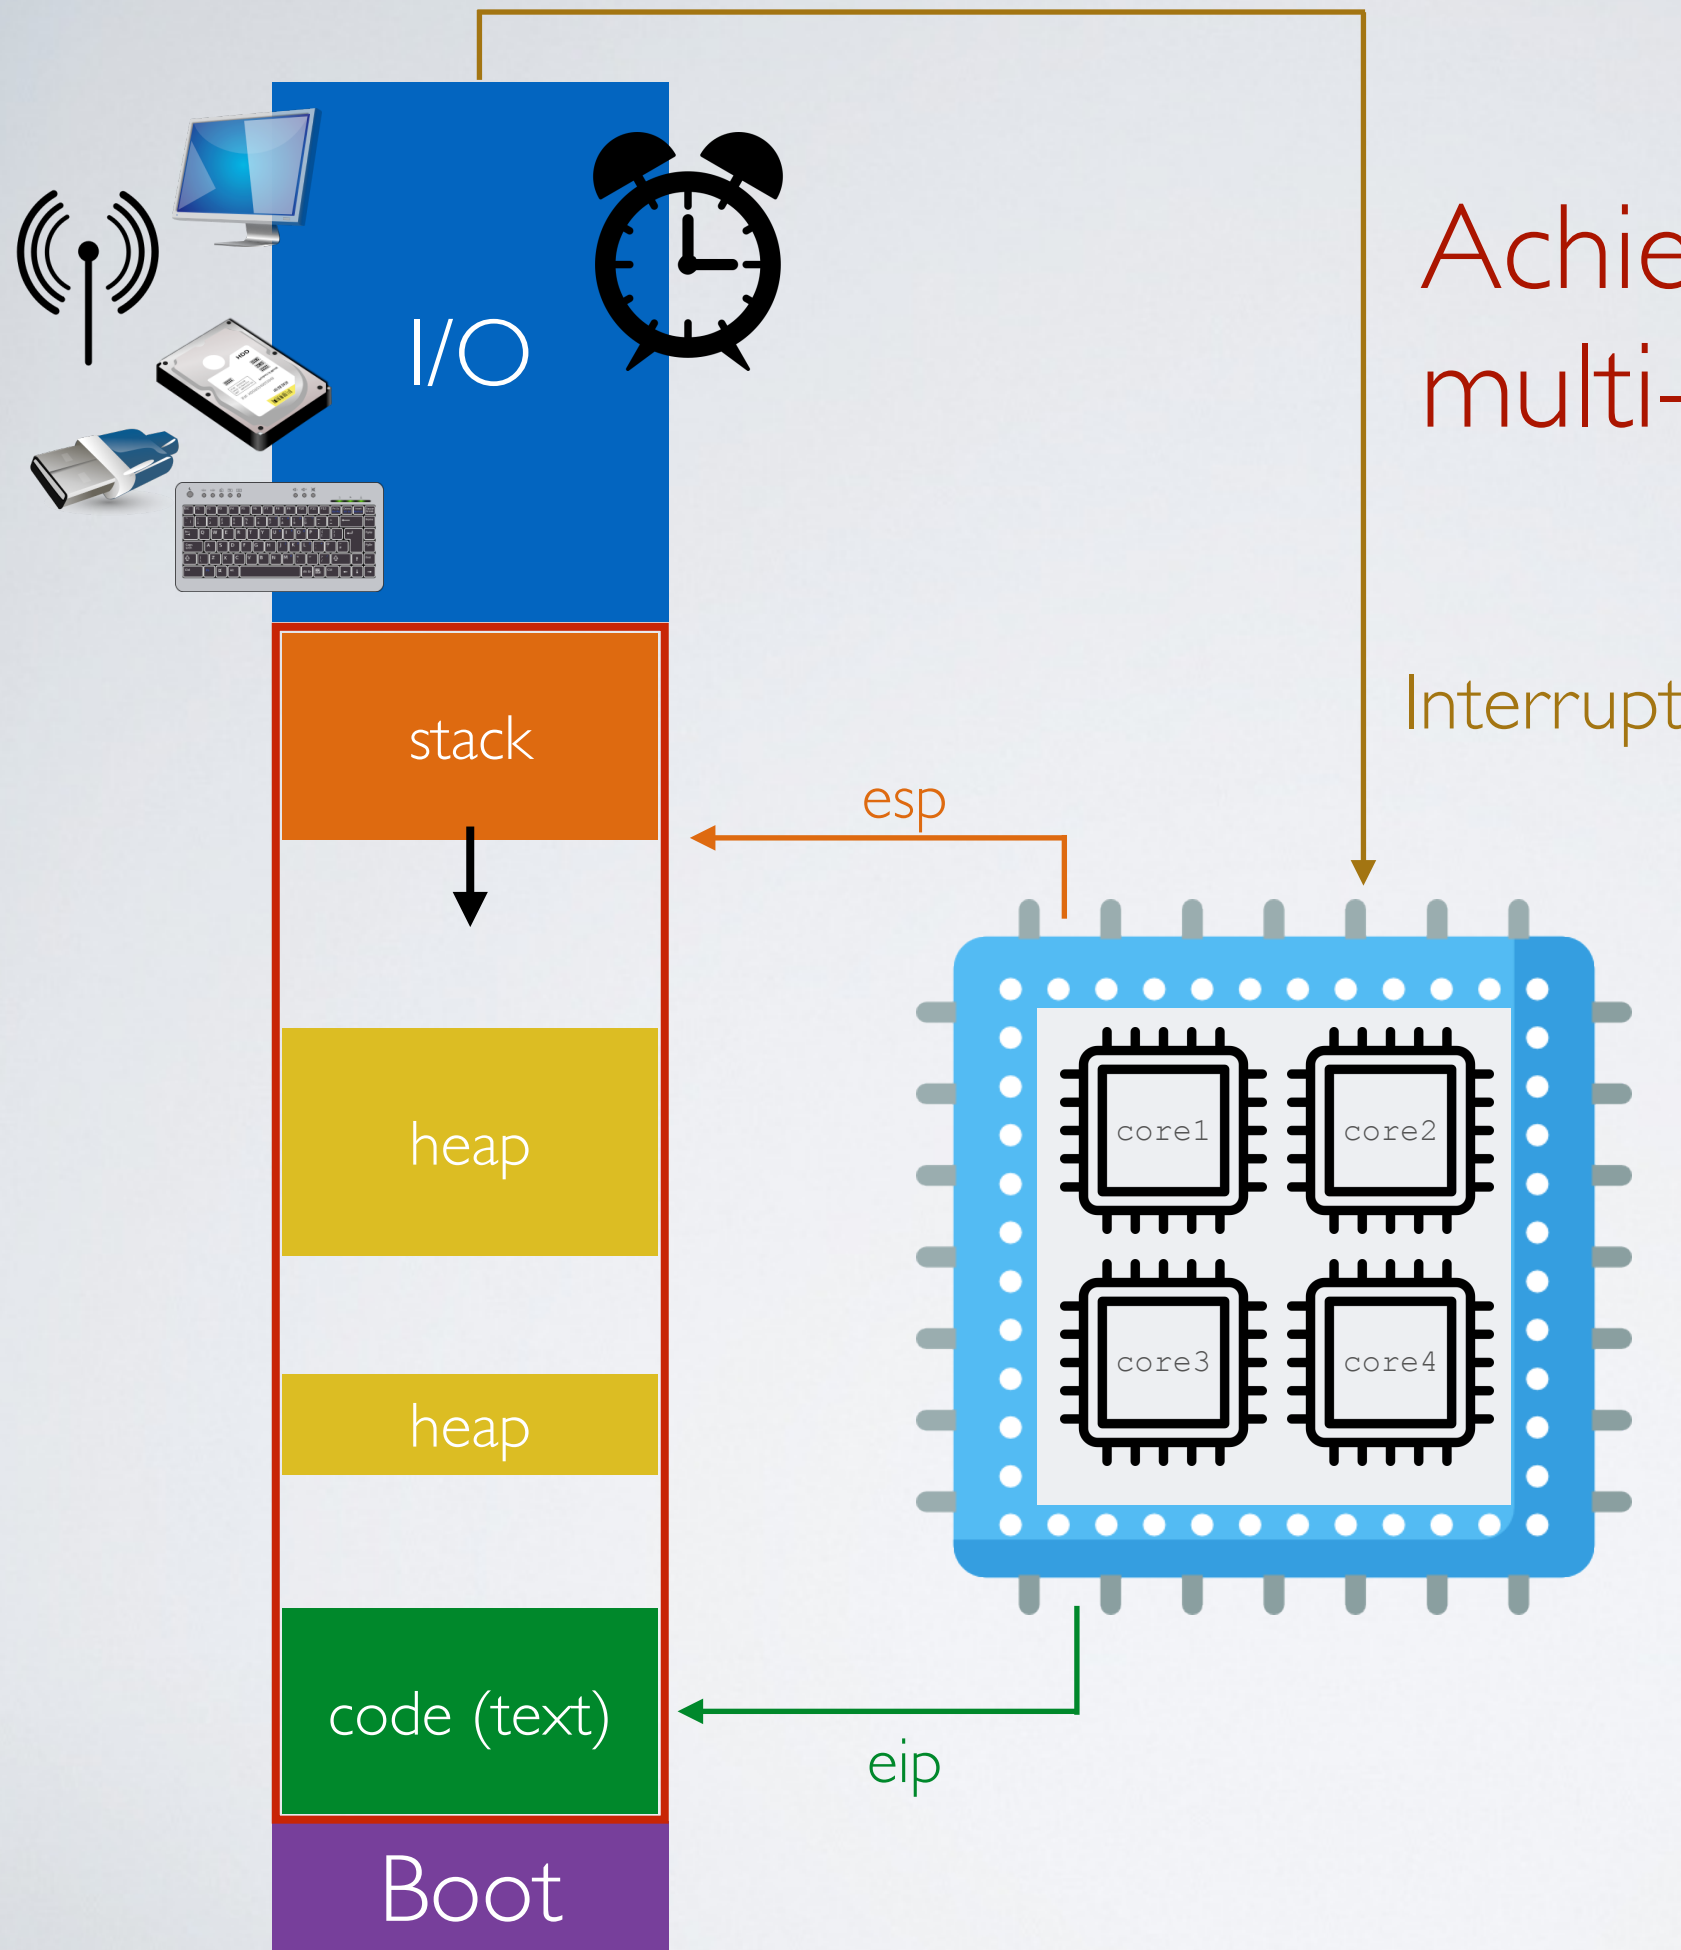
Manage concurrent access to resources using semaphores, locks, monitors

- **Communication**

Exchange messages between processes using IPC (sockets & signals)

- **Threads**

Lightweight concurrency within a process



Achieving parallelism with multi-core processors

The need for **user programs**

The need for abstraction for user programs

How to write a user program like the *Bash* shell that reads keyboard inputs from the user?

➡ Read input data from the I/O device directly? But which one?

- The one connected to the PS2 port?
- The one connected to the USB?
- The one connected to the bluetooth?
- The remote one connected to the network?

⦿ User programs do not operate I/O devices directly

✓ The OS abstracts those functionalities and provide them as **system calls**

System Calls

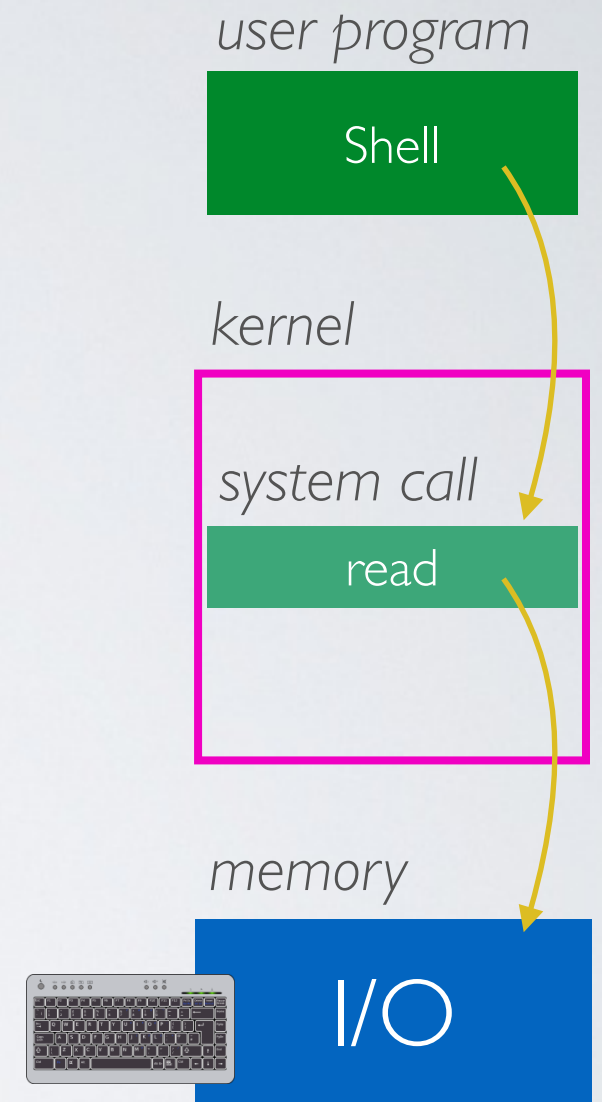
➔ Provide user programs with an API to use the services of operating system

There are 5 categories of system calls

- Process control
- File management
- Device management
- Information/maintenance (system configuration)
- Communication (IPC)
- Protection

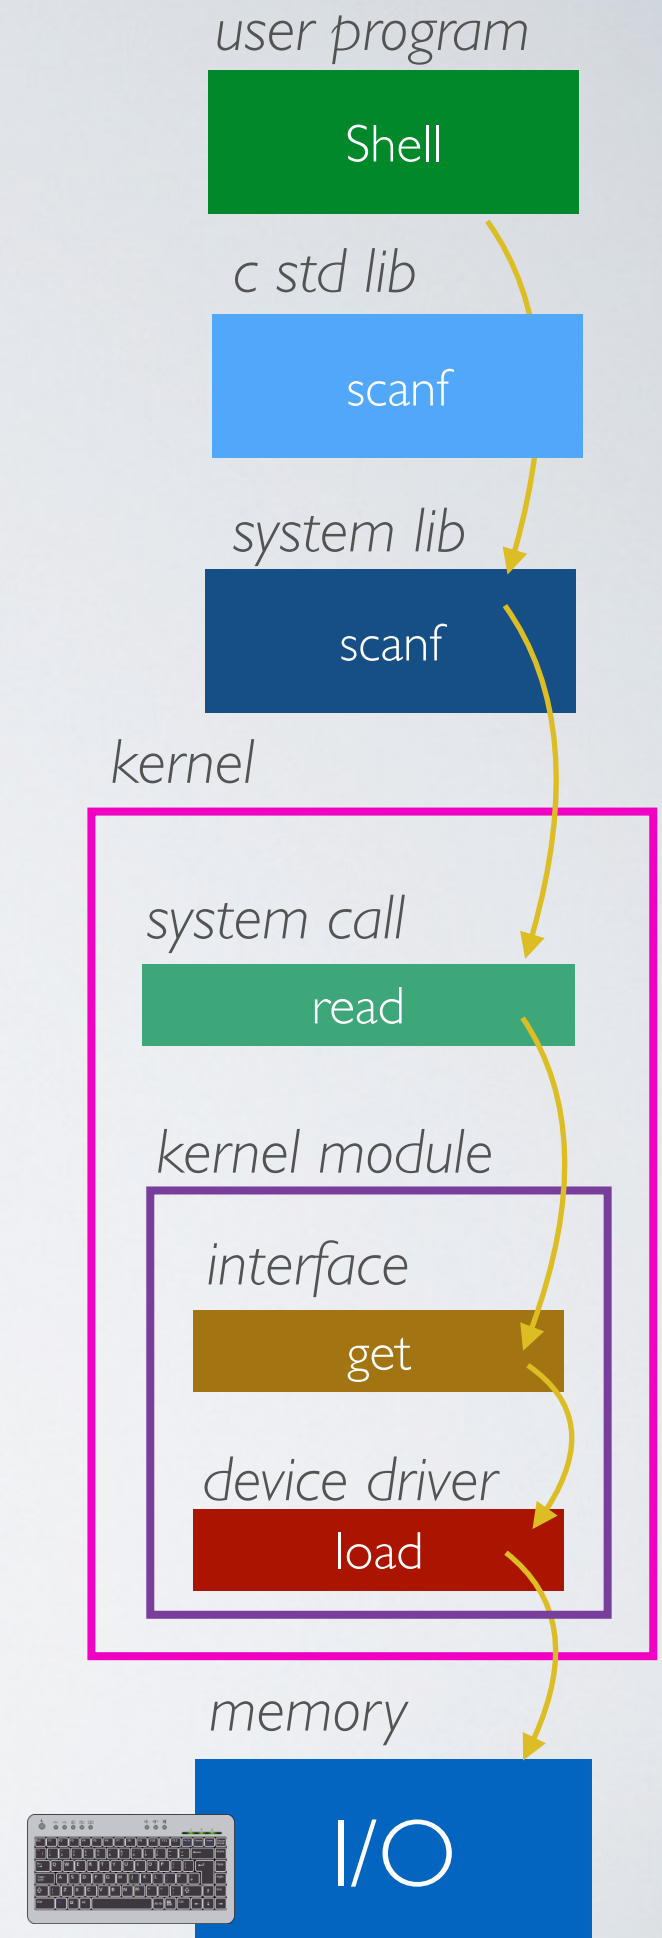
✓ There are 393 system calls on Linux 3.7

http://www.cheat-sheets.org/saved-copy/Linux_Syscall_quickref.pdf



In reality, many (many) level of abstraction and modularity

➔ This is what makes developing OS very challenging (CSCB07)



With concurrency

- ✓ **From the system perspective**

better CPU usage resulting in a faster execution overall
(but not individually)

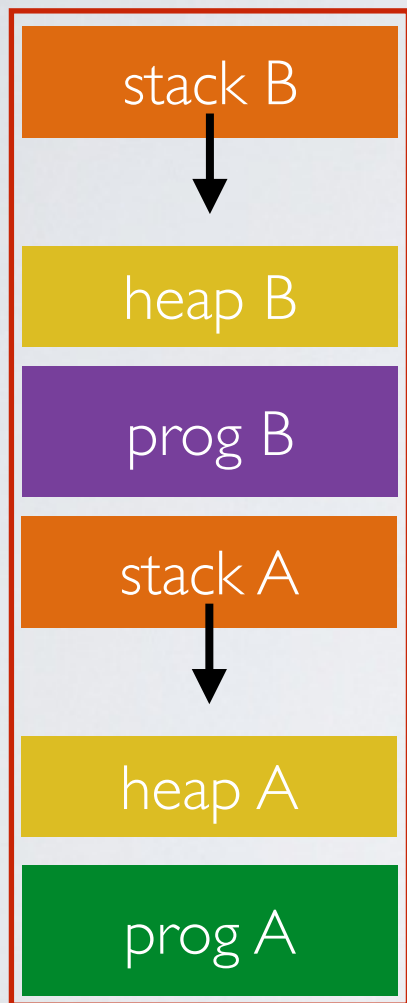
- ✓ **From the user perspective**

programs seem to be executed in parallel

➡ But it requires scheduling, synchronization and some protection mechanisms

The need for **virtual memory**

The problem of managing the memory



How to make programs and execution contexts co-exists in memory?

- ✓ Placing multiple execution contexts (stack and heap) at random locations in memory is not a problem ...
... well, as long as you have enough memory
- ⦿ However having programs placed at random locations is problematic

Let's look at some C code and its binary

```
#include <stdio.h>

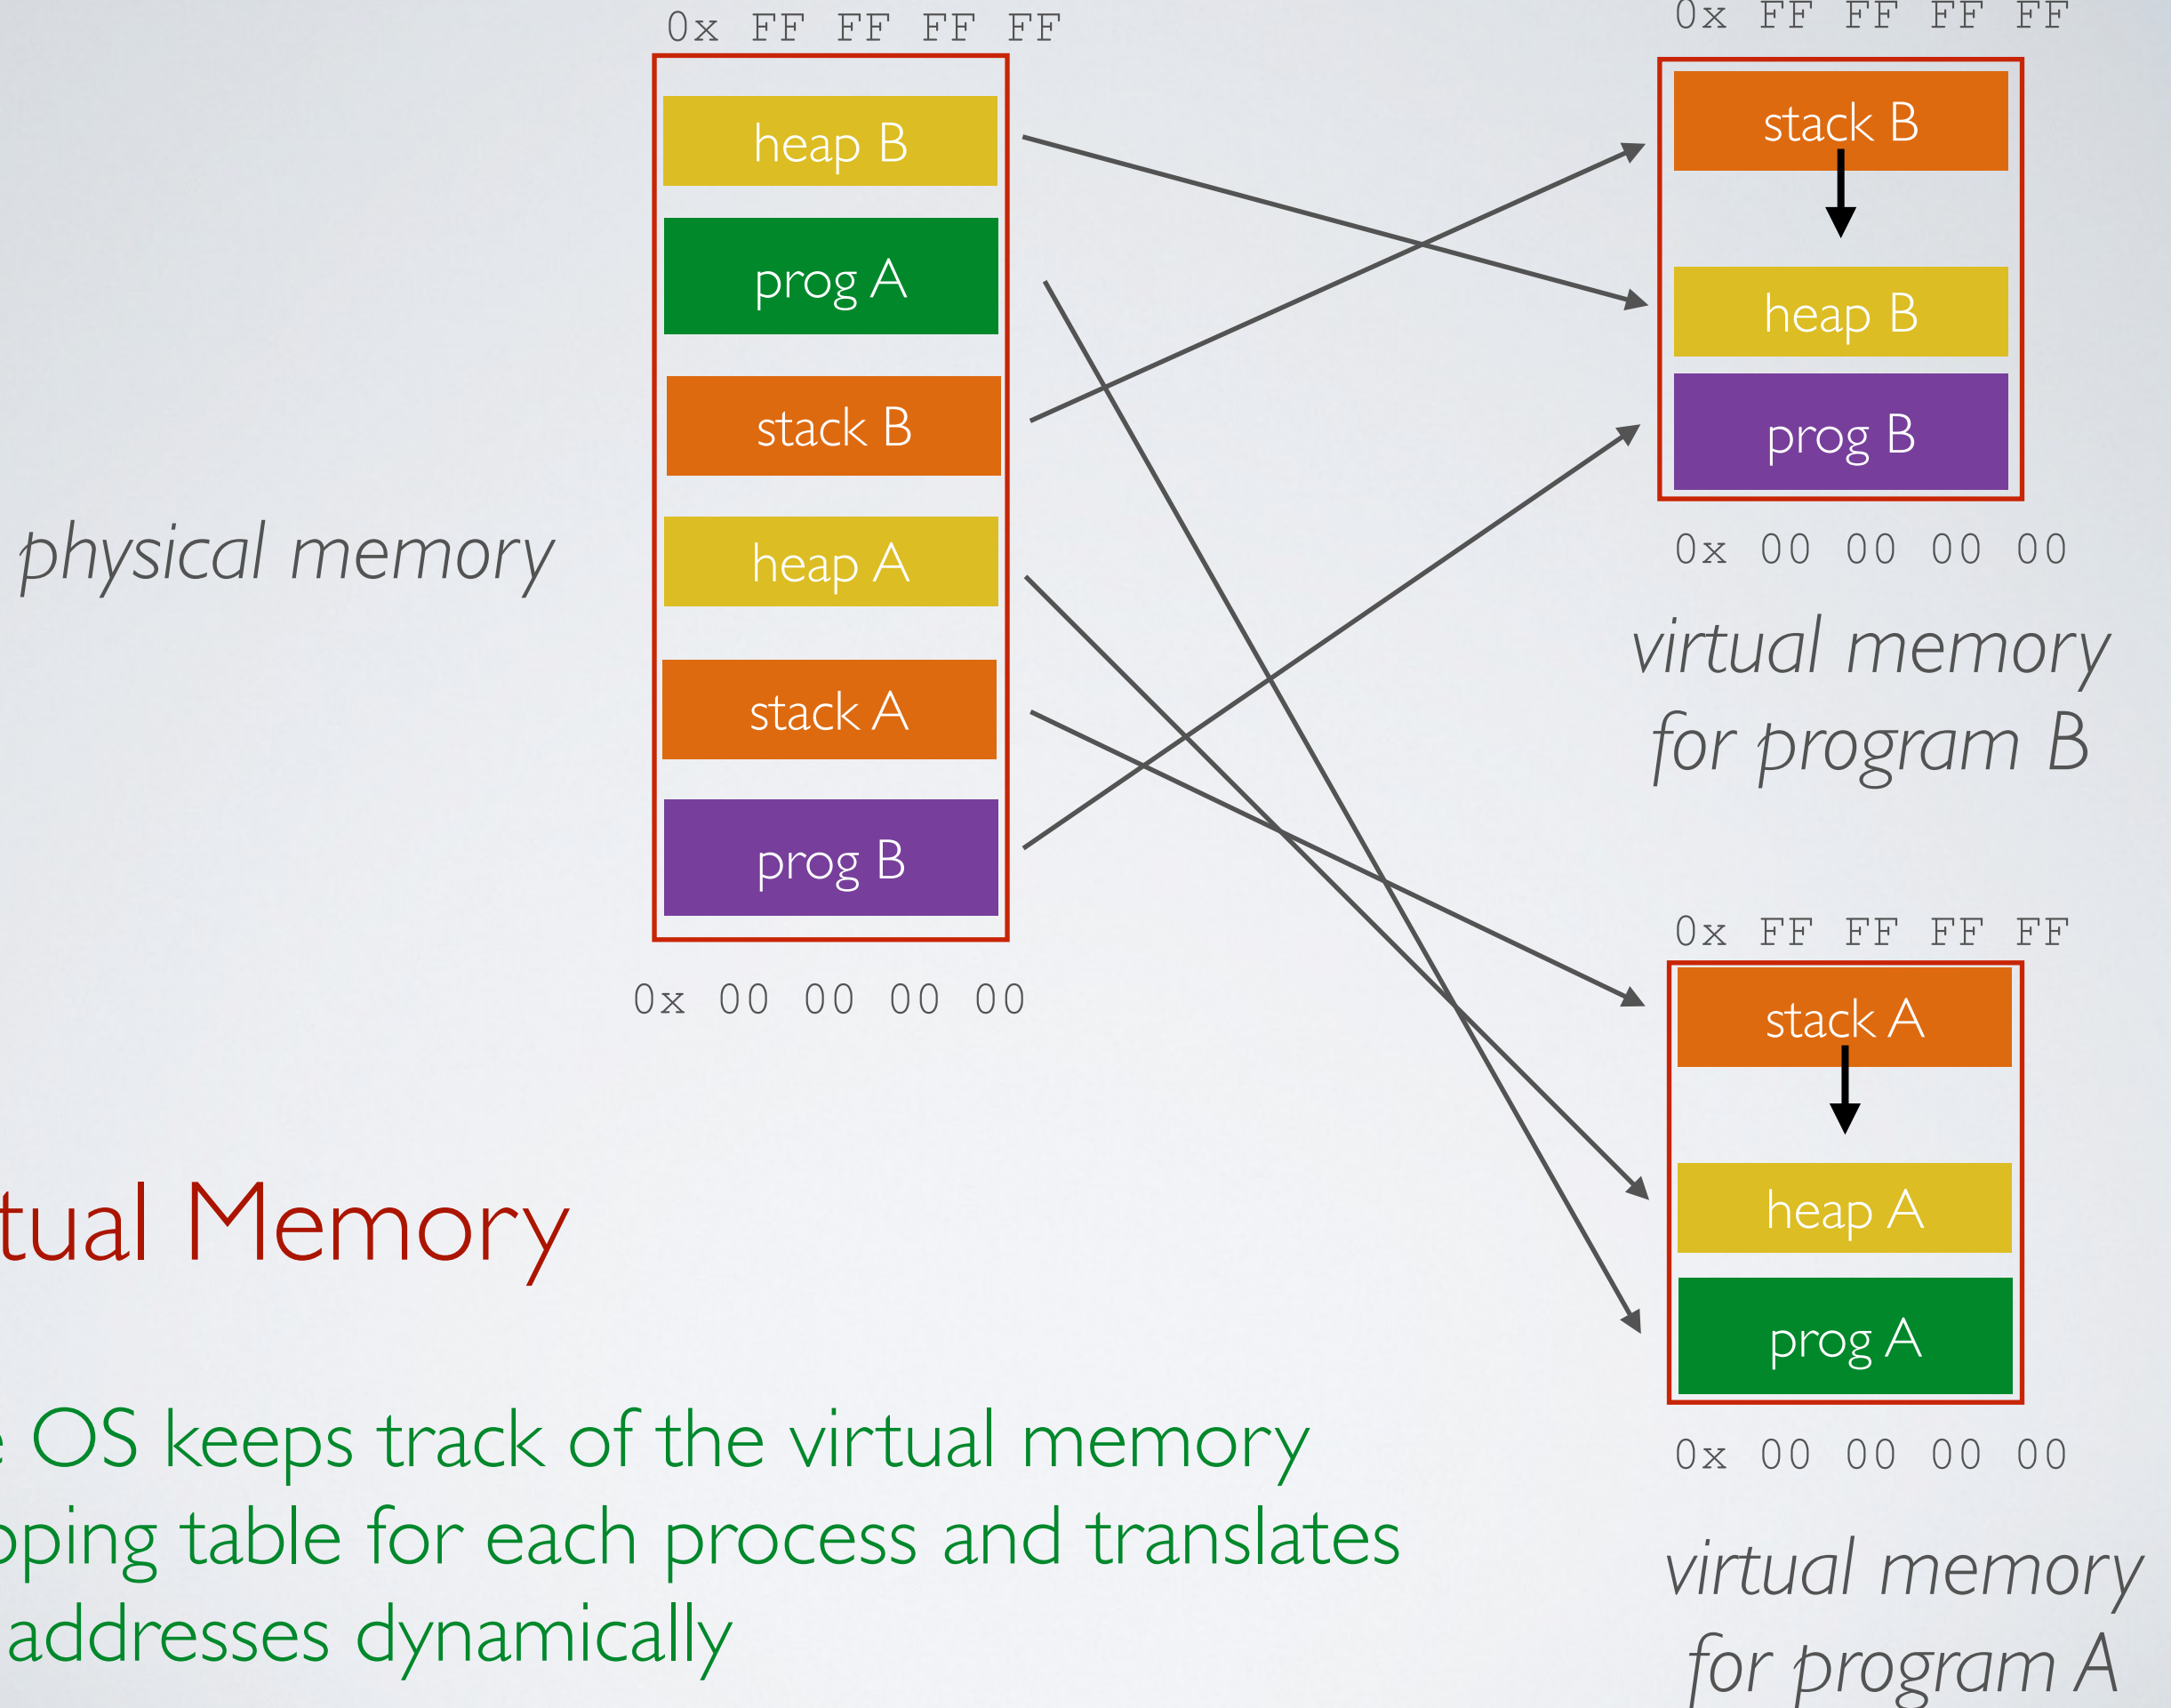
int foo(){
    printf("hello world!");
}

int main(int argc, char **argv){
    foo();
}
```

Since function addresses and others are hard-encoded in the binary, the program cannot be placed at random locations in memory

```
0804840b <foo>:
804840b: 55                push    ebp
804840c: 89 e5            mov     ebp,esp
804840e: 83 ec 08        sub     esp,0x8
8048411: 83 ec 0c        sub     esp,0xc
8048414: 68 d0 84 04 08   push    0x80484d0
8048419: e8 c2 fe ff ff   call    80482e0 <printf@plt>
804841e: 83 c4 10        add     esp,0x10
8048421: 90              nop
8048422: c9              leave
8048423: c3              ret

08048424 <main>:
8048424: 8d 4c 24 04      lea     ecx,[esp+0x4]
8048428: 83 e4 f0        and     esp,0xffffffff
804842b: ff 71 fc        push    DWORD PTR [ecx-0x4]
804842e: 55              push    ebp
804842f: 89 e5            mov     ebp,esp
8048431: 51              push    ecx
8048432: 83 ec 04        sub     esp,0x4
8048435: e8 d1 ff ff ff   call    804840b <foo>
804843a: b8 00 00 00 00   mov     eax,0x0
804843f: 83 c4 04        add     esp,0x4
8048442: 59              pop     ecx
8048443: 5d              pop     ebp
8048444: 8d 61 fc        lea     esp,[ecx-0x4]
8048447: c3              ret
8048448: 66 90           xchg    ax,ax
804844a: 66 90           xchg    ax,ax
804844c: 66 90           xchg    ax,ax
804844e: 66 90           xchg    ax,ax
```



Another problem

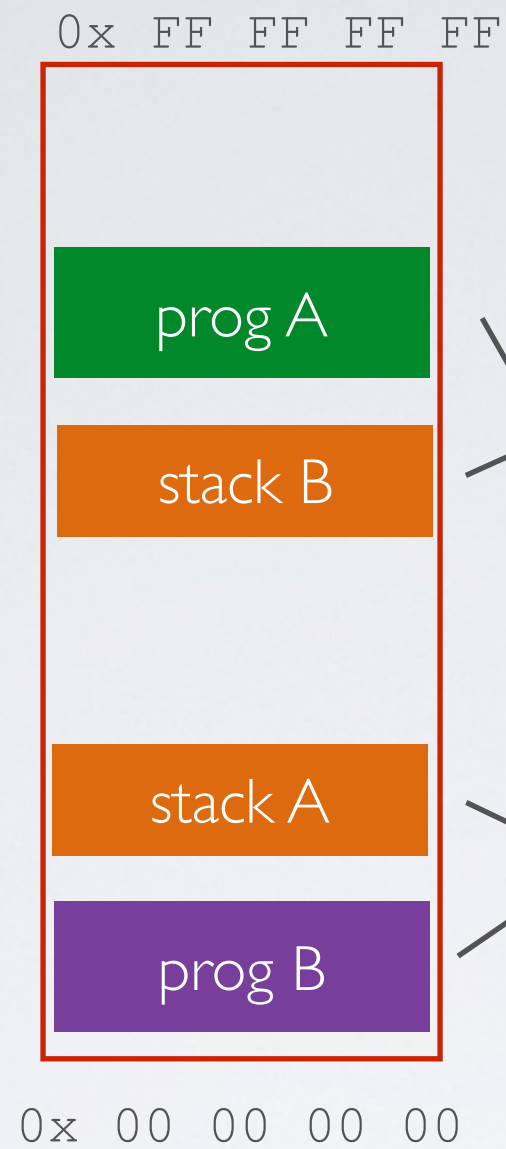
What if we run out of memory because of too many concurrent programs?

✓ Swap memory
move some data to the disk

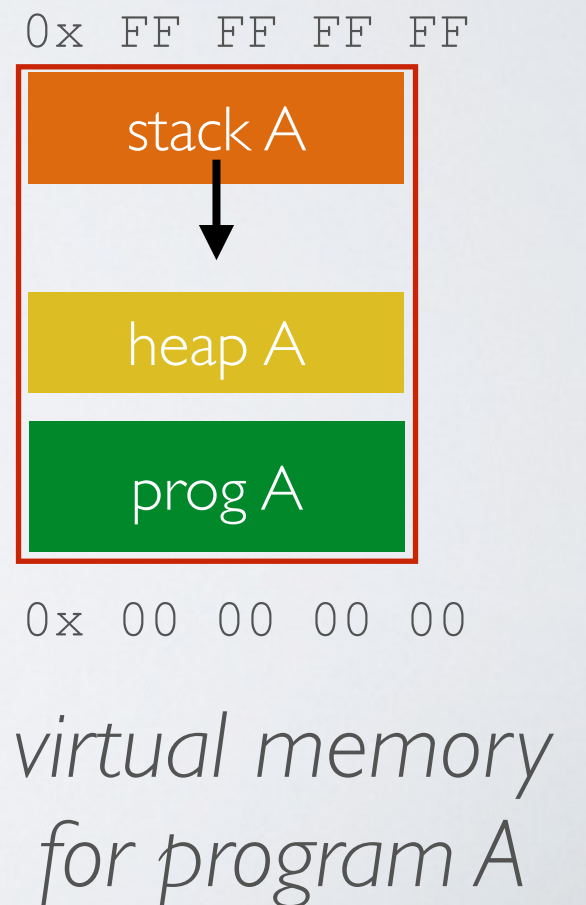
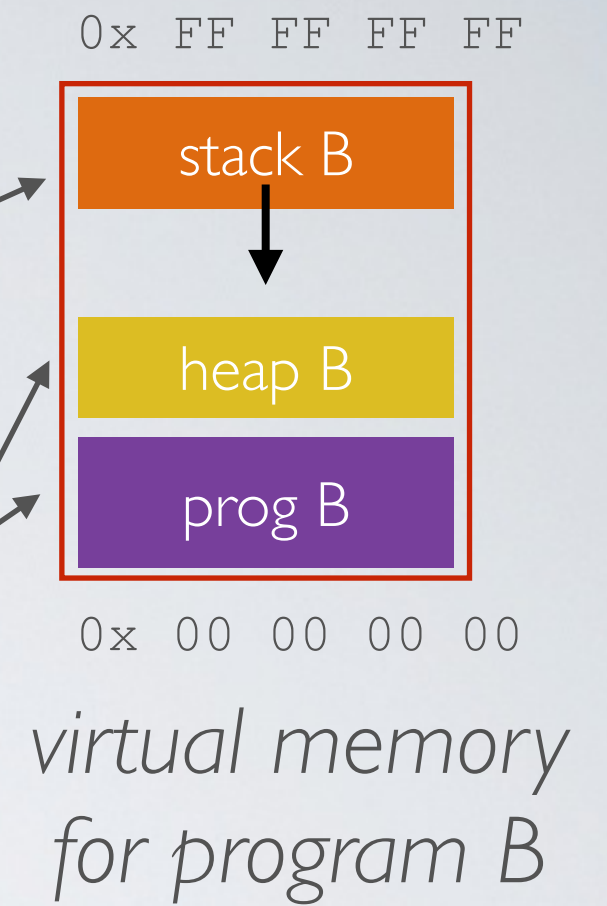
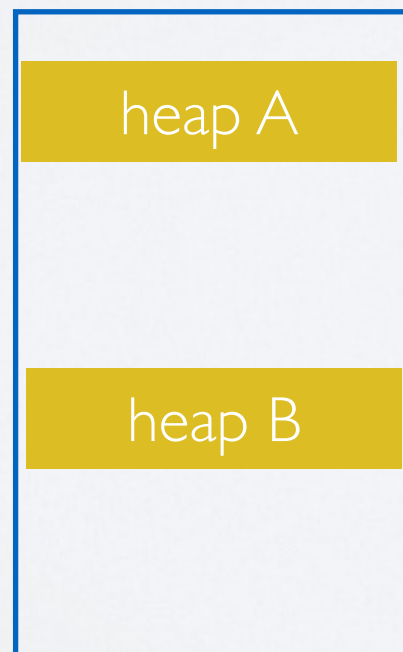
➡ Managing memory becomes very complex
but necessary

Swap

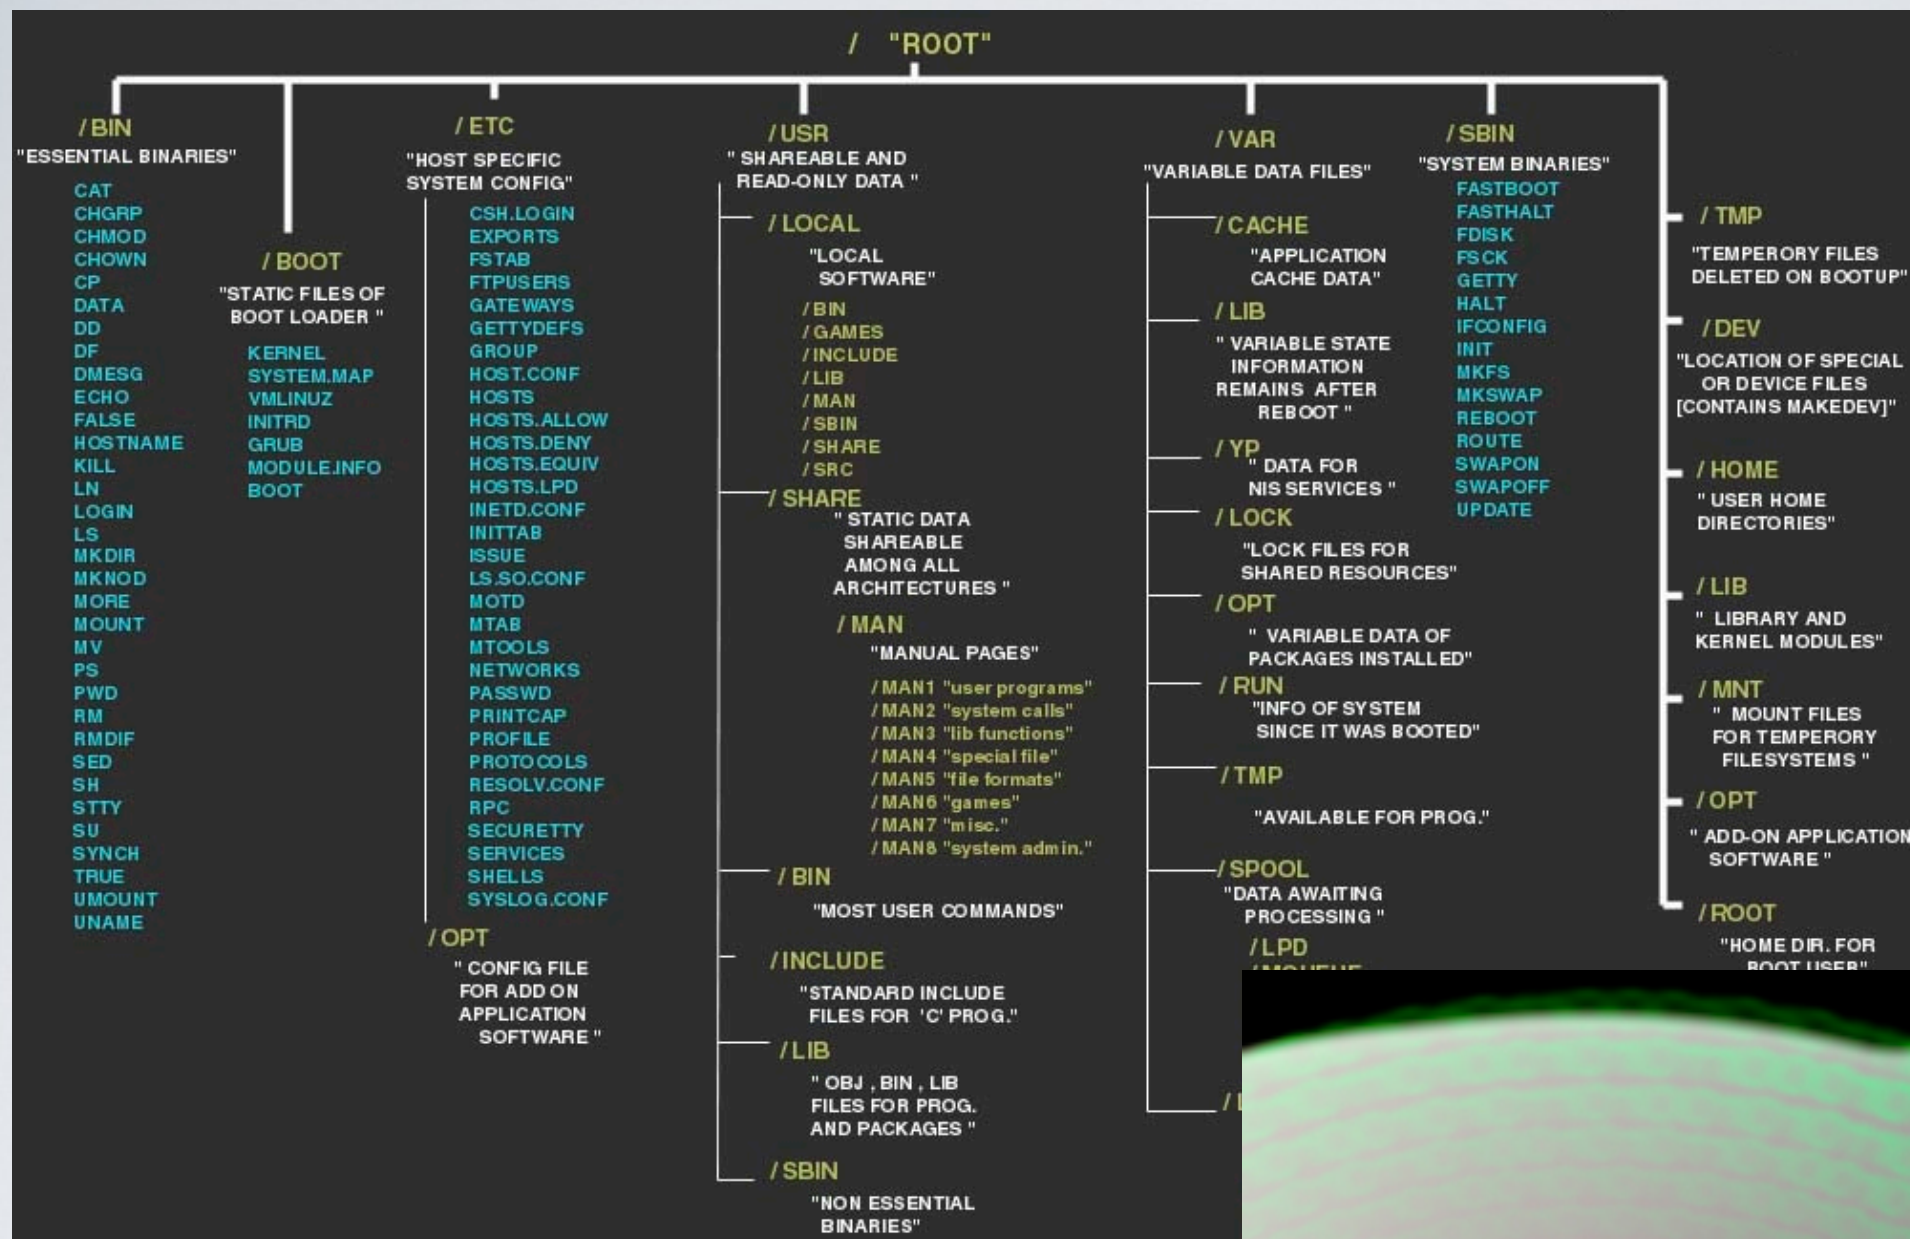
physical memory



hard drive



The need for a **file system**



Files and Directories

versus

Reality



So, what is an **operating system**?

Operating System

- ➔ In a nutshell, an OS manages hardware and runs programs
 - creates and manages processes
 - manages access to the memory (including RAM and I/O)
 - manages files and directories of the filesystem on disk(s)
 - enforces protection mechanisms for reliability and security
 - enables inter-process communication

Acknowledgments

Some of the course materials and projects are from

- Ryan Huang - teaching CS 318 at *John Hopkins University*
- David Mazière - teaching CS 140 at *Stanford*