

Threads and Synchronization

Thierry Sans

(recap) Processes

- A process is defined by its Process Control Block (PCB) that defines:
 - The execution state (running, waiting, ready)
 - The address space with code and data
 - The execution context (PC, SP, registers)
 - The resources (open files)
 - and so others ...

The cost of multi-processing

```
while (1) {  
    int sock = accept();  
    if ((child_pid = fork()) == 0) {  
        // Handle client request  
    } else {  
        // Close socket  
    }  
}
```

Recall our Web Server example
we need to fork a child process for each request

- Create a new PCB
- Copy the address space and the resources
- Have the OS execute this child process
(context switching)
- Use signals and pipes if the child wants to send information back to the parent process

A good but costly abstraction

- ✓ Good to avoid processes interfering with each other but ...
 - Creating a process is costly (space and time)
 - Context switching is costly (time)
 - Inter-process communication is costly (time)

The need for cooperation

An application could have some sort of cooperating processes

- that all share the same code and data (address space)
- that all share the same resources (files, sockets, etc.)
- that all share the same privileges

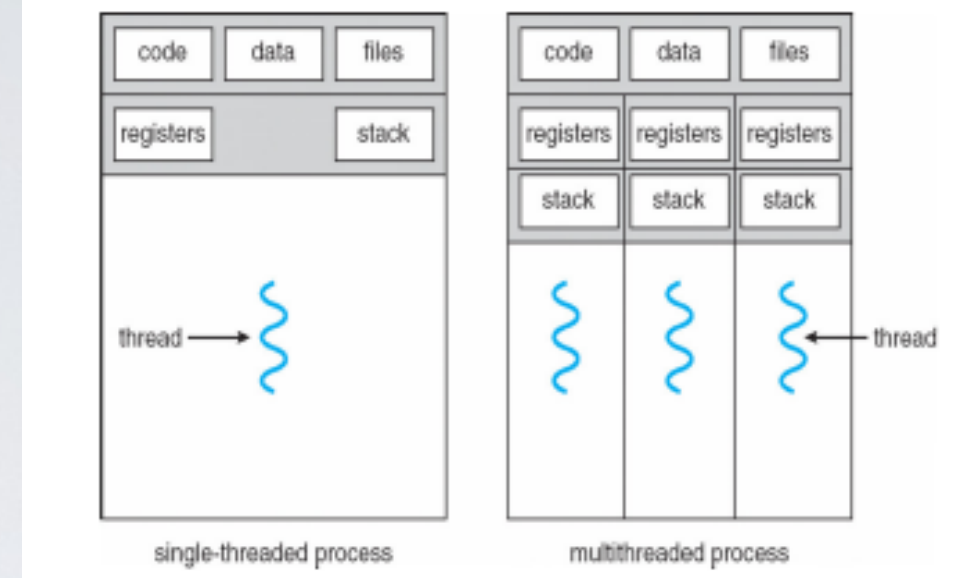
... while having different execution context (PC, SP, registers)

Rethinking process

Why not separate the process concept from its execution state?

- **Process** : address space, privileges, resources, etc
- **Thread** : PC, SP, registers

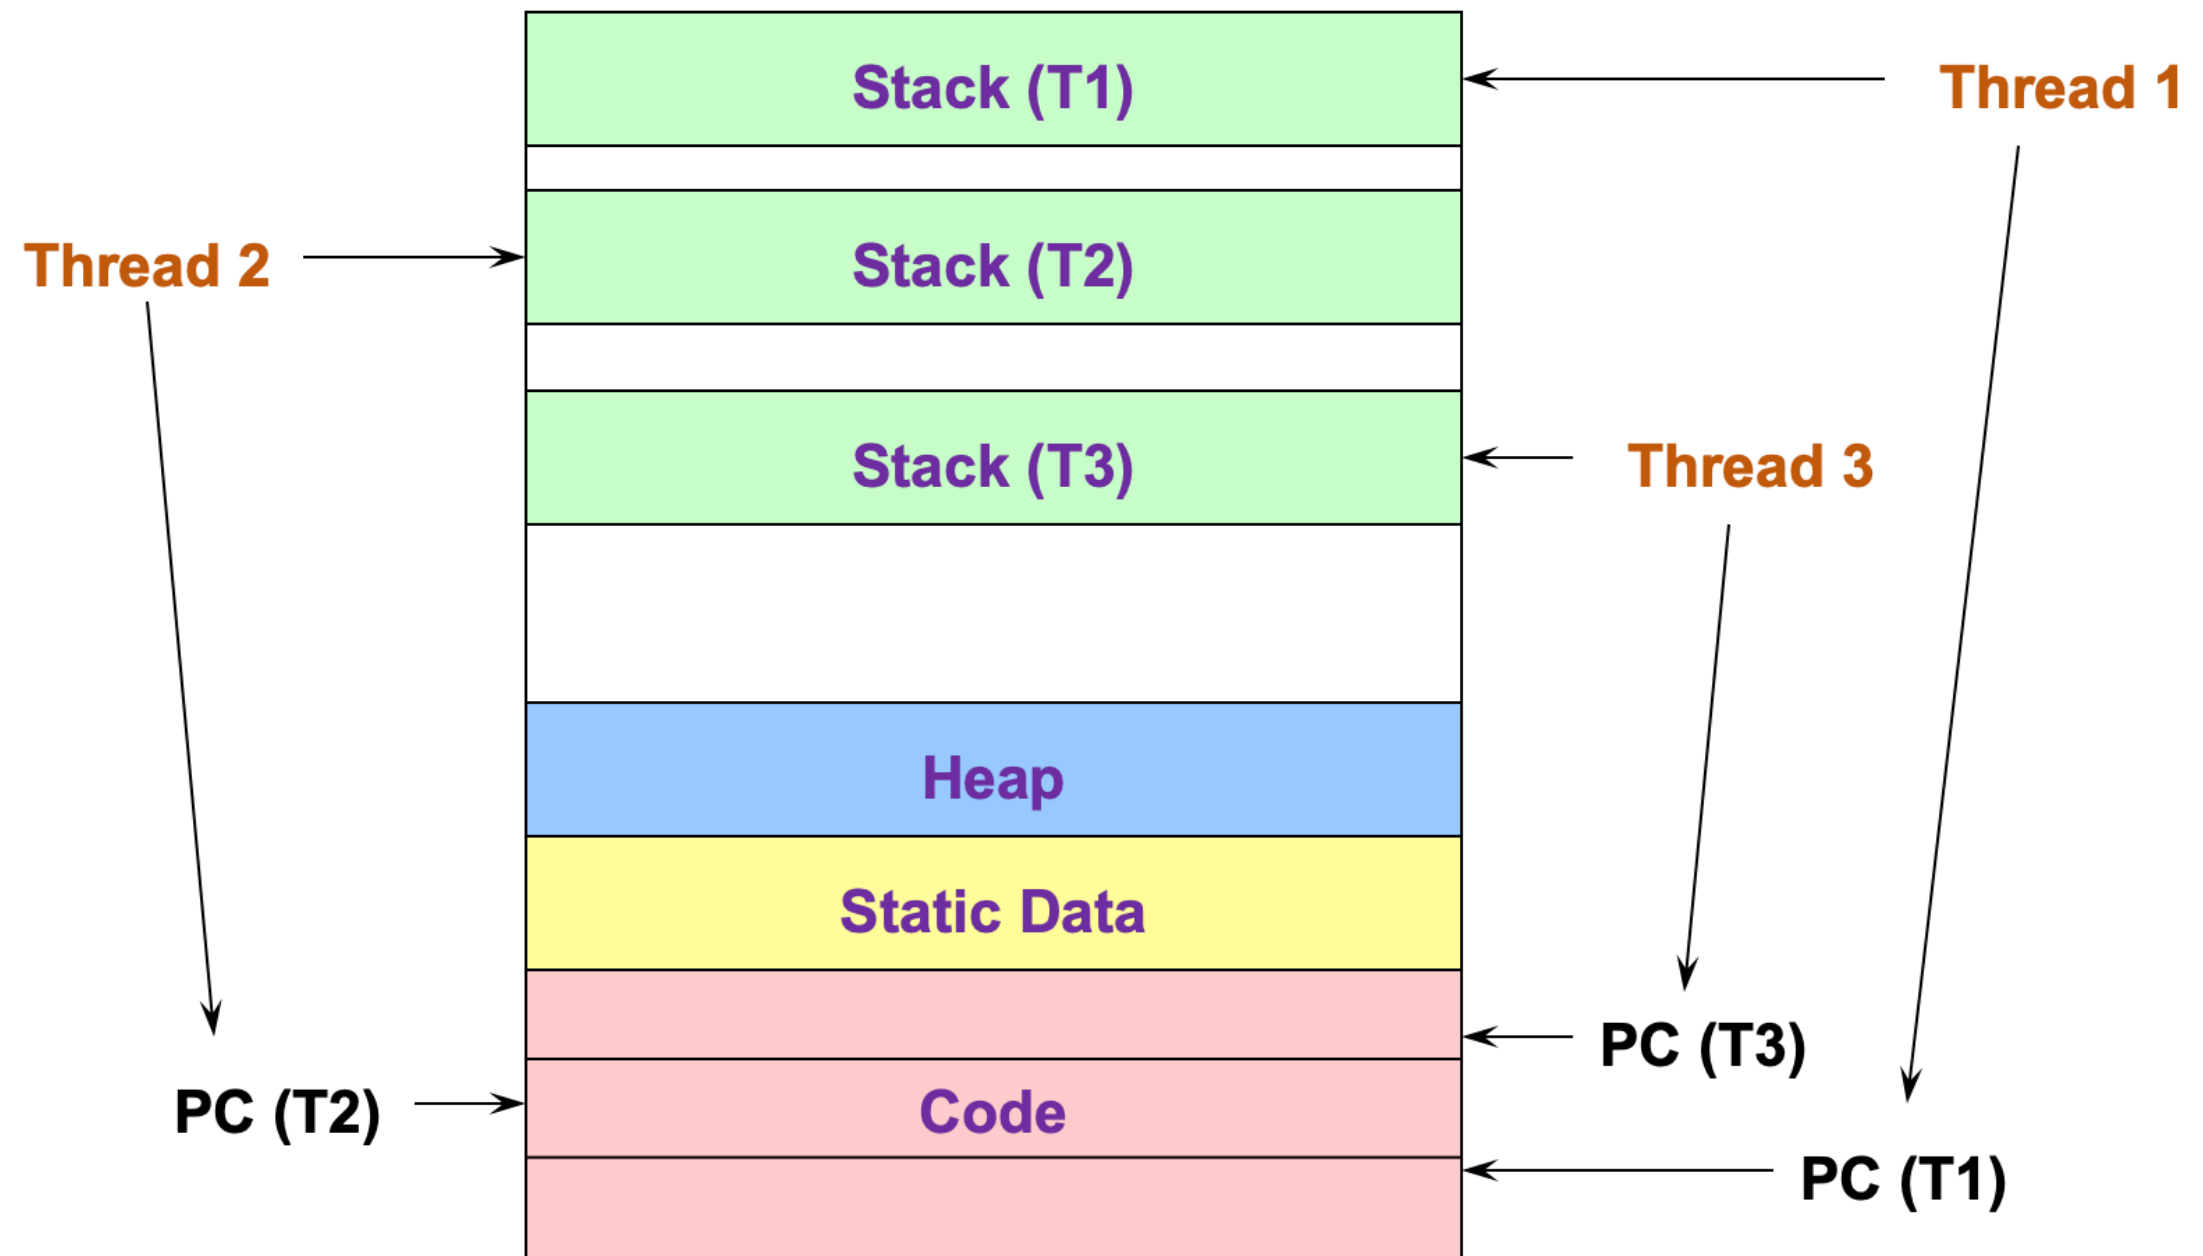
Threads



Modern OSes separate the concepts of processes and threads

- The thread defines a sequential execution stream within a process (PC, SP, registers)
 - The process defines the address space and general process attributes (everything but threads of execution)
- ➡ Most popular abstraction for concurrency
threads become the unit of scheduling
while processes are now the containers in which threads execute
- ✓ A thread is bound to a single process but a process can have multiple threads

Threads within a process



Our web server becomes

```
web_server() {  
    while (1) {  
        int sock = accept();  
        thread_fork(handle_request, sock);  
    }  
}  
  
handle_request(int sock) {  
    Process request  
    close(sock);  
}
```

Benefits

- **Responsiveness**

an application can continue running while it waits for some events in the background

- **Resource sharing**

threads can collaborate by reading and writing the same data in memory (instead of asking the OS to pass data around)

- **Economy of time and space**

no need to create a new PCB and switch the entire context (only the registers and the stack)

- **Scalability in multi-processor architecture**


the same application can run on multiple cores

Multithreading models

- One-to-one model
Kernel-level threads (a.k.a native threads)
- Many-to-one model
User-level threads (a.k.a green threads)
- Many-to-many model
Hybrid threads (a.k.a n:m threading)

One-to-one model

Kernel-level threads (a.k.a native threads)



The kernel manage and schedule threads

- e.g Windows threads
- e.g POSIX pthreads `PTHREAD_SCOPE_SYSTEM`
- e.g (new) Solaris lightweight processes (LWP)

➡ All thread operations are managed by the kernel

✓ good for scheduling

✓ bad for speed

POSIX Thread API

- Create a new thread, run fn with arg

```
tid thread_create (void (*fn) (void *), void *);
```

- Allocate Thread Control Block (TCB)
- Allocate stack
- Put func, args on stack
- Put thread on ready list

- Destroy current thread


```
void thread_exit ();
```

- Wait for thread thread to exit

```
void thread_join (tid thread);
```


Many-to-one model

User-level threads (a.k.a green threads)



One kernel thread per process
thread management and scheduling is delegated to a library

- e.g pthreads `PTHREAD_SCOPE_PROCESS`
- e.g Java threads


➡ The kernel is not involved

- ✓ Very lightweight and fast
- ✓ All threads can be blocked if one of them is waiting on an event
- ✓ Cannot be scheduled on multiple cores

Many-to-many model

Hybrid threads

(a.k.a n:m threading)



User threads implemented on kernel threads

- e.g (old) Solaris

➡ Multiple kernel-level threads per process

Now threads can collaborate but ...

What are these two threads printing?

Ping thread

```
while(1) {  
    printf("ping\n");  
};
```

Pong thread

```
while(1) {  
    printf("pong\n");  
};
```

Too much milk

	Alice	Bob
12:30	Look in the fridge. Out of milk.	
12:35	Leave for store	
12:40	Arrive at store	Look in the fridge. Out of milk.
12:45	Buy milk	Leave for store
12:50	Arrive home, put milk away	Arrive at store
12:55		Buy milk
1:00		Arrive home, put milk away ... oh no!

Beyond milk

X is a global variable initialized to 0

thread 1

```
void foo() {  
    x++;  
};
```

thread 2

```
void bar() {  
    x--;  
};
```

What is the value of x after thread 1 and 2?

CPU instruction level

Incrementing (or decrementing) x is **not** an atomic operation

thread 1 (foo function)

```
LOAD X  
INCR  
STORE X
```

thread 2 (bar function)

```
LOAD X  
DECR  
STORE X
```

Non-deterministic execution

Execution scenario #1

LOAD X
INCR
STORE X
LOAD X
DECR
STORE X

➡ X is equal to 0

Execution scenario #2

LOAD X
LOAD X
INCR
DECR
STORE X
STORE X

➡ X is equal to -1

Execution scenario #3

LOAD X
LOAD X
INCR
DECR
STORE X
STORE X

➡ X is equal to 1

... and many other possible scenarios with the outcome of x being equal to either 0, -1 or 1

Race-condition problem

The system behaviours depends on the sequence or timing of events that is non-deterministic

- Not desirable in most cases (hard to catch bug)

Mutual exclusion

We want to use **mutual exclusion** to synchronize access to shared resources

Code that uses mutual exclusion to synchronize its execution is called a **critical section**

- Only one thread at a time can execute in the critical section
- All other threads are forced to wait on entry
- When a thread leaves a critical section, another can enter

A classic example

Identify a critical section that lead to a race condition

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

critical
section

Requirements

1. **Mutual exclusion**

If one thread is in the critical section, then no other is

➔ Mutual exclusion ensures **safety property** (nothing bad happen)

2. **Progress**

If some thread T is not in the critical section, then T cannot prevent some other thread S from entering the critical section. A thread in the critical section will eventually leave it.

3. **Bounded waiting** (no starvation)

If some thread T is waiting on the critical section, then T will eventually enter the critical section

➔ Progress and bounded waiting ensures the **liveness property** (something good happen)

4. **Performance**

The overhead of entering and exiting the critical section is small with respect to the work being done within it

Mechanisms for building critical sections

- **Locks**

Primitive, minimal semantics, used to build others

- **Semaphores**

Basic, easy to get the hang of, but hard to program with

- **Monitors**

High-level, requires language support, operations implicit

Locks

- A lock is an object in memory providing two operations
 - `acquire()`
wait until lock is free, then take it to enter a C.S
 - `release()`
release lock to leave a C.S, waking up anyone waiting for it
- ➡ Threads pair calls to `acquire` and `release`
We say that the thread *holds the lock* in between `acquire/release`
- ✓ Locks can spin (a spinlock) or block (a mutex)

Using locks

code

```
Withdraw(acct, amt) {  
    acquire(lock);  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    release(lock);  
    return balance;  
}
```

execution of thread 1 and thread 2

```
acquire(lock);  
balance = get_balance(account);  
balance = balance - amount;
```

```
acquire(lock);
```

```
put_balance(acct, balance);  
release(lock);
```

```
balance = get_balance(acct);  
balance = balance - amt;  
put_balance(acct, balance);  
release(lock);
```


Implementing a spin lock (naive but wrong attempt)

```
struct lock {  
    int held = 0;  
}  
  
void acquire (lock) {  
    while (lock->held);  
    lock->held = 1;  
}  
  
void release (lock) {  
    lock->held = 0;  
}
```

What is the context switch happens in between?
➡ We have a race condition

The hardware to the rescue

- test-and-set (TAS x86 CPU instruction)
atomically writes to the memory location
and returns its old value in a **single indivisible step**
- ➔ the caller is responsible for testing if the operation has succeeded or not

```
bool test_and_set(bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old;  
}
```

This is pseudo-code!
The hardware execute this atomically

Implementing a spin lock

```
struct lock {  
    int held = 0;  
}  
  
void acquire (lock) {  
    while test-and-set(&lock->held);  
}  
  
void release (lock) {  
    lock->held = 0;  
}
```

Busy wait (a.k.a spin)

- Waste of CPU time
- Unfair access to lock

Implementing a lock by disabling interrupt

```
struct lock {  
}  
  
void acquire (lock) {  
    disable_interrupts();  
}  
  
void release (lock) {  
    enable_interrupts();  
}
```

- ➔ Disabling interrupts blocks notification of external events that could trigger a context switch
- Can miss or delay important events

Our lock implementations so far

- **Goal** : Use mutual exclusion to protect critical sections of code that access shared resources
- ✓ **Method** : Use locks (spinlocks or disable interrupts)
- ◉ **Problem** : Critical sections (CS) can be long
 - ◉ spinlocks waste CPU time
and do not provide fair access to lock
 - ◉ disabling interrupts can delay important events

Blocking Lock Implementation

```
struct lock {
    int held = 0;
    queue Q;
}

void acquire (lock) {
    disable_interrupts();
    while (lock->held) {
        enqueue(lock->Q, current_thread);
        thread_block(current_thread);
    }
    lock->held = 1;
    enable_interrupts();
}

void release (lock) {
    disable_interrupts();
    if (!isEmpty(lock->Q)) {
        thread_unblock(dequeue(lock->Q));
    }
    lock->held = 0;
    enable_interrupts();
}
```

Semaphores

An abstract data type to provide mutual exclusion described by *Dijkstra* in the "*THE multiprogramming system*" in 1968

➔ Semaphores are “integers” that support two operations:

- Semaphore::P() decrement, block until semaphore is open
a.k.a wait(), or sem_wait(), or sema_down()
- Semaphore::V() increment, allow another thread to enter
a.k.a signal(), or sem_post(), or sema_up()

✓ Semaphore safety property
the semaphore value is always greater than or equal to 0

Blocking mechanism

Associated with each semaphore is a queue of waiting threads

➡ When $P()$ is called by a thread:

- If semaphore is open, thread continue
- If semaphore is closed, thread blocks on queue

➡ Then $V()$ opens the semaphore

- If a thread is waiting on the queue, the thread is unblocked
- If no threads are waiting on the queue, the signal is remembered for the next thread

Using semaphores

code

```
Withdraw(acct, amt) {  
    P(s);  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    V(s);  
    return balance;  
}
```

execution of
thread 1, thread 2 and thread 3

```
P(s);  
balance = get_balance(account);  
balance = balance - amount;
```

```
P(s);
```

```
P(s);
```

```
put_balance(acct, balance);  
V(s);
```

```
...  
V(s);
```

```
...  
V(s);
```


Semaphore Implementation

```
struct semaphore {
    int value;
    queue Q;
}

void init(sema, value) {
    sema->value = value;
}

void P (sema) {
    disable_interrupts();
    while (sema->value == 0) {
        enqueue(sema->Q, current_thread);
        thread_block(current_thread);
    }
    sema->value--;
    enable_interrupts();
}

void V (sema) {
    disable_interrupts();
    if (!isEmpty(sema->Q)) {
        thread_unblock(dequeue(sema->Q));
    }
    sema->value++;
    enable_interrupts();
}
```

Two types of semaphores provided by OSes

Binary semaphore (a.k.a *mutex*)

controls access to a single resource (mutual exclusion)

➔ behave exactly like a lock

General semaphore (a.k.a counting semaphore)

controls access to a finite number of resources n available

- The semaphore is initialized with a positive integer n
- $P()$ blocks when this number is equal to 0

➔ at most n threads can be in the critical section

Acknowledgments

Some of the course materials and projects are from

- Ryan Huang - teaching CS 318 at *John Hopkins University*
- David Mazière - teaching CS 140 at *Stanford*