

Copy on Write

- OSes spend a lot of time copying data
 - System call arguments between user/kernel space
 - Entire address spaces to implement `fork()`
- ➔ Use Copy on Write (CoW) to defer large copies as long as possible, hoping to avoid them altogether
 - Create shared mappings of parent pages in child virtual address space (instead of copying pages)
 - Shared pages are protected as read-only in parent and child
Any write operation generates a protection fault, trap to OS, copy page, change page mapping in client page table, restart write instruction

Example - Fork step

