

# I/O and Disks

Thierry Sans

I/O

# I/O management



- I/O devices vary greatly and new types of I/O devices appear frequently
  - Various methods to control them and to manage their performances
- ➔ Ports, buses, device controllers connect to various devices

# I/O Device Interfaces

**Port** - connection point for device (e.g. serial port)

**Bus** - daisy chain or shared direct access

e.g. Peripheral Component Interconnect Bus (PCI)

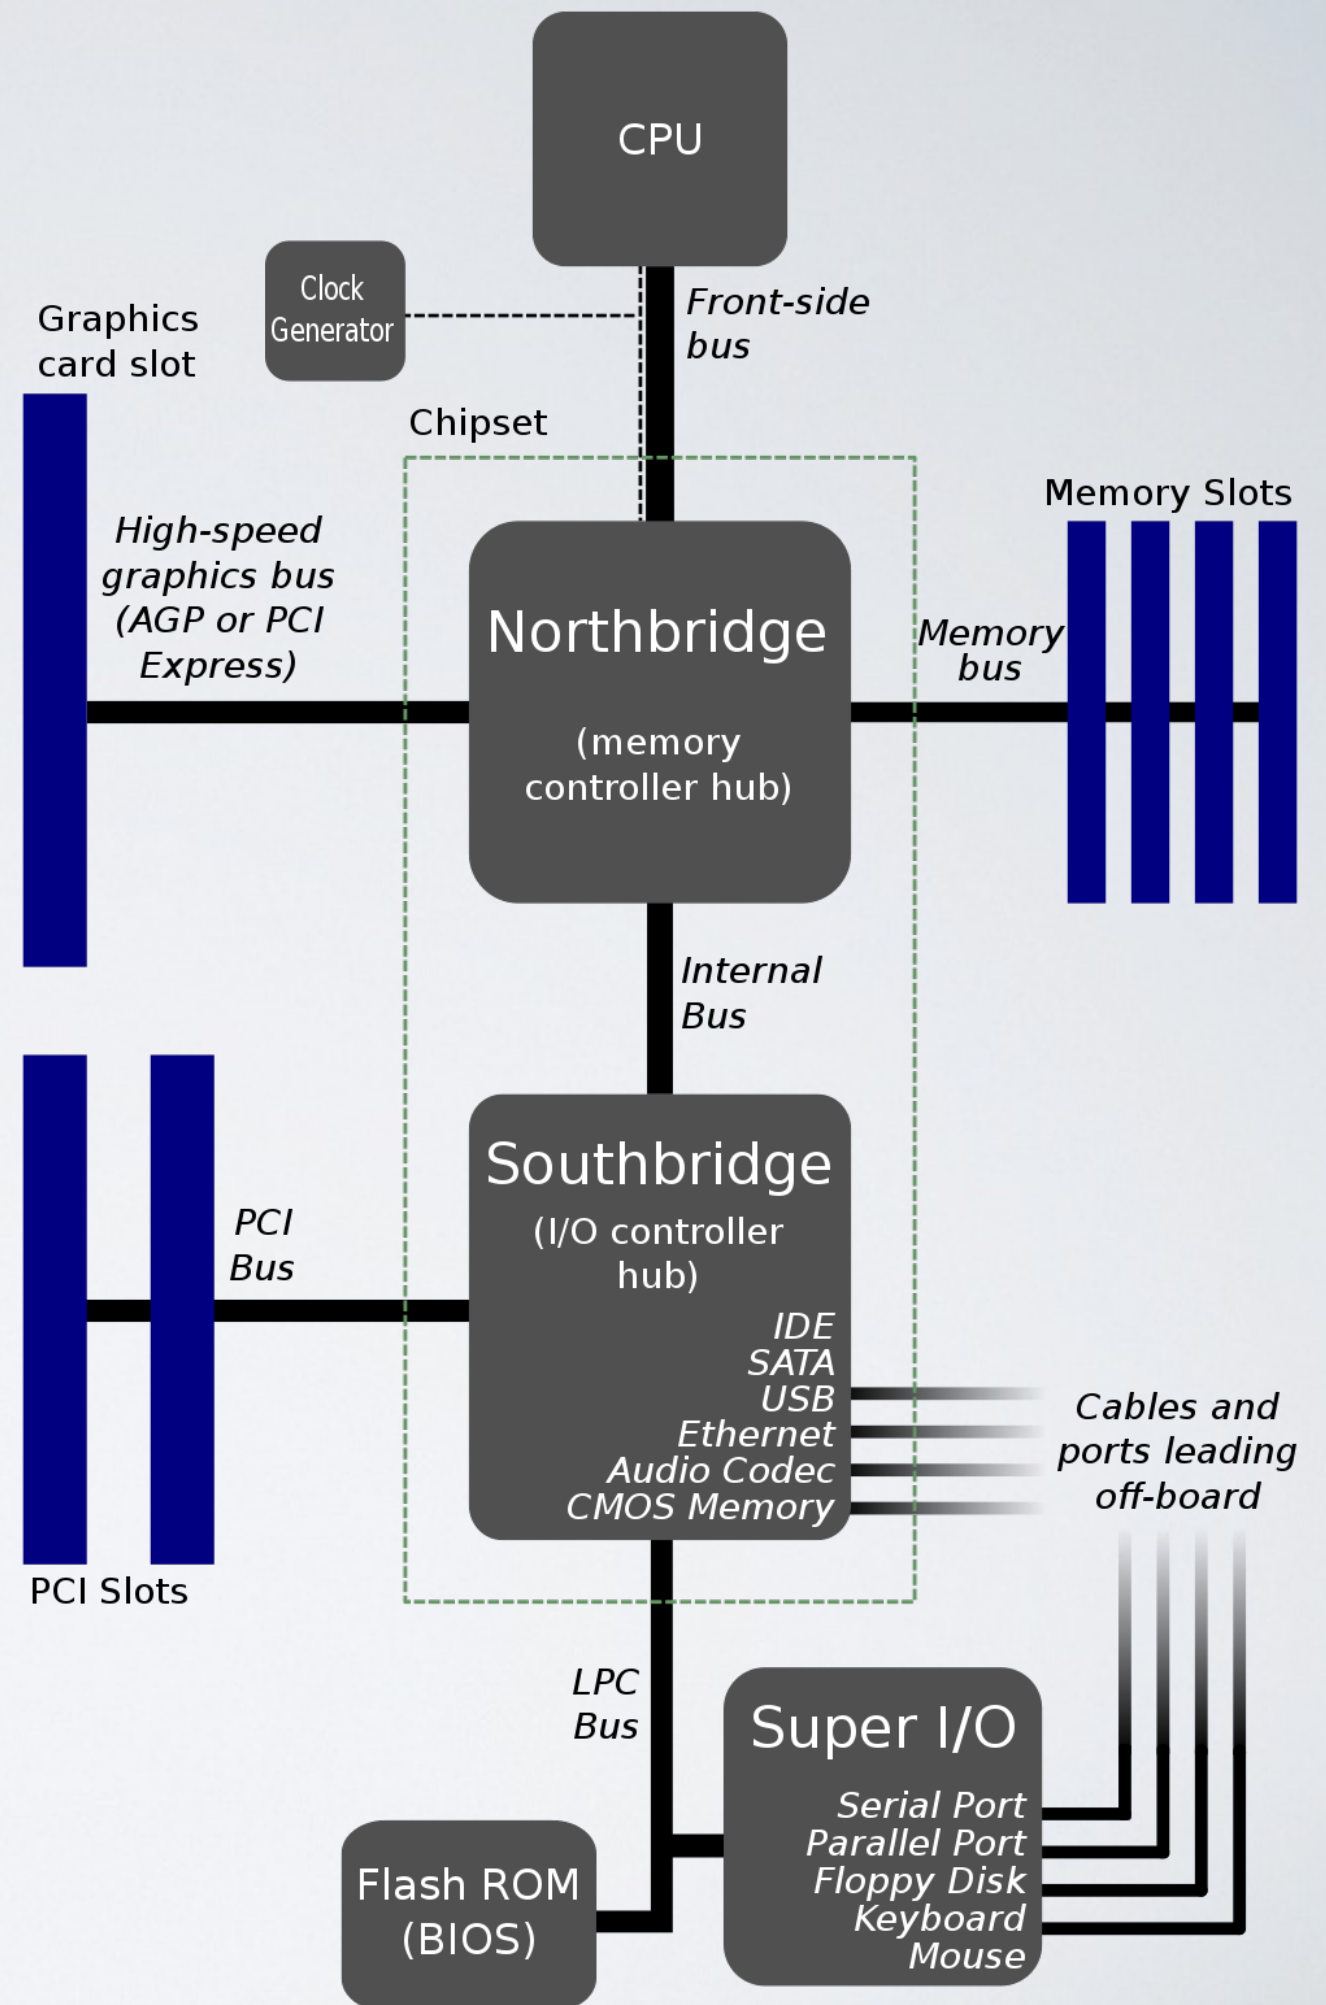
e.g. Universal Serial Bus (USB)

**Controller** (host adapter) - electronics that operate port, bus, device (e.g. Northbridge, Southbridge, graphics controller, DMA, NIC, ...)

- Can be integrated or separated (host adapter)
- Contains processor, microcode, private memory, bus controller, etc



# I/O architecture



# How the OS communicates with the device?

- ➔ Each device has three types of registers  
and the OS controls the device by reading or writing these registers

**status** register

See the current status of the device

**command** register

Tell the device to perform a certain task

**data** register

Pass data to the device, or get data from the device

# Two ways to read/write those registers

## **I/O ports**

`in` and `out` instructions on x86 to read and write devices registers

## **Memory-mapped I/O**

Device registers are available as if they were memory locations and the OS can `load` (to read) or `store` (to write) to the device



# I/O Ports on PC

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



# Reading/Writing to I/O ports

Pintos threads/io.h

```
static inline uint8_t inb (uint16_t port)
{
    uint8_t data;
    asm volatile ("inb %w1, %b0" : "=a" (data) : "Nd" (port));
    return data;
}

static inline void outb (uint16_t port, uint8_t data)
{
    asm volatile ("outb %b0, %w1" : : "a" (data), "Nd" (port));
}
```

# Device driver

```
while (STATUS == BUSY)  
    ; //wait until device is not busy
```

```
write data to data register
```

```
write command to command register
```

```
    Doing so starts the device and executes the command
```

```
while (STATUS == BUSY)  
    ; //wait until device is done with your request
```

# Example : parallel port (LPT1)

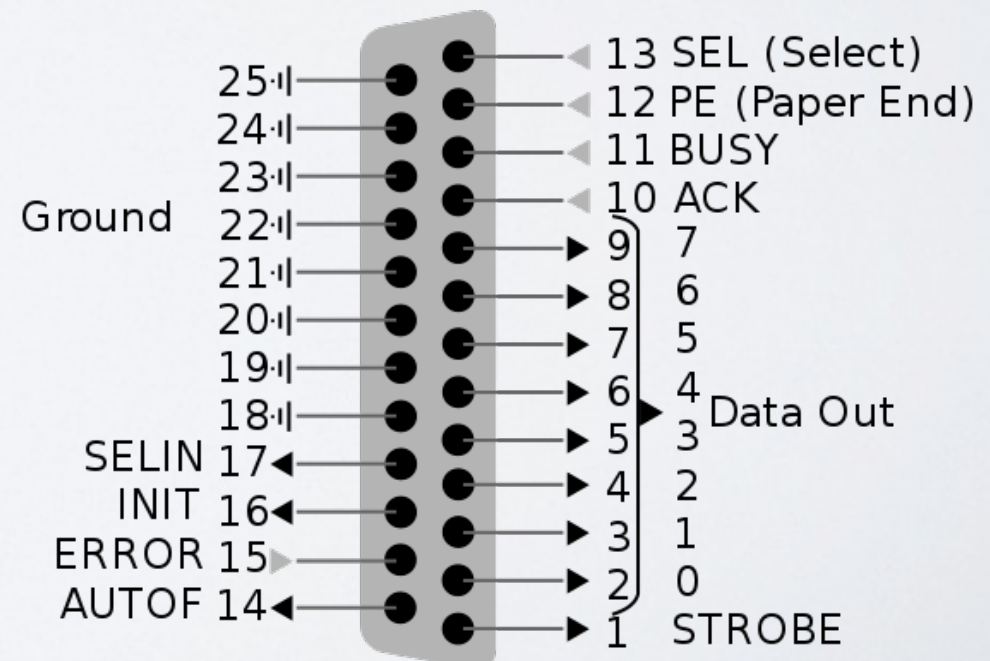
- Three controllers

$D_7$	$D_6$	$D_5$	$D_4$	$D_3$	$D_2$	$D_1$	$D_0$
read/write data register (port 0x378)							

$\overline{BSY}$	$\overline{ACK}$	PAP	OFON	$\overline{ERR}$	-	-	-
read-only status register (port 0x379)							

-	-	-	IRQ	DSL	$\overline{INI}$	ALF	STR
read/write control register (port 0x37a)							

- Every bits (except IRQ) corresponds to a pin on 25-pin connector



# Parallel Port Driver

```
void
sendbyte(uint8_t byte)
{
    /* Wait until  $\overline{\text{BSY}}$  bit is 1. */
    while ((inb (0x379) & 0x80) == 0)
        delay ();

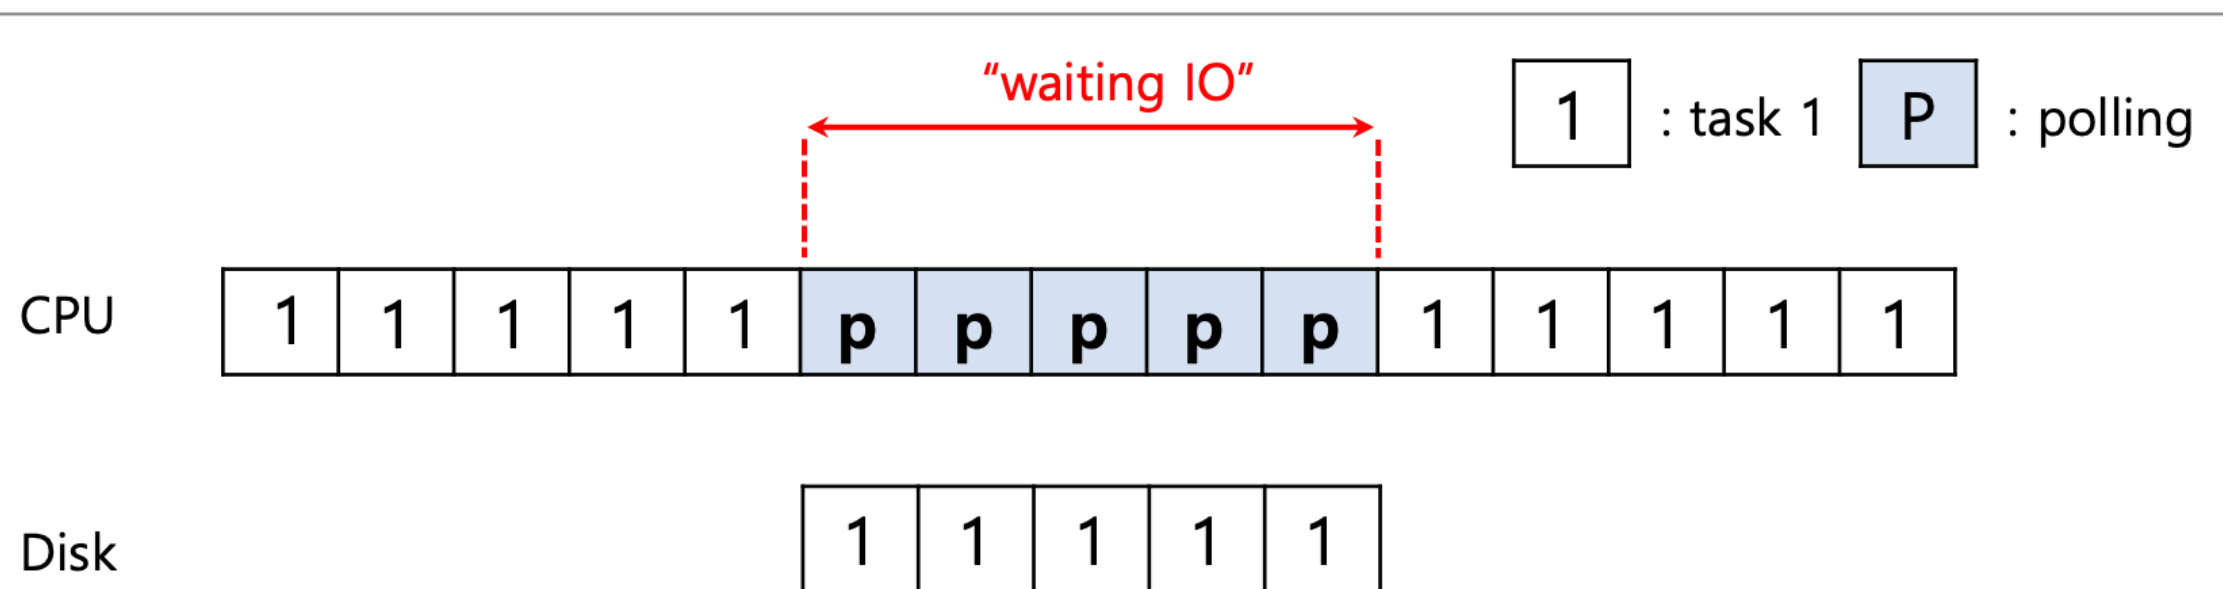
    /* Put the byte we wish to send on pins D7-0. */
    outb (0x378, byte);

    /* Pulse STR (strobe) line to inform the printer
       * that a byte is available */
    uint8_t ctrlval = inb (0x37a);
    outb (0x37a, ctrlval | 0x01);
    delay ();
    outb (0x37a, ctrlval);
}
```



# Polling


- ➔ OS waits until the device is ready by repeatedly reading the status register
- ✓ Simple and working
- ⦿ Wastes CPU time just waiting for the device



**Diagram of CPU utilization by polling**

# Interrupts

1. Put the I/O request process to sleep and switch context
  2. When the device is finished, send an interrupt to wake the process waiting for the I/O
- ✓ CPU is properly utilized



**Diagram of CPU utilization by interrupt**

# Polling vs Interrupts

## ➔ **Interrupts is not always the best solution**

If, device performs very quickly, interrupt will slow down the system

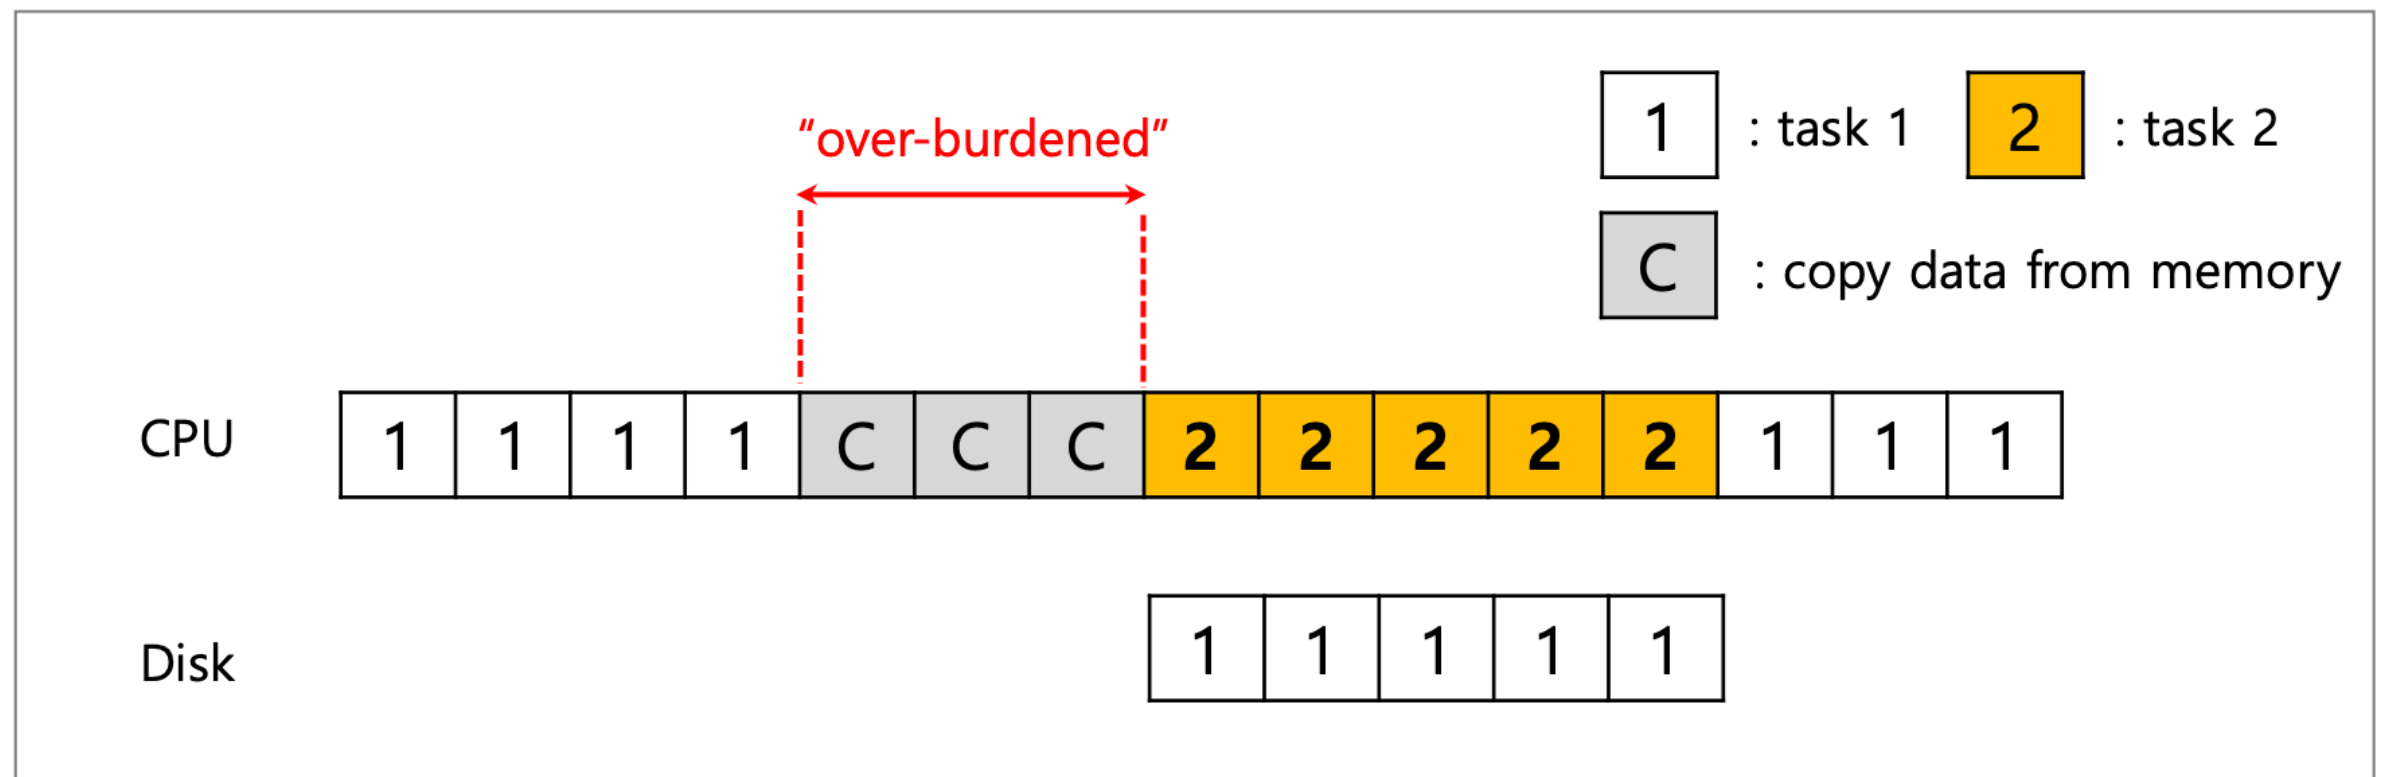
E.g., high network packet arrival rate

- Packets can arrive faster than OS can process them
- Interrupts are very expensive (context switch)
- Interrupt handlers have high priority
- In worst case, can spend 100% of time in interrupt handler and never make any progress a.k.a receive livelock

✓ Best - adaptive switching between interrupts and polling

# One More Problem : Data Copying

- CPU wastes a lot of time in copying a large chunk of data from memory to the device

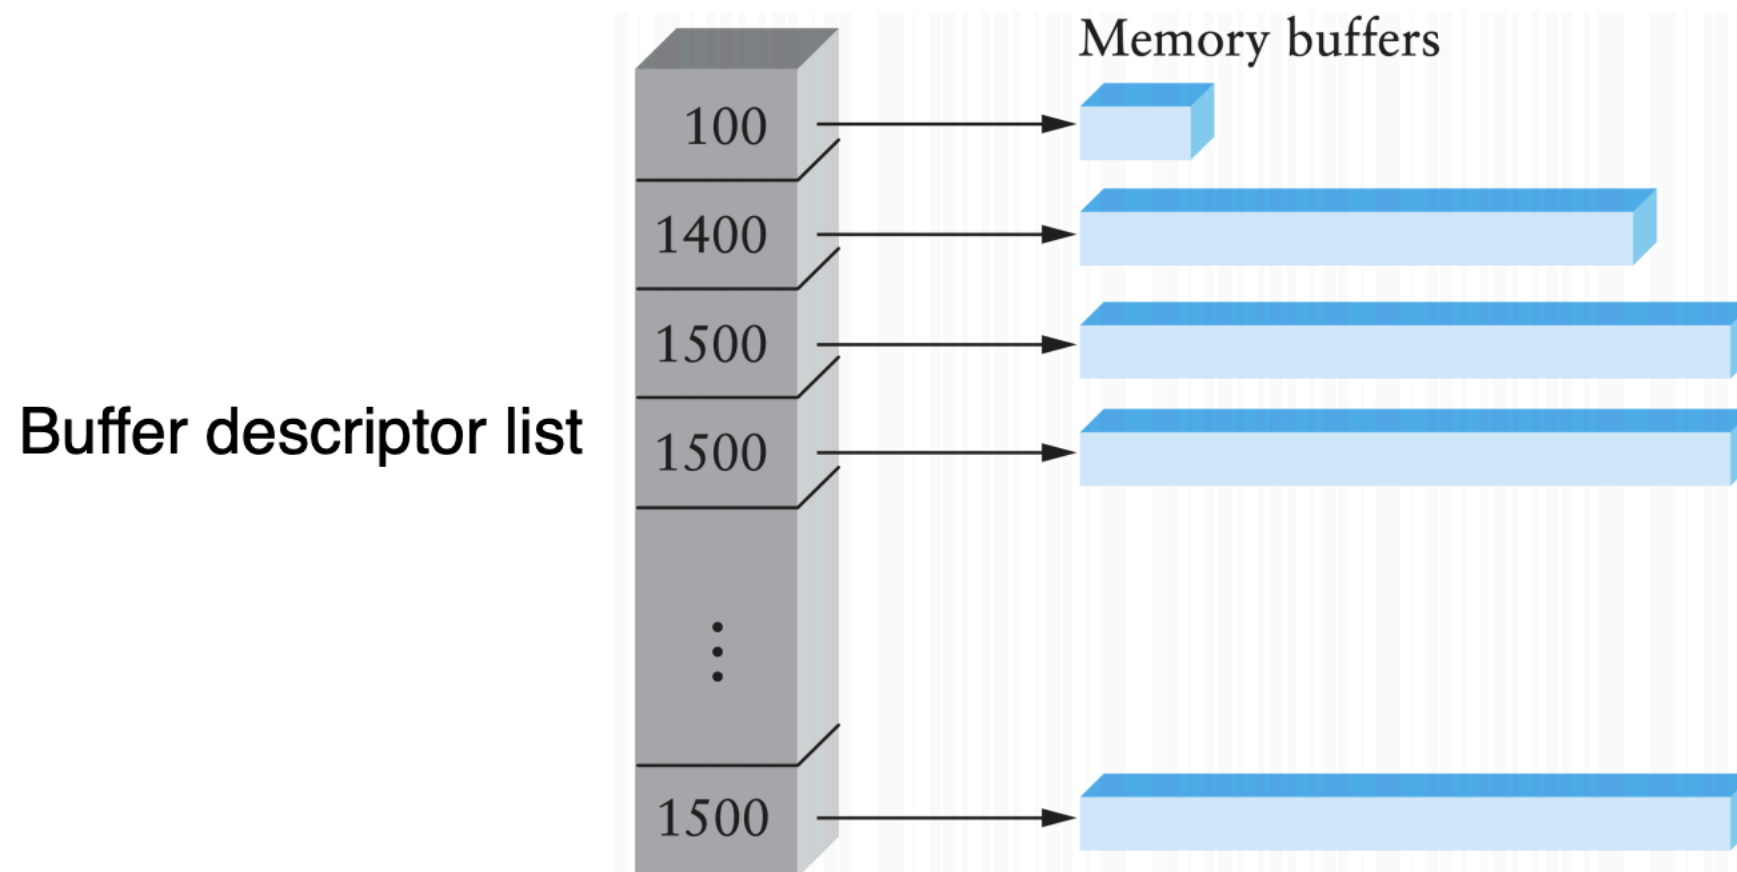


**Diagram of CPU utilization**



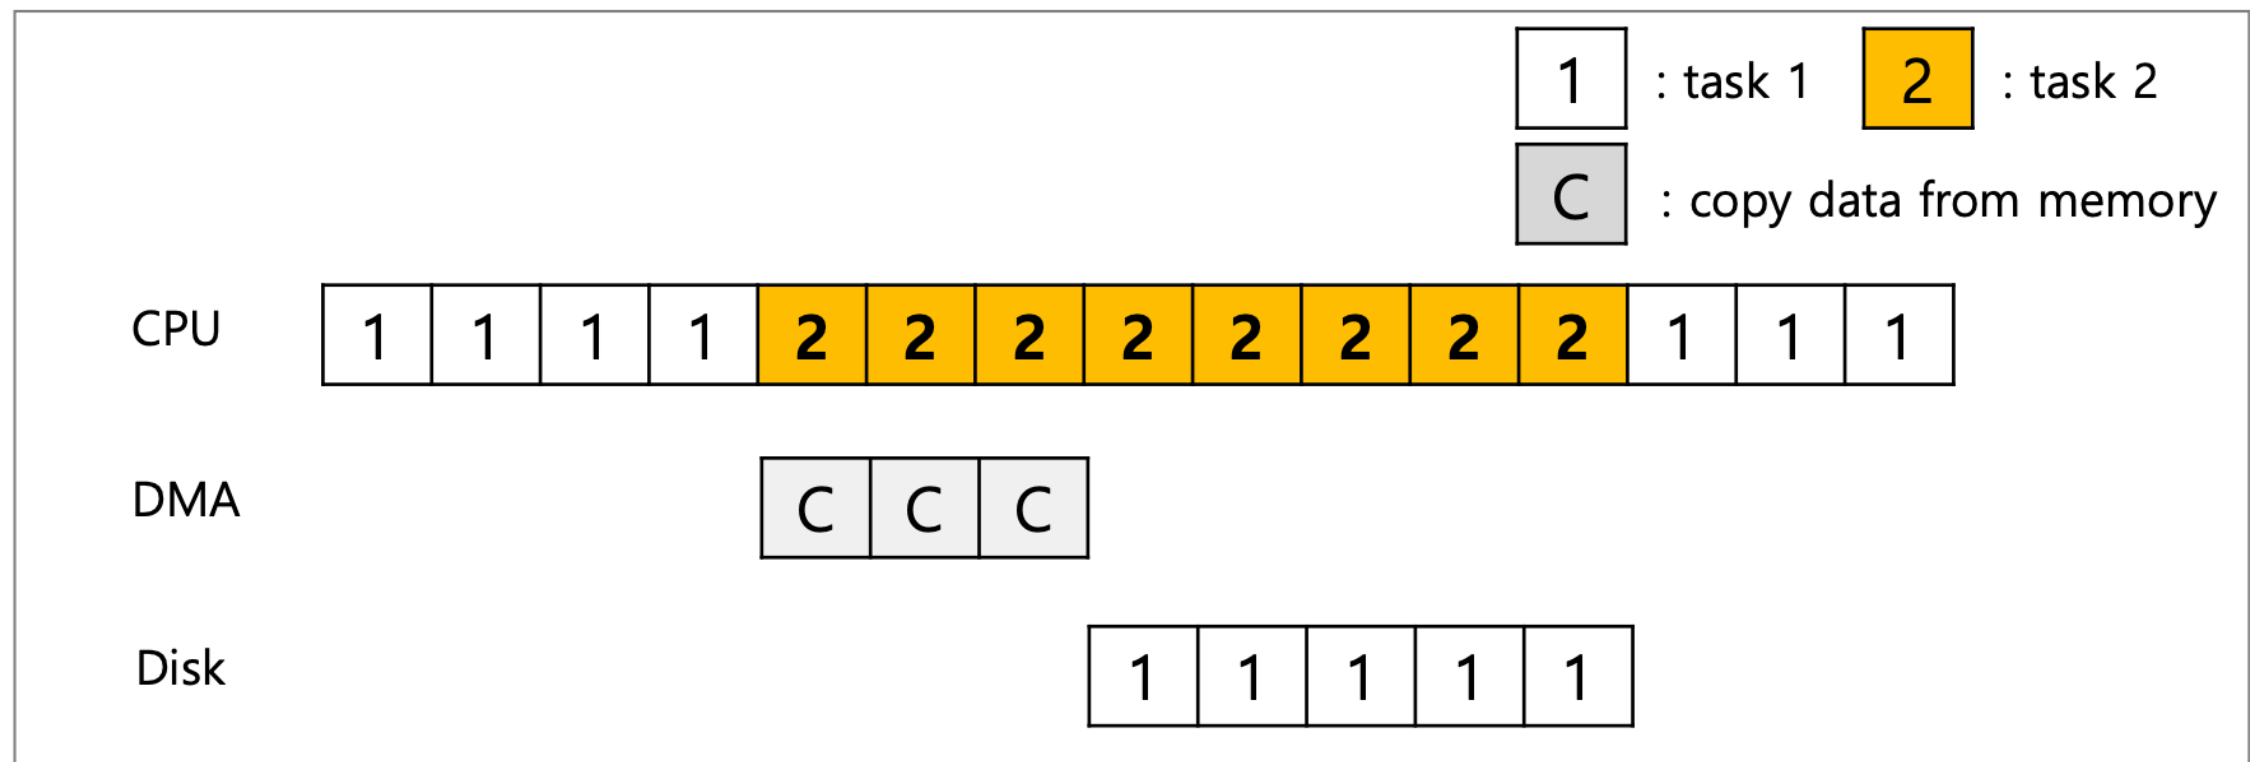
# DMA (Direct Memory Access)

- ➔ Only use CPU to transfer control requests, not data, by passing buffer locations in memory
  - Device reads list and accesses buffers through DMA
  - Descriptions sometimes allow for scatter/gather I/O



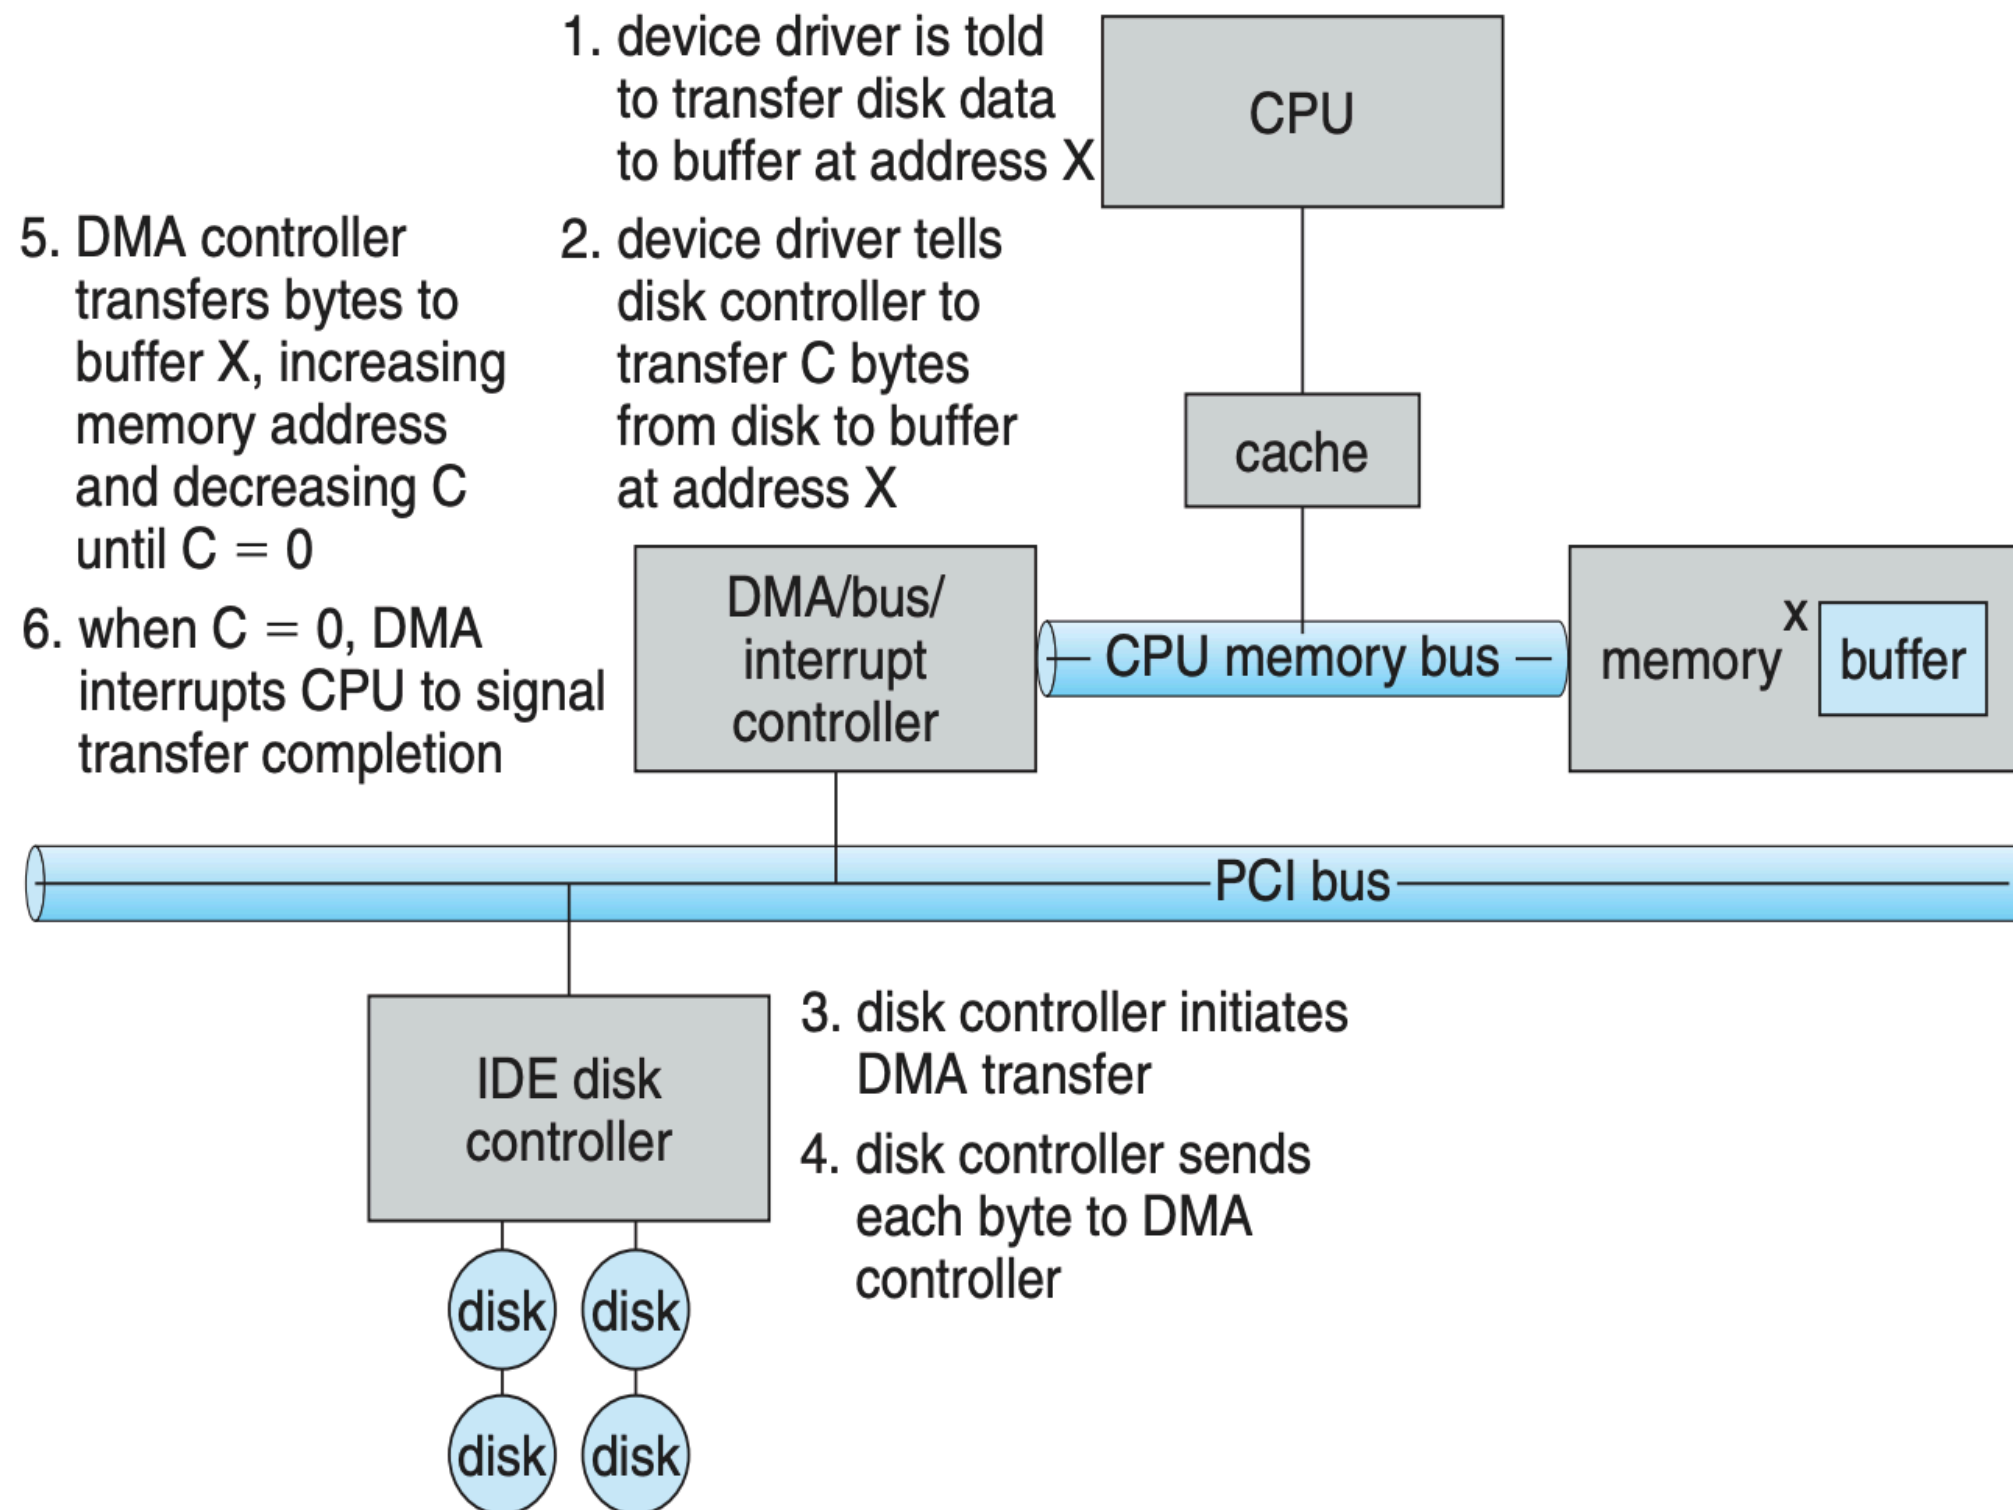
# DMA (Direct Memory Access)

1. OS writes DMA command block into memory
2. DMA bypasses CPU to transfer data directly between I/O device and memory
3. When completed, DMA raises an interrupt



**Diagram of CPU utilization by DMA**

# Example : IDE disk read with DMA



# I/O instruction using DMA

Pintos threads/io.h

```
static inline void insw (uint16_t port, void *addr, size_t cnt)
{
    asm volatile ("rep insw" : "+D" (addr), "+c" (cnt)
                  : "d" (port) : "memory");
}
```




# Example : IDE Disk Driver

```
void IDE_ReadSector(int disk, int off, void *buf)
{
    outb(0x1F6, disk == 0 ? 0xE0 : 0xF0); // Select Drive
    IDEWait();
    outb(0x1F2, 1); // Read length (1 sector = 512 B)
    outb(0x1F3, off); // LBA low
    outb(0x1F4, off >> 8); // LBA mid
    outb(0x1F5, off >> 16); // LBA high
    outb(0x1F7, 0x20); // Read command
    insw(0x1F0, buf, 256); // Read 256 words
}

void IDEWait()
{
    // Discard status 4 times
    inb(0x1F7); inb(0x1F7);
    inb(0x1F7); inb(0x1F7);
    // Wait for status BUSY flag to clear
    while ((inb(0x1F7) & 0x80) != 0)
        ;
}
```

# Example : Network Interface Card



- Link interface talks to wire/fiber/antenna
- FIFOs on card provide small amount of buffering
- Bus interface logic uses DMA to move packets to and from buffers in main memory

# Variety is a challenge

## ● Problem : there are many devices and each has its own protocol

- Some devices are accessed by I/O ports or memory mapping or both
- Some devices can interact by polling or interrupt or both
- Some device can transfer data by programmed I/O or DMA or both

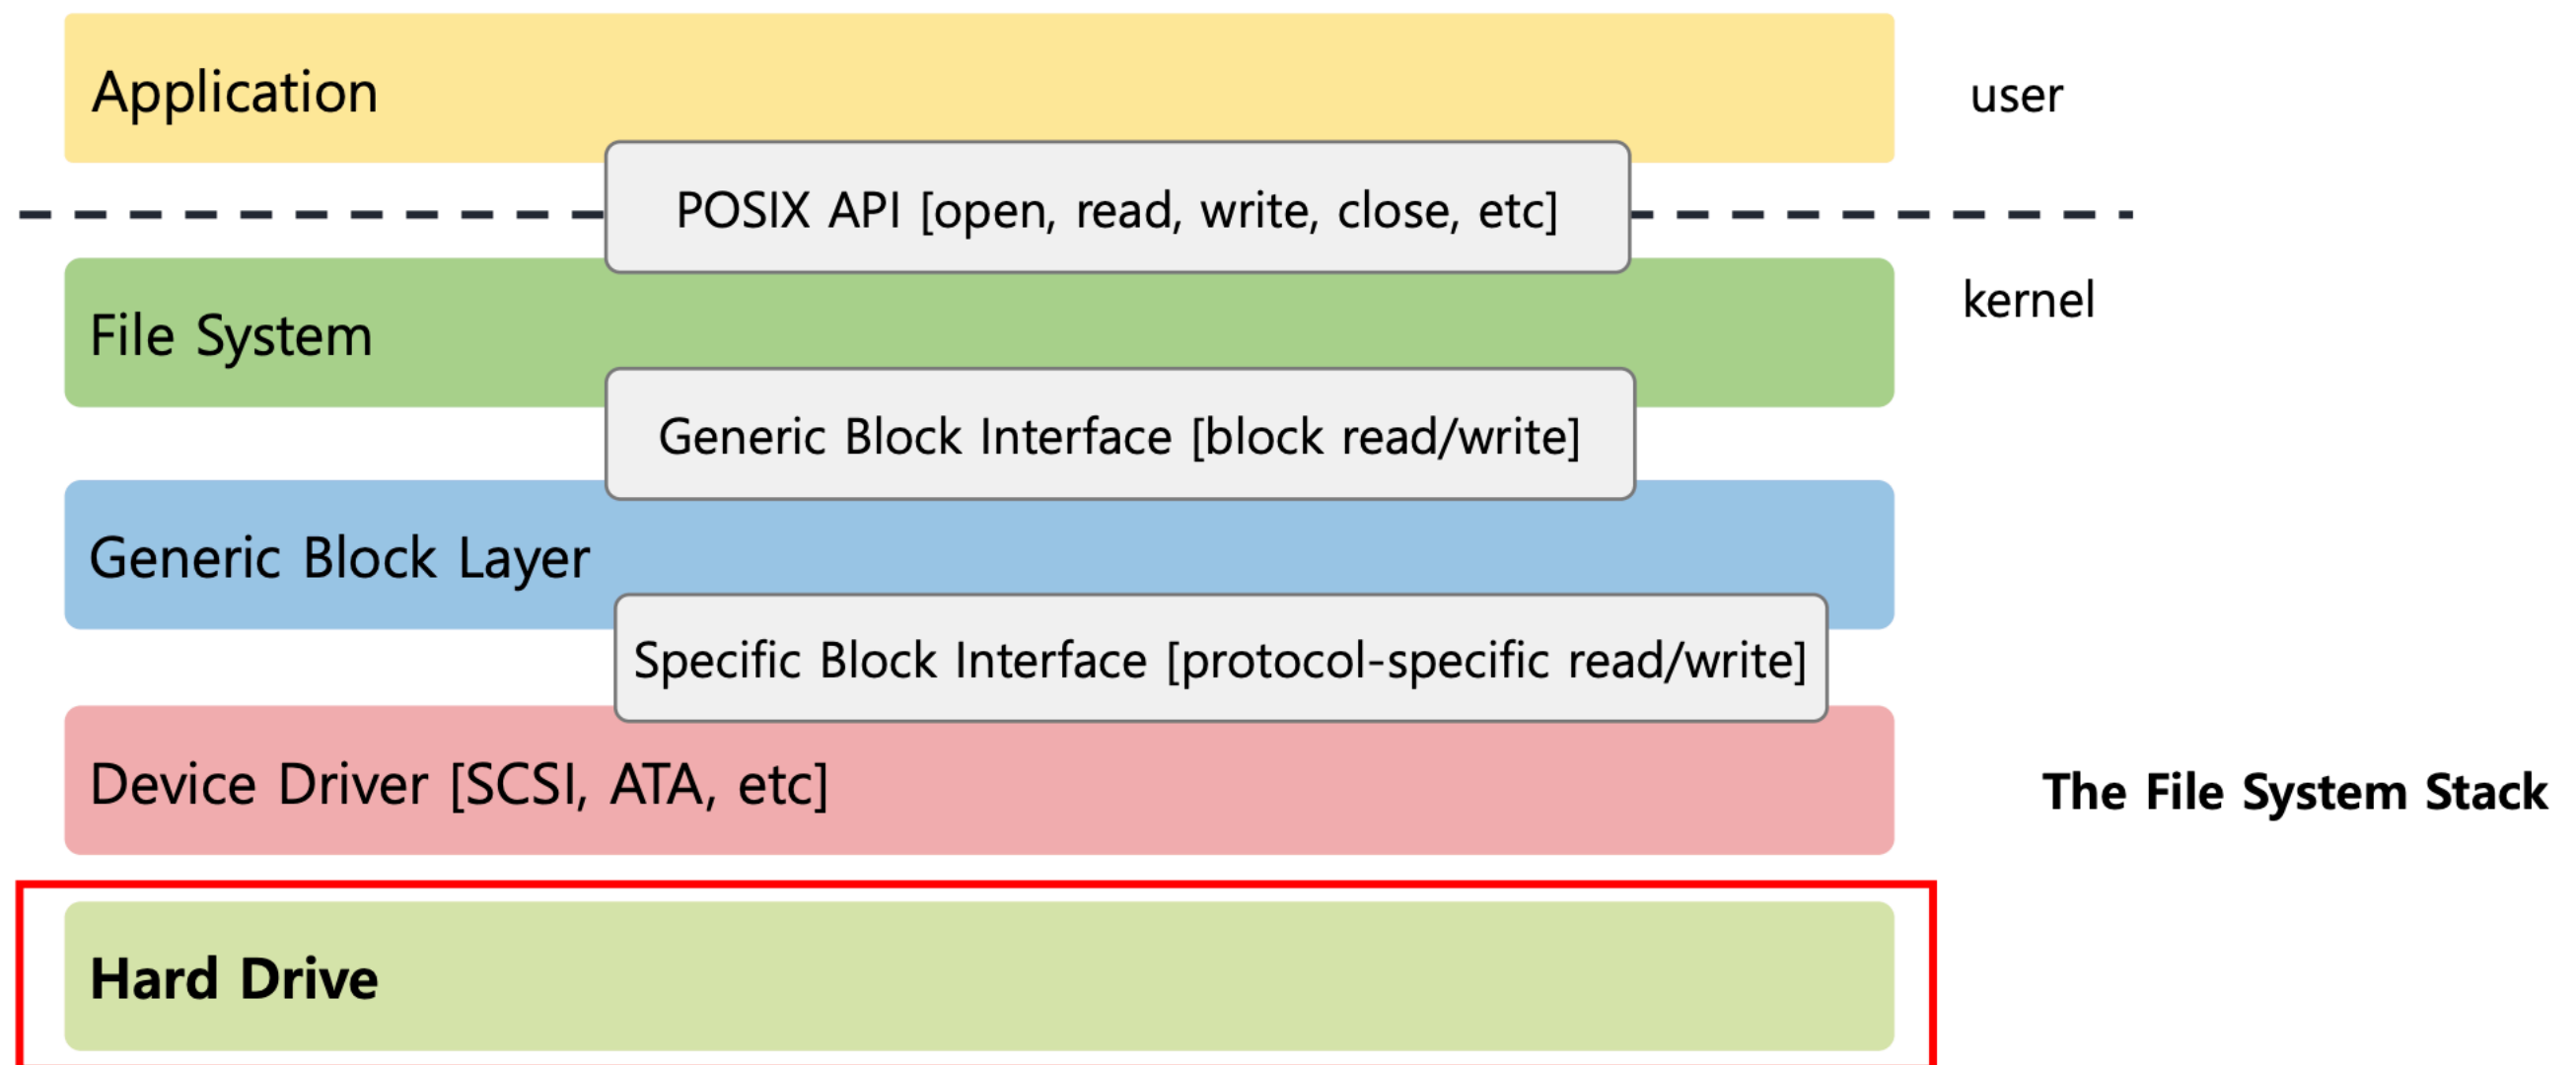
## ✓ Solution : abstraction

- Build a common interface
- Write device driver for each device

➡ Drivers are 70% of Linux source code

# File System Abstraction

- File system specifics of which disk class it is using  
It issues block read and write request to the generic block layer





# Disks

# Hard Disk Drive (HDD)

**Platter** (aluminum coated with a thin magnetic layer)

- A circular hard surface
- Data is stored persistently by inducing magnetic changes to it
- Each platter has 2 sides, each of which is called a surface

**Spindle**


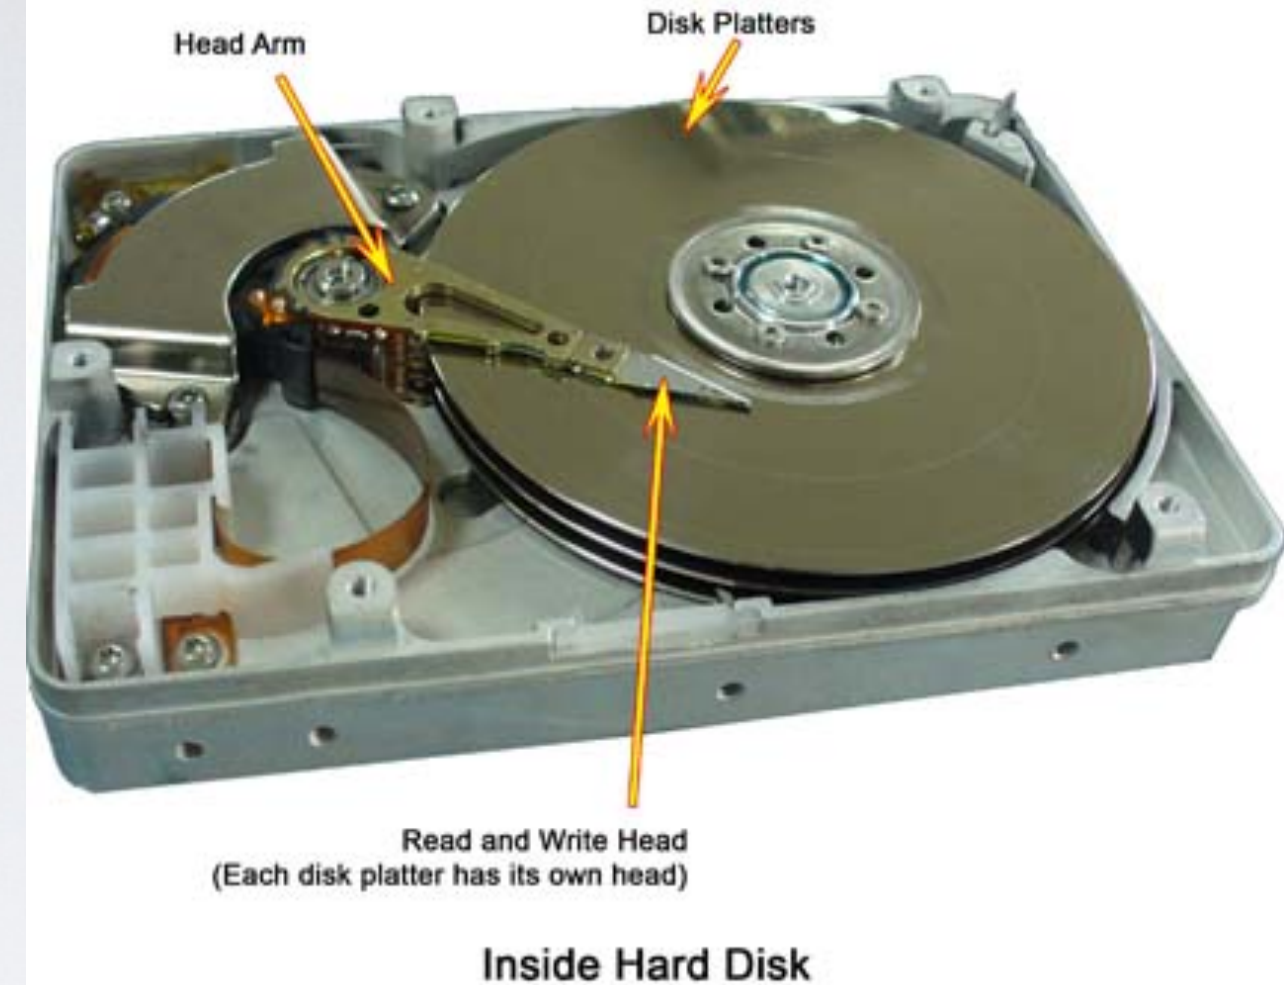
- Spindle is connected to a motor that spins the platters around
- The rate of rotations is measured in RPM (Rotations Per Minute)  
Typical modern values : 7,200 RPM to 15,000 RPM

**Track**

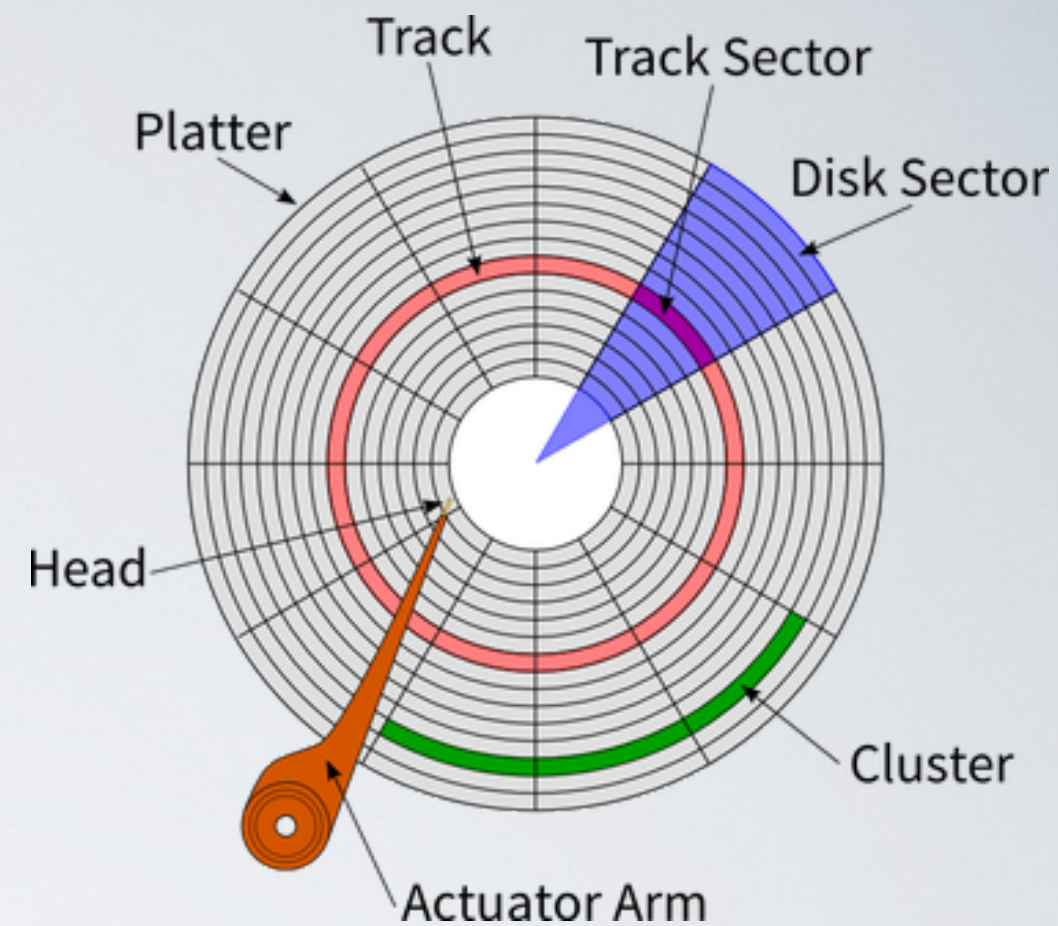
- Concentric circles of sectors
- Data is encoded on each surface in a track
- A single surface contains many thousands and thousands of tracks

**Cylinder**

- A stack of tracks of fixed radius
- Heads record and sense data along cylinders
- Generally only one head active at a time



# HDD Interface



- ➡ Disk interface presents linear array of sectors
  - Historically 512 Bytes but 4 KiB in "advanced format" disks
  - Written atomically (even if there is a power failure)
- ✓ Disk maps logical sector #s to physical sectors
- ✓ OS doesn't know logical to physical sector mapping



# Seek, Rotate, Transfer

Seek - move head to above specific track

1. speedup – accelerate arm to max speed
2. coast – at max speed (for long seeks)
3. slowdown – stops arm near destination
4. settle – adjusts head to actual desired track

- Seeks is slow

- settling alone can take 0.5 to 2ms
- entire seek often takes 4 - 10 ms



# Seek, **Rotate**, Transfer

Rotate disk until the head is above the right sector

- ➔ Depends on rotations per minute (RPM)  
With typical 7200 RPM it takes 8.3 ms / rotation
- Average rotation is slow (4.15 ms)

# Seek, Rotate, **Transfer**

Data is either read from or written to the surface.

- ➔ Depends on RPM and sector density  
With typical 100+ MB/s it takes  $5\mu\text{s}$  / sector (512 bytes)

✓ Pretty Fast

# Workload

So ...

- seeks are slow
- rotations are slow
- transfers are fast

What kind of workload is fastest for disks?

- Sequential : access sectors in order (transfer dominated)
- Random : access sectors arbitrarily (seek+rotation dominated)

➡ Disk Scheduler decides which I/O request to schedule next

- First Come First Served (FCFS)
- Shortest Seek Time First (SSTF)
- Elevator Scheduling (SCAN) commonly used on Unix

# Solid State Drive (SSD)

- ➔ Completely solid state (no moving parts), remembers data by storing charge (like RAM)
- ✓ Same interface as HDD (linear array of sectors)
- ✓ No mechanical seek and rotation times to worry about (SSD are way faster than HDD)
- ✓ Lower power consumption and heat (better for mobile devices)
- More expensive than HDD yet (but getting cheaper)
- Limited durability as charge wears out over time (but improving)
- Limited # overwrites possible
  - Blocks wear out after 10,000 (MLC) – 100,000 (SLC) erases
  - Requires Flash Translation Layer (FTL) to provide wear levelling, so repeated writes to logical block don't wear out physical block
  - FTL can seriously impact performance



# Acknowledgments

Some of the course materials and projects are from

- Ryan Huang - teaching CS 318 at *John Hopkins University*
- David Mazière - teaching CS 140 at *Stanford*