

Clustering in FFS

Access one block, probably access next

➡ Let's try to put sequential blocks in adjacent sectors


If you look at inode, most likely will look at data too

➡ Let's try to keep inode in same cylinder as file data

Access one name, frequently access many, e.g., “ls -l”

➡ Let's try to keep all inodes in a dir in same cylinder group

What Does Disk Layout Look Like Now?



How to keep inode close to data block?

- ➔ Use groups across disks and allocate inodes and data blocks in same group
- ✓ Each cylinder group basically a mini-Unix file system