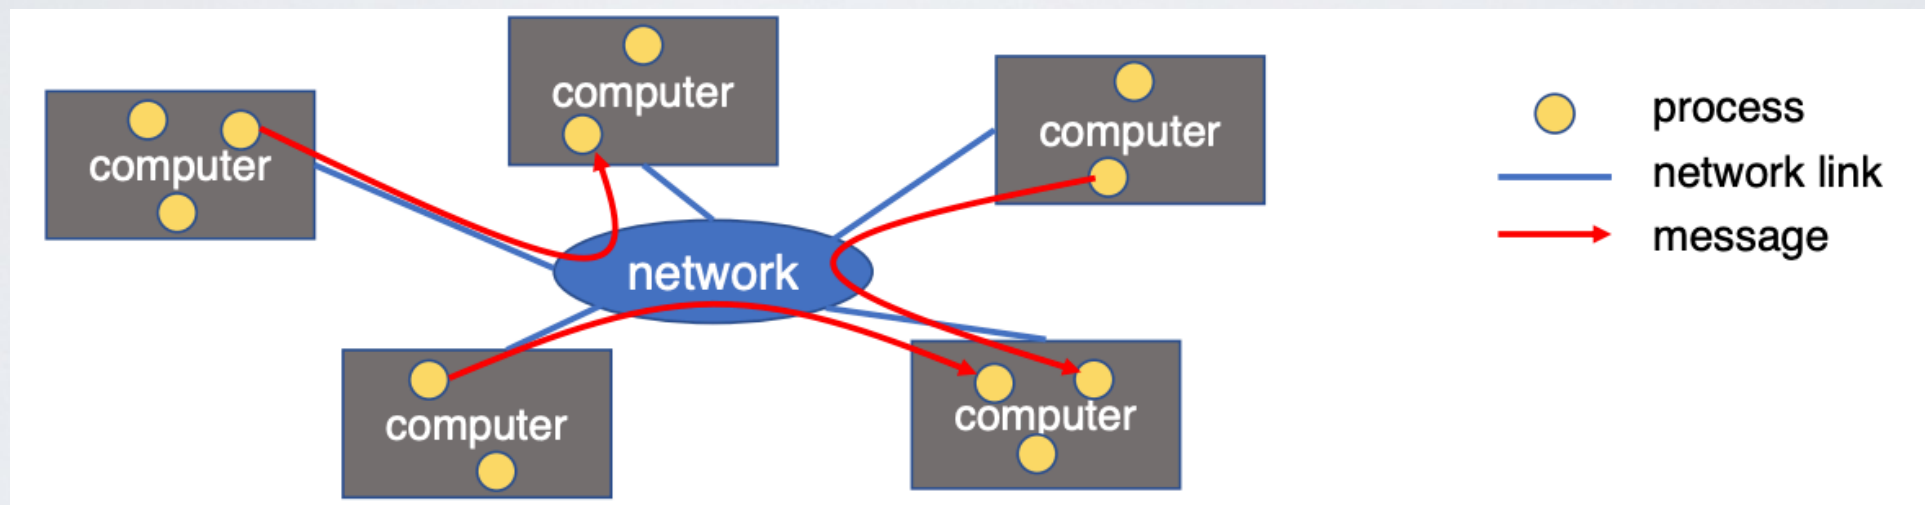


# Distributed Systems

Thierry Sans

# What is a distributed system?

➡ Cooperating processes in a computer network



"A distributed system is one where I can't do work because some machine I've never heard of isn't working!" *Leslie Lamport*



Popular distributed systems today:

Google file systems, BigTable, MapReduce, Hadoop, ZooKeeper, etc.

# Forms & models of distributed systems?

Degree of integration

- **Loosely-coupled**  
internet applications (e.g email, web, FTP, SSH)
- **Mediumly-coupled**  
remote execution (e.g RPC), remote file system (e.g NFS)
- **Tightly-coupled**  
distributed file systems (e.g. AFS)



# Why distributed systems?

Why do we want distributed systems?

- Performance - parallelism across multiple nodes
- Scalability - by adding more nodes
- Reliability - leverage redundancy to provide fault tolerance
- Cost - cheaper and easier to build lots of simple computers
- Control - users can have complete control over some components
- Collaboration - much easier for users to collaborate through network resources



# The promise of distributed systems

## The promise of distributed systems

- Higher availability - one machine goes down, use another
- Better durability - store data in multiple locations
- More security - each piece easier to make secure

# The reality of distributed systems

Reality has been disappointing

- Worse availability - depend on every machine being up
  - Worse reliability - can lose data if any machine crashes
  - Worse security - anyone in world can break into system
- Coordination is more difficult - must coordinate multiple copies of shared state information (using only a network)

# Requirements

**Transparency** - the ability of the system to mask its complexity behind a simple interface

Possible transparencies

- Location - cannot tell where resources are located
  - Migration - resources may move without the user knowing
  - Replication - cannot tell how many copies of resource exist
  - Concurrency - cannot tell how many users there are
  - Parallelism - may speed up large jobs by splitting them into smaller pieces
  - Fault Tolerance - system may hide various things that go wrong
- ➔ Transparency and collaboration require some way for different processors to communicate with one another


# Clients and Servers

The prevalent model for structuring distributed computation is the client/server paradigm

- ➔ A **server** is a program (or collection of programs) that provide a service (file server, name service, etc.)
  - The server may exist on one or more nodes
  - Often the node is called the server, too, which is confusing
- ➔ A **client** is a program that uses the service
  - A client first binds to the server (locates it and establishes a connection to it)
  - A client then sends requests, with data, to perform actions, and the servers sends responses, also with data



# Naming

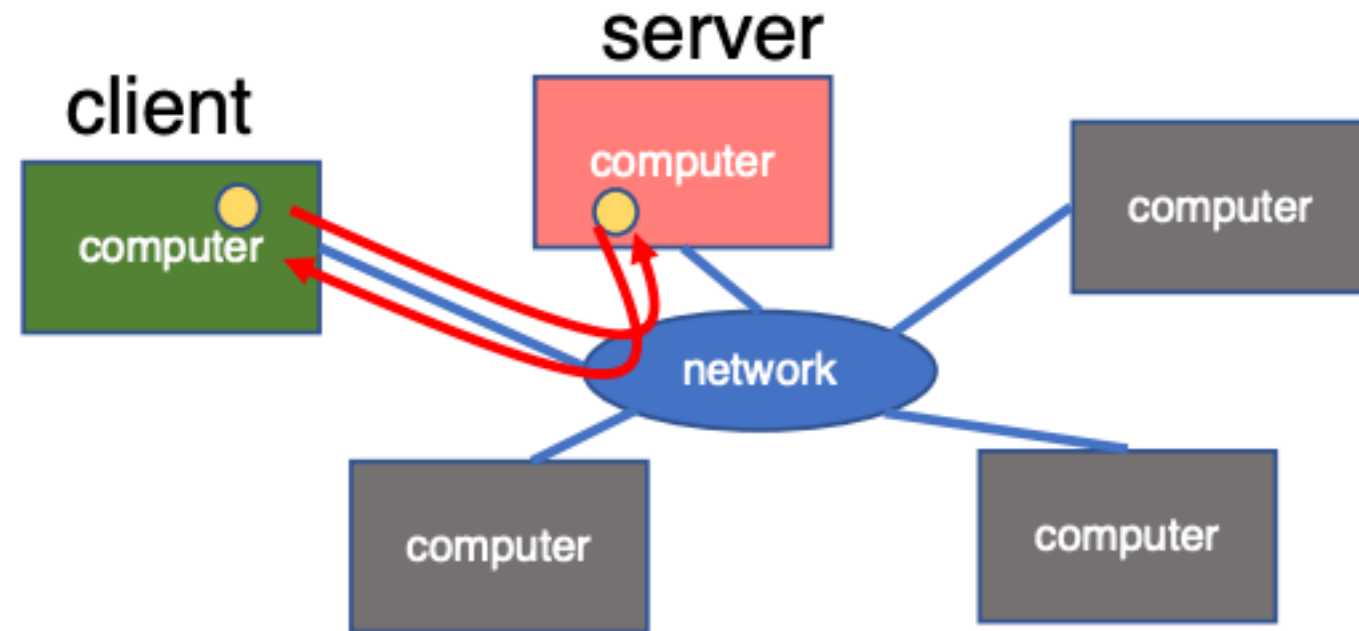


How to refer to a node in a distributed system?

Essentially naming systems in network

- Address processes/ports within system (host, id) pair
- Physical network address (Ethernet address)
- Network address (Internet IP address)
- Domain Name Service (DNS) provides resolution of canonical names to network address

# Communication



How can one computer communicate with another?

- Raw Message - UDP
- Reliable Message - TCP
- Remote Procedure Call (RPC)  
and Remote Method Invocation(RMI)

# Raw messaging

➔ Network programming = raw messaging (socket I/O)  
programmers hand-coded messages to send requests and responses

- Too low-level and tiresome

- Need to worry about message formats
- Must wrap up information into message at source
- Must decide what to do with message at destination
- Have to pack and unpack data from messages
- May need to sit and wait for multiple messages to arrive

Messages are not a very natural programming model

- Could encapsulate messaging into a library
- Just invoke library routines to send a message
- Which leads us to RPC...

# Procedure calls

Procedure calls are a more natural way to communicate

- Every language supports them
- Semantics are well-defined and understood
- Natural for programmers to use

➔ Idea - let servers export procedures that can be called by client programs

- Similar to module interfaces, class definitions, etc.
- Clients just do a procedure call as if they were directly linked with the server
- Under the covers, the procedure call is converted into a message exchange with the server



# Remote Procedure Calls (RPC)

So, we would like to use procedure call as a model for distributed (remote) communication

Lots of issues

- How do we make this invisible to the programmer?
- What are the semantics of parameter passing?
- How do we bind (locate, connect to) servers?
- How do we support heterogeneity (OS, arch, language)?
- How do we make it perform well?

# Why is RPC interesting?

**Remote Procedure Call (RPC)** is the most common means for remote communication

It is used both by operating systems and applications

- DCOM, CORBA, Java RMI, etc., are all basically just RPC
  - NFS is implemented as a set of RPCs
- ➔ Someday you will most likely have to write an application that uses some form of RPC for remote communication (or you already have)

# RPC example

## Client Program:

```
...  
sum = server->Add(3,4);  
...
```

## Server Interface:

```
int Add(int x, int y);
```

## Server Program:

```
int Add(int x, int y) {  
    return x + y;  
}
```

# RPC model

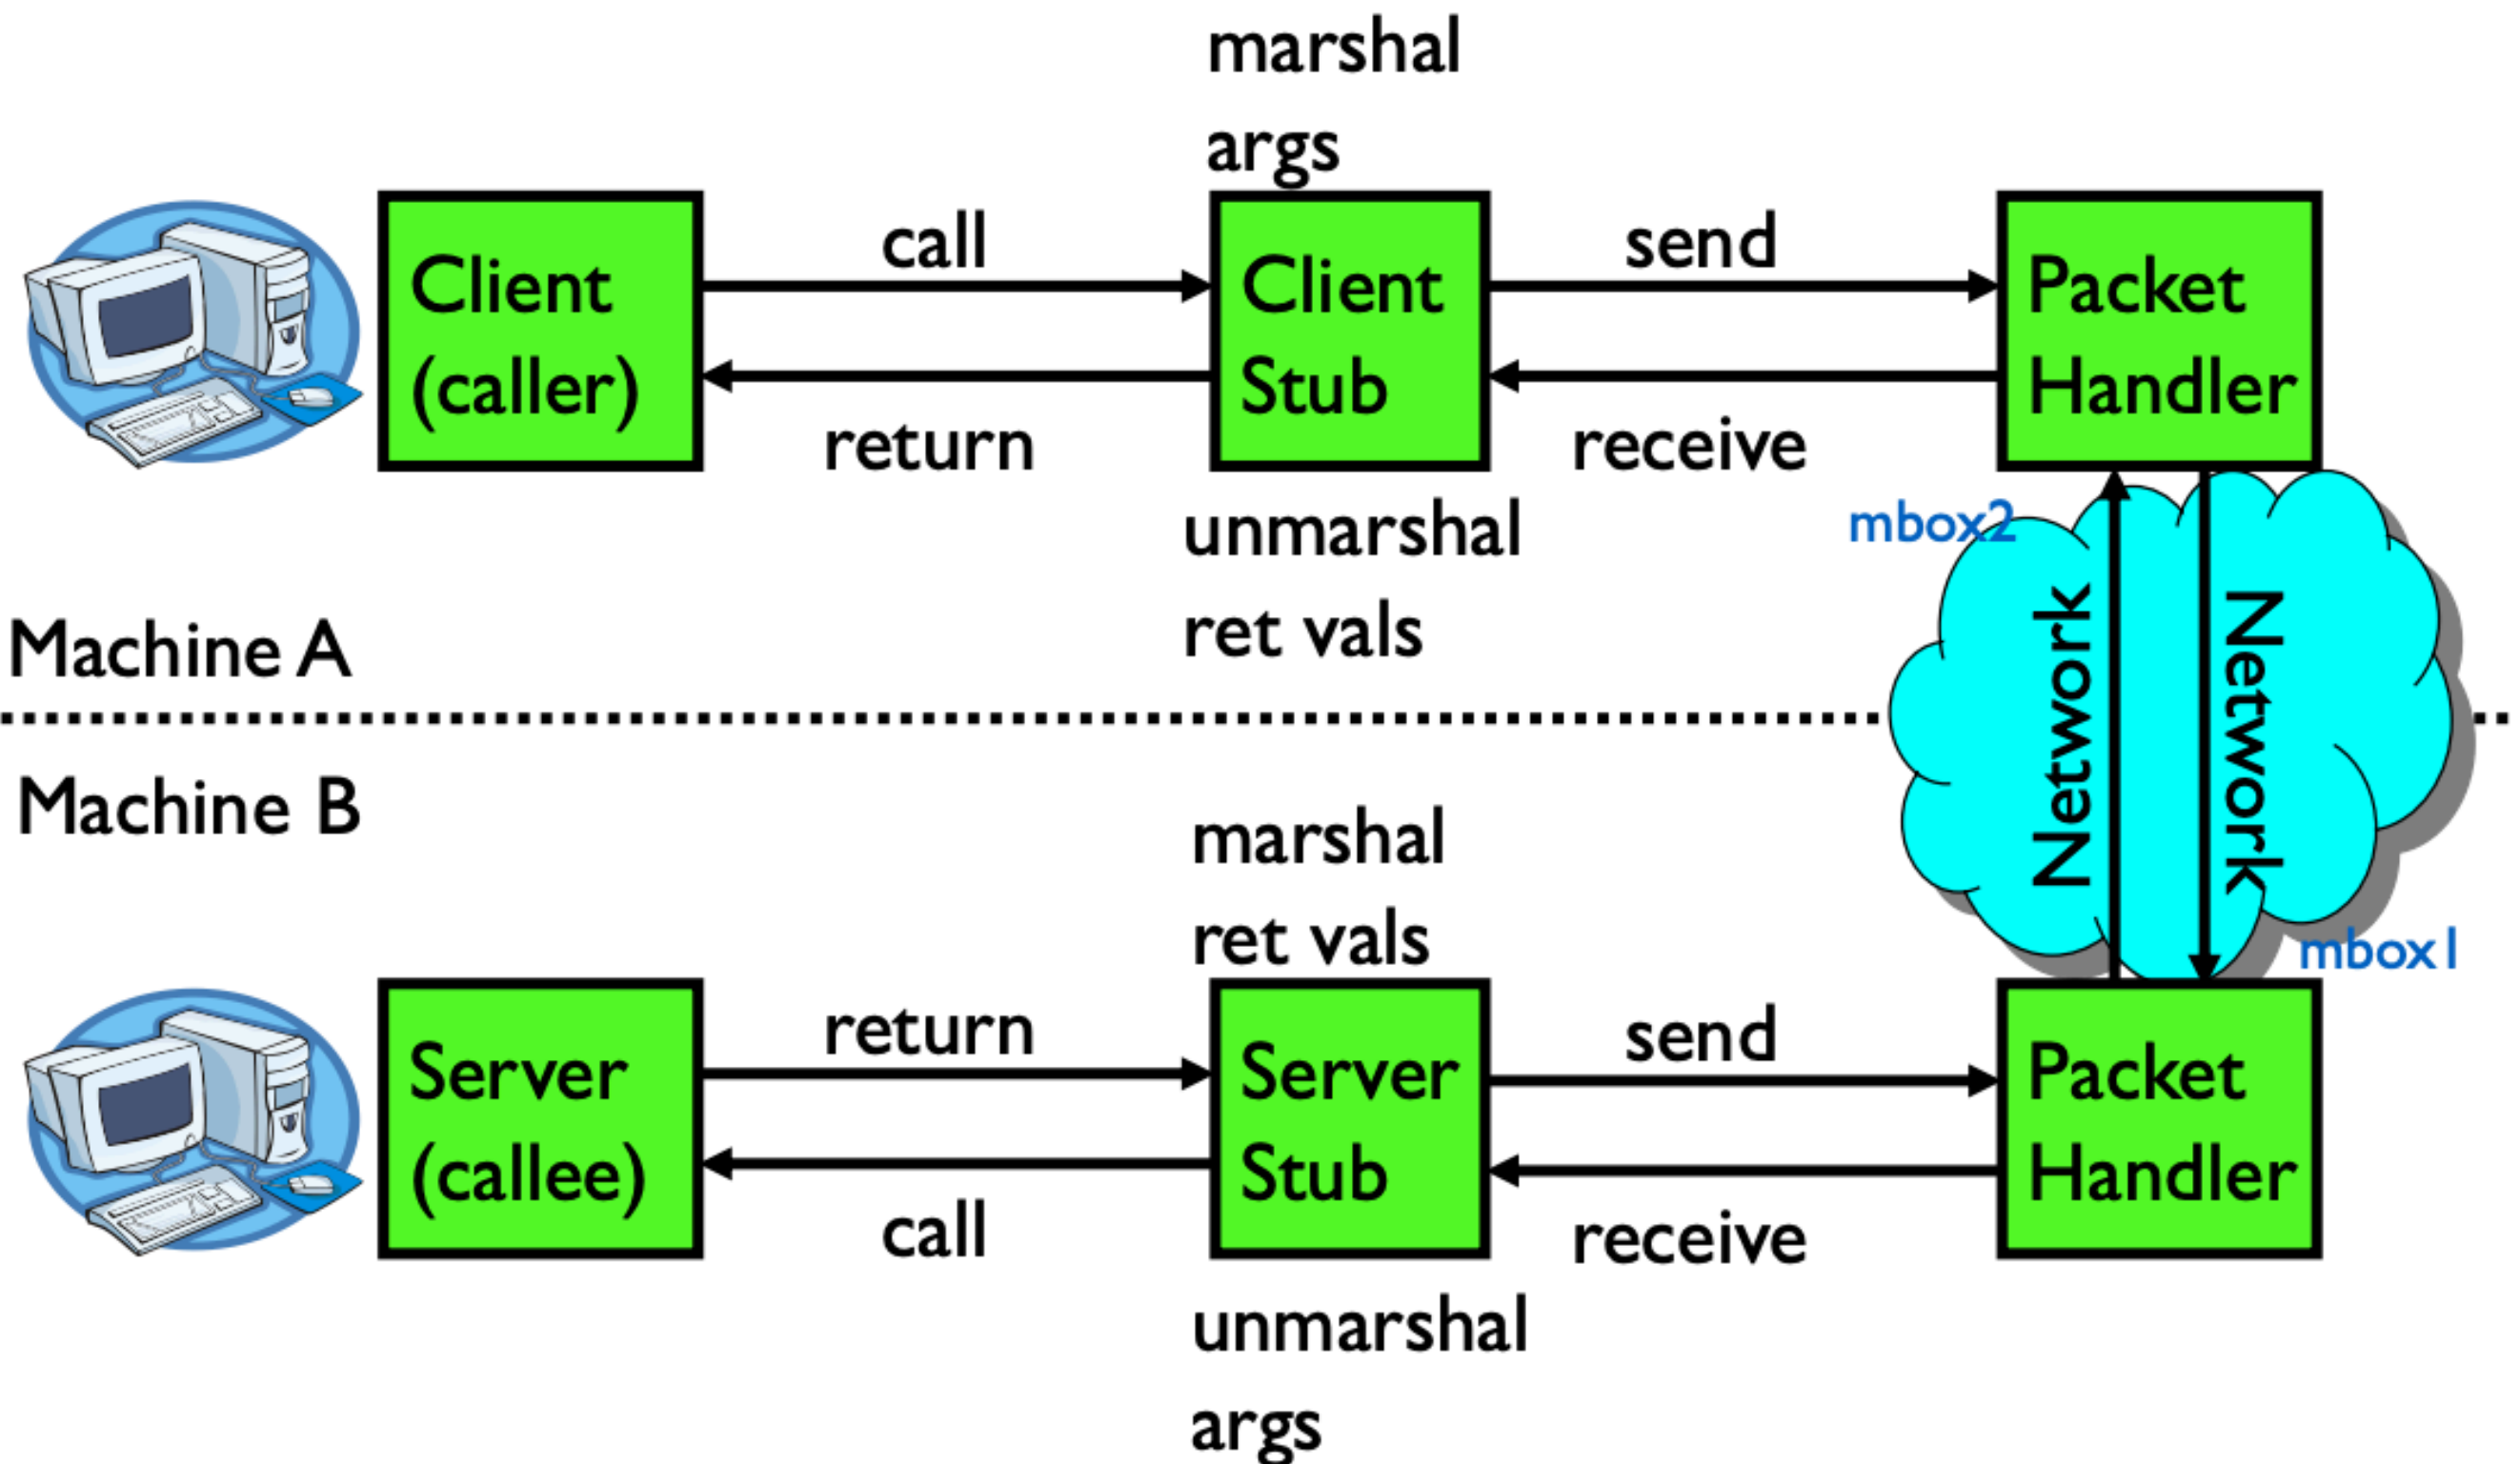
- ➡ A server defines the server's interface using an Interface Definition Language (IDL) that specifies the names, parameters, and types for all client-callable server procedures

A stub compiler reads the IDL and produces two stub procedures for each server procedure (client and server)

- Server programmer implements the server procedures and links them with server-side stubs
- Client programmer implements the client program and links it with client-side stubs
- ➡ The stubs are the “glues” responsible for managing all details of the remote communication between client and server



# RPC information flow



# RPC stubs

- ➡ The stubs send messages to each other to make RPC happen transparently
  - A client-side stub packs message, send it off, wait for result, unpack result and return to caller
  - A server-side stub unpack message, call procedure, pack results, send them off

# RPC marshallling

**Marshalling** is the packing of procedure parameters into a message packet

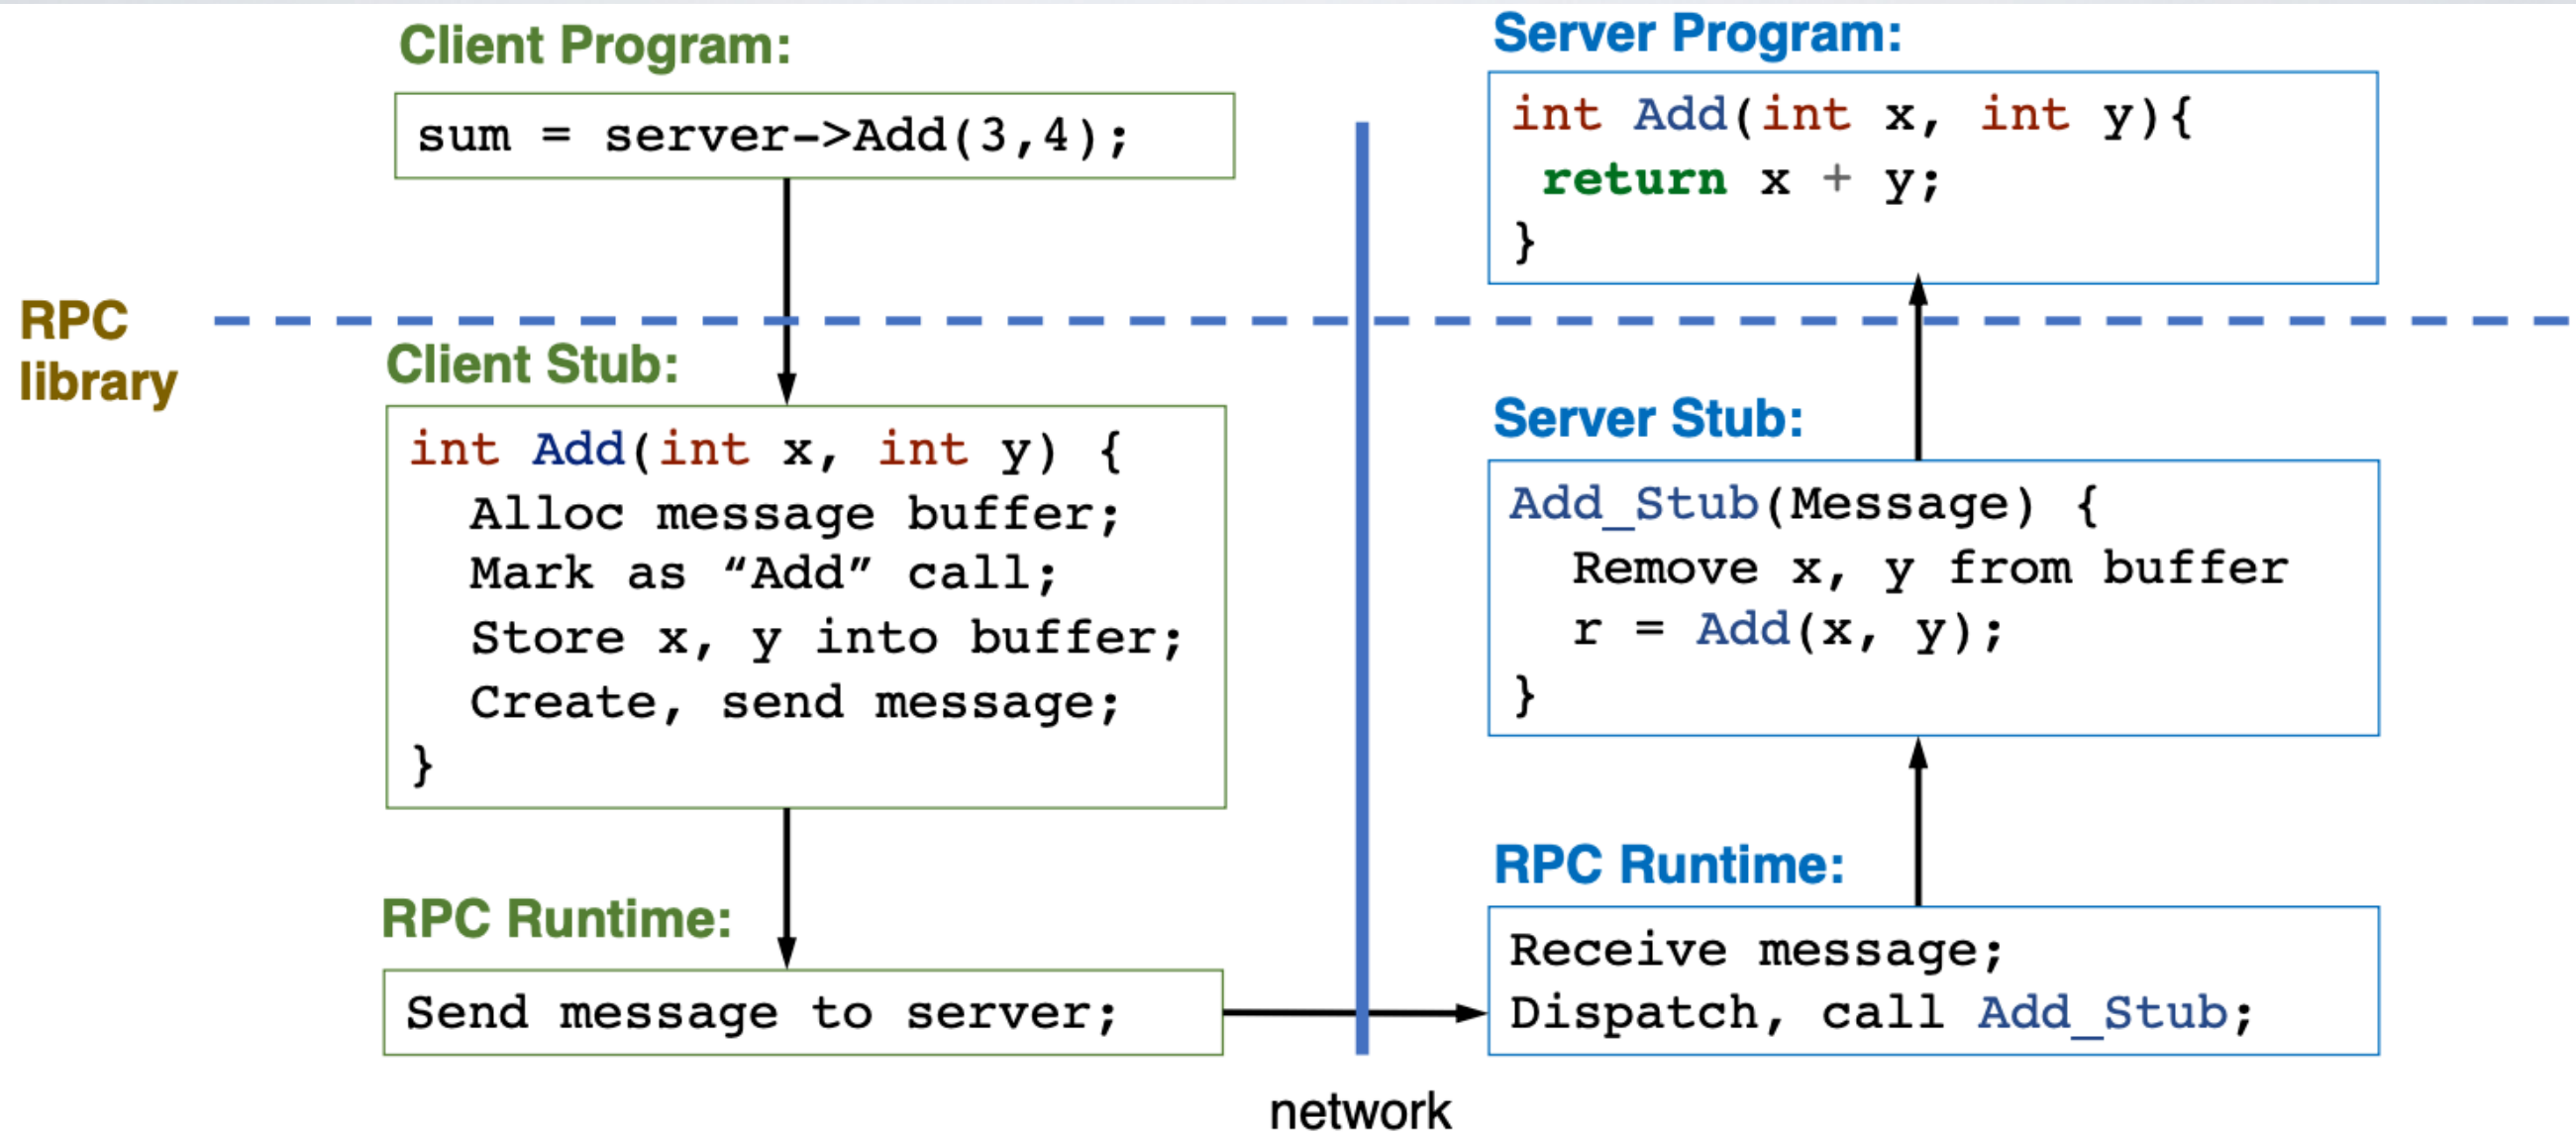
The RPC stubs call type-specific procedures to marshal (or unmarshal) the parameters to a call

- The client stub marshals the parameters into a message
- The server stub unmarshals parameters from the message and uses them to call the server procedure

On return

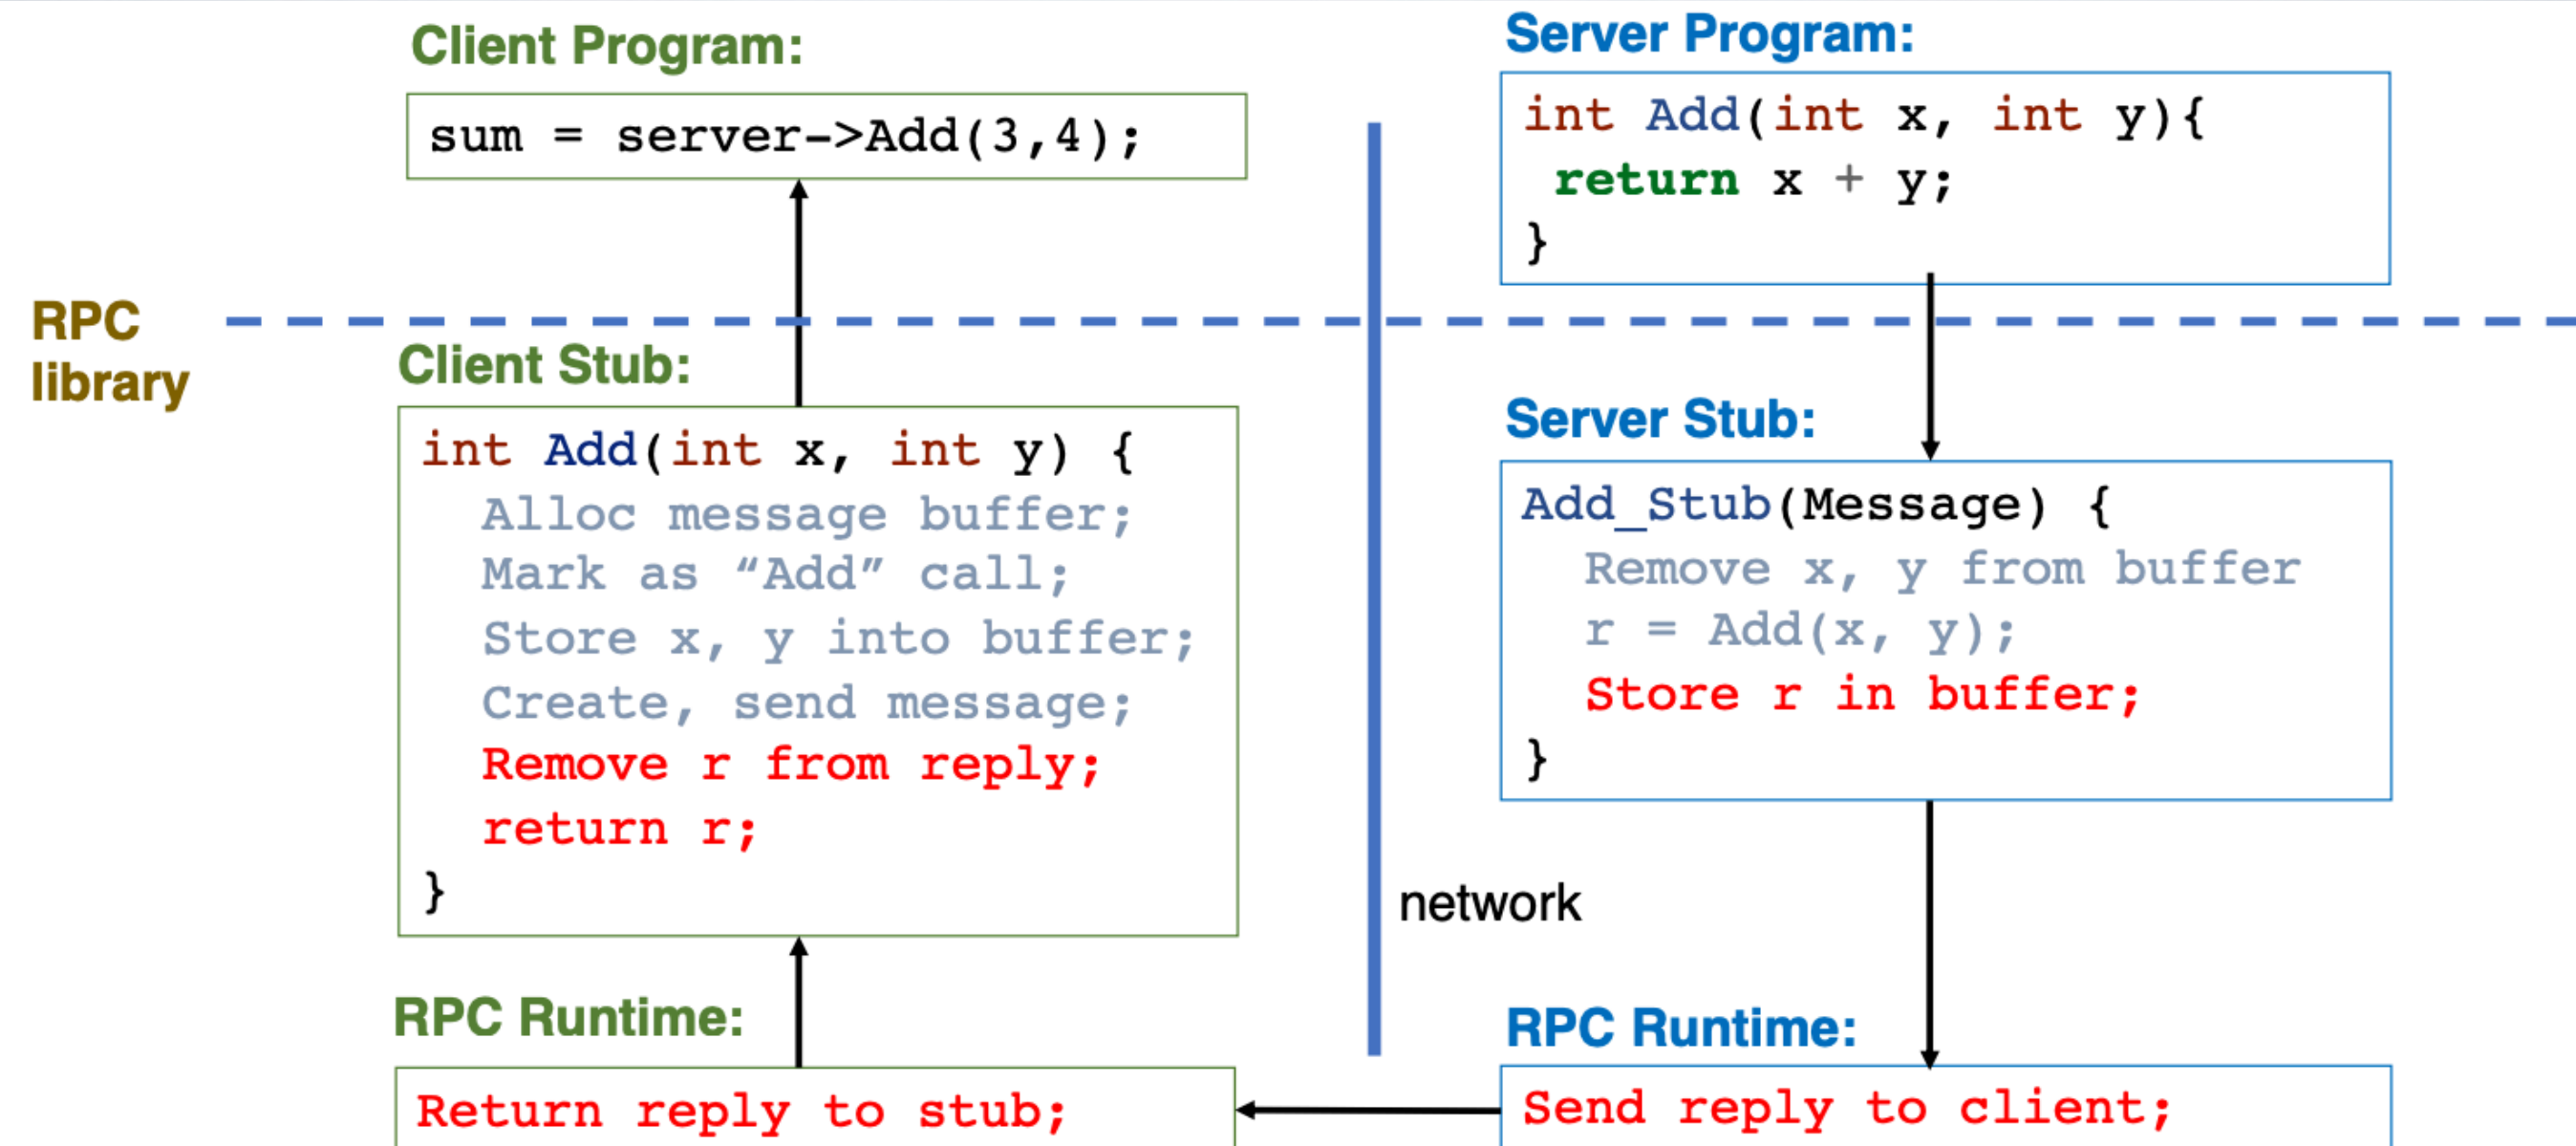
- The server stub marshals the return parameters
- The client stub unmarshals return parameters and returns them to the client progra

# RPC example - call





# RPC example - return



# RPC implementation details

## **What if client/server machines are different architectures and/or languages?**

Need to convert everything to/from some canonical form and tag every item with an indication of how it is encoded (avoids unnecessary conversions)

➔ Abstract Syntax Notation One (ASN.1)

## **How does client know which server to send to?**

Need to translate name of remote service into network endpoint (IP, port)

- ➔ Binding - the process of converting a user-visible name into a network endpoint
  - Static - fixed at compile time
  - Dynamic - performed at runtime

# RPC transparency

One goal of RPC is to be as transparent as possible

- ➡ Make remote procedure calls look like local procedure call although binding can break transparency

What else?

- Failures – remote nodes/networks can fail in more ways than with local procedure calls
- Performance – remote communication is inherently slower than local communication

# RPC failure semantic - at-least-once

What does a failure look like to the client RPC library?

- Client never sees a response from the server
- Client does not know whether the server processed the request

Simplest scheme - **at-least-once behavior**

- RPC library waits for response for time  $T$ , if none arrives, re-send the request
- Possibly repeat this a few times
- If still no response then return an error to the application



# RPC failure semantic - at-most-once

- Problem with at-least-once behavior

What if the request is "deduct \$100 from bank account" ?

➔ At-least-once works well with idempotent requests

Another (better) RPC behavior - **at-most-once**

- ➔ Having Server RPC code detects duplicate requests returns previous reply instead of re-running handler
- How to detect a duplicate request?
  - Client includes unique ID (XID) with each request, and uses the same XID for re-send
  - Server checks an incoming XID in a table, if an entry is found, directly returns the reply

# Problems with RPC - performance

Cost of Procedure Call  $\ll$  same-machine RPC  $\ll$  network RPC

➔ Means programmers must be aware that RPC is not free

# RPC summary

RPC is the most common model for communication in distributed applications

- Some popular libraries such as *gRPC*
- "Cloaked" as DCOM, CORBA, Java RMI, etc.

➔ RPC is essentially language support for distributed programming

# Acknowledgments

Some of the course materials and projects are from

- Ryan Huang - teaching CS 318 at *John Hopkins University*
- David Mazière - teaching CS 140 at *Stanford*