

Project I

Thierry Sans

Overview

```
devices/timer.c          |    42  ++++++-  
threads/fixed-point.h    |   120  ++++++++  
threads/synch.c          |    88  ++++++++  
threads/thread.c         |   196  ++++++++  
threads/thread.h         |    23  +++  
5 files changed, 440 insertions(+), 29 deletions(-)
```

➡ Most changes in threads and devices

➡ Look in lib/kernel for useful data structures: list, hash, bitmap

Pintos thread implementation

➔ Pintos implements user processes on top of its own threads

- Per-thread state in thread control block structure

```
struct thread {  
    ...  
    uint8_t *stack; /* Saved stack pointer. */  
    ...  
};  
uint32_t thread_stack_ofs = offsetof(struct thread, stack);
```

- C declaration for asm thread-switch function

```
struct thread *switch_threads (struct thread *cur, struct thread *next);
```

- Thread initialization function to create new stack

```
void thread_create (const char *name, thread_func *function, void *aux);
```

Desirable timeline

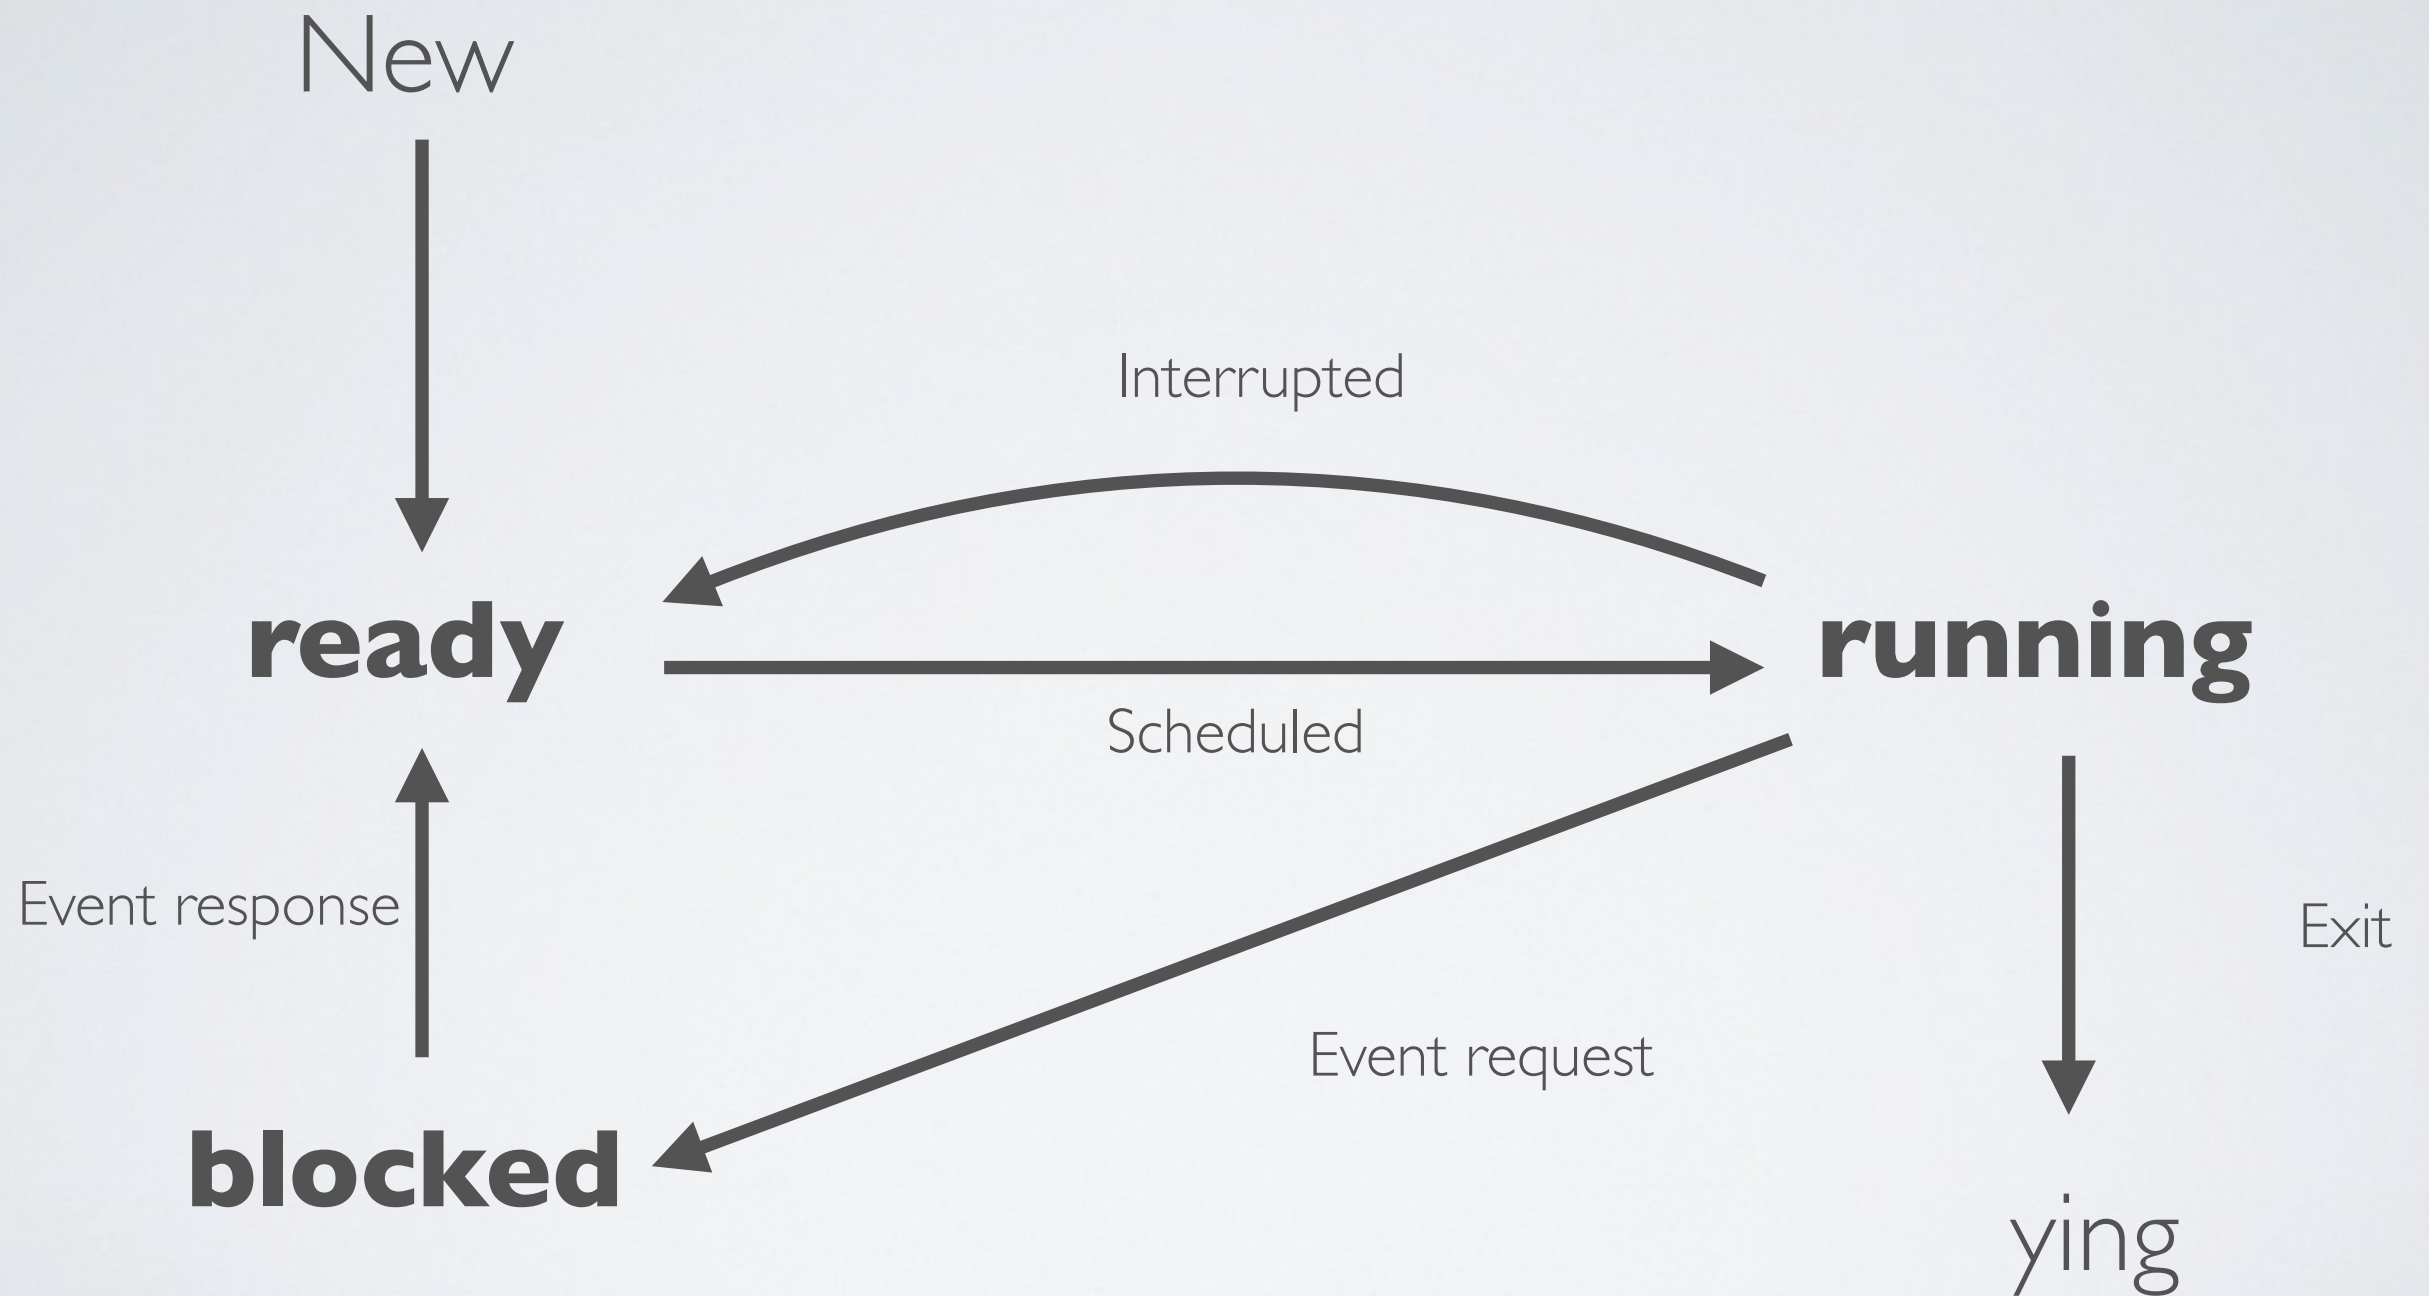
- Week 1 : From polling to interrupts
- Week 2 : Priority Scheduling
- Week 3 : Multilevel Feedback Queue Scheduler

From polling
to interrupts

Alarm Clock

- ➡ Reimplement `void timer_sleep(int64_t ticks)` in `devices/timer.c`
- ⦿ Existing implantation uses “busy waiting” (hardware polling)
- ✓ Suspends execution of the calling thread until time has advanced by at least `ticks` timer ticks (interrupt)

Thread States



Synchronization

Two techniques to serialize access to shared resource

1. **Disabling interrupts**

turns off thread preemption; only one thread can run

⦿ Undesirable unless absolutely necessary

2. **Synchronization primitives** (`threads/synch.h`)

- Semaphores
- Locks
- Condition variables

Priority Scheduling

Priority Scheduling

- ➔ Replace round-robin with priority-based scheduler
 - always run a thread with higher priority
 - yield immediately when higher priority thread is ready
- ◎ Priority inversion problem
- ◎ Starvation problem

Priority Inversion Problem

- ➔ A low priority threads (L) holds a resource needed by a higher priority thread (H)
 - H is blocked because L has locked the resource
 - L is blocked because a medium-priority thread (M) is running
- ✓ Fixed by priority donation

Priority donation

- ➔ A higher priority thread “donates” its priority to the lower priority thread it is blocked on
 - H "donates" its high-priority to L
 - When releases the resources lock, L returns to low priority
 - H runs immediately (since lower priority-threads should yield to higher-priority threads)

Multiple Priority Donation

If multiple threads need a resource, the priority of the thread holding the resource is the max of all priorities

➡ Effective priority is the max of donated priorities

Chained Priority Donation

- ➔ Donated priorities propagate through a chain of dependencies
 - H donates to M
 - M donates priority to L

Priority Scheduling

```
void thread_set_priority (int new_priority)
```

- Set the current thread's priority to new_priority
- Yield if the thread no longer has the highest priority
- If thread has donated priority, it still operates at the donated priority

```
int thread_get _priority ()
```

- Returns the current thread's priority
- With priority donation returns the higher (donated) priority

Multilevel Feedback Queue Scheduler

Principles

- ➔ Multilevel feedback queue scheduler tries to be fair with CPU time
 - No priority donation
 - Give highest priority to thread that has used the least CPU time recently
 - Prioritizes interactive and I/O-bound threads
 - De-prioritizes CPU-bound threads
- ✓ Details in section 2.2.4 and Appendix B

Priority

```
priority = PRI_MAX - recent_cpu/4 - nice*2
```

✓ Details in Appendix B.1

Nice

`nice` allows threads to declare how generous they want be with there own CPU time

➡ Integer value between -20 and 20

- `nice > 0`

lower effective priority, gives away CPU time

- `nice < 0`

higher effective priority, takes away CPU time from other threads

✓ Details in Appendix B.1

recent_cpu

recent_cpu : CPU time a thread has “recently” received

- Exponentially weighted moving average
- Incremented every clock tick when a thread is running
- Recomputed for all threads every second:

$$\text{recent_cpu} = (2 * \text{load_avg}) / (2 * \text{load_avg} + 1) * \text{recent_cpu} + \text{nice}$$

✓ Details in Appendix B.3

load_avg

load_avg : average number of ready threads in the last minute

- Single value system wide
- Initialized to zero
- Recomputed every second

$$\text{load_avg} = (59/60) * \text{load_avg} + (1/60) * \text{ready_threads}$$

✓ Details in Appendix B.4

Implementation

Add `-mlfqs` kernel option to allow the scheduling algorithm to be configured at startup time

➔ add to `parse_options()`

No priority donation

- `thread_set_priority()` should do nothing
- `thread_get_priority()` returns priority calculated by scheduler

✓ Details in section 2.2.4 and Appendix B

Acknowledgments

Some of the course materials and projects are from

- Ryan Huang - teaching CS 318 at *John Hopkins University*
- David Mazière - teaching CS 140 at *Stanford*